Bantam Java Compiler Project

Lab Manual

Extended Version 1.3

Marc L. Corliss	David Furcy	E Christopher Lewis
Hobart and William Smith Colleges	University of Wisconsin Oshkosh	VMware, Inc.
corliss@hws.edu	*furcyd@uwosh.edu*	*lewis@vmware.com*
http://math.hws.edu/mcorliss	*http://www.uwosh.edu/faculty_staff/furcyd*	*http://www.eclewis.net/*

About the Authors

Marc L. Corliss is an Assistant Professor in the Mathematics and Computer Science Department at Hobart and William Smith Colleges. In 2006, Professor Corliss received his PhD from the University of Pennsylvania in computer science. His research interests are in system design, and in particular the design of compilers and processors. He is also interested in computer science education and building new tools for teaching computer systems courses.

David Furcy is an Assistant Professor in the Computer Science Department at the University of Wisconsin Oshkosh. Professor Furcy obtained his PhD in computer science from the Georgia Institute of Technology in 2004. His research interests include artificial intelligence (heuristic search) and computer science education. Most recently, his passion for teaching has led him to investigate the design of effective algorithm and program visualization tools for education.

E Christopher Lewis is a Staff Engineer in the Advanced Development Group at VMware, Inc. He is exploring new technologies in the implementation and application of virtual machines. Before joining VMware in 2007, Dr. Lewis was a Professor in the Department of Computer and Information Science at the University of Pennsylvania. He received his PhD in computer science from the University of Washington in 2001.

Acknowledgements

Contents

1 Introduction

This manual describes a new programming language called Bantam Java and an infrastructure from which a student can build a Bantam Java compiler. The Bantam Java language and compiler are designed specifically for the classroom (*i.e.*, a compilers course). As the name suggests, Bantam Java contains a subset of the Java programming language, which is the language used in many computer science programs. As such, Bantam Java will seem familiar to most students. The compiler infrastructure, as well as the language, are designed so that an undergraduate student can complete the compiler in a one semester course.

Before describing the Bantam Java language and compiler, this manual motivates compiler construction for an undergraduate computer science student, which, unfortunately, some computer science programs have removed from their curriculum. There are many benefits in having an undergraduate student construct a compiler. First, constructing a compiler requires a deep understanding of the source programming language, the target machine, and everything in between. Very few aspects of computer programming seem like *magic* to a student after successfully implementing a compiler.

A compiler also makes for a challenging software engineering project. Modern compilers are incredibly complex. To deal with this complexity, most compilers make use of modular design, well-organized layers of abstraction, and design patterns such as the visitor pattern [13]. Implementing a compiler gives students experience applying all of these techniques. It also provides students with an opportunity to take part in designing a large software infrastructure.

A compiler is also a wonderful application of many concepts from other areas of computer science including programming languages, computational theory, data structures, algorithms, and complexity analysis. For instance, a compiler relies on finite automata and context-free grammars (usually covered in a computational theory course) for performing lexical and syntactic analysis. Building a compiler helps students appreciate the importance of these topics via their application in a practical context. In addition, a compilers course can help strengthen the understanding of those students who may have had some initial trouble with these often challenging concepts.

For all of these reasons, incorporating compiler construction into an undergraduate computer science curriculum is important. Unfortunately, there are very few compilers that are suitable for an undergraduate course project. Commercial compilers are too complex to be undertaken in a one semester course. At the other extreme, there are a number of programming languages and

compilers [4, 21], which do not include some of the important features of modern languages such as objects and inheritance. Although these languages and compilers abstract aspects of real programming languages, they appear toy-like (and possibly irrelevant) to students. The challenge is building a compiler infrastructure, which has most of the important features of modern programming languages, but which can be designed and implemented in a one semester course.

This manual presents a new programming language, Bantam Java, and a compiler infrastructure that are designed specifically for use in the classroom. The toolset will run on any Linux/x86 machine. Bantam Java (or Bantam for short) is a small subset of the Java language, making it reasonable to implement a compiler for it in a one semester course. (Bantam is a city in the north of the Indonesian island of Java and connotes smallness.) Although small, Bantam is sufficiently practical that it can easily be used to write any text-based application. The same cannot be said for many languages compiled by instructional compilers.

Bantam also has several other virtues that make it well-suited for an undergraduate compiler course. First and foremost, the project uses a language similar to Java, which is commonly used in introductory programming courses. This approach leverages the student's existing intuitions developed in earlier programming courses. The instructor also does not have to argue the relevance of the compiled language; the student instantly believes that the language is relevant and that the leap to a commercial compiler is in their grasp. Moreover, the insights that the student develops in constructing the compiler can be applied directly to Java.

A second virtue of this project is that several aspects of the project are customizable. For instance, this project supports several tools (*e.g.*, lexer and parser generators) and target architectures (*e.g.*, MIPS, x86, JVM). This project also allows instructors to choose between a base language with a small set of core features or an extended language. In addition, the components of the project can be easily swapped in and out so that an instructor could, for example, choose to skip lexing and parsing and instead have students implement a garbage collector. The project is also easily extended if an instructor wanted to add additional language features or targets.

A third virtue of this project is that it includes a free, comprehensive student manual (the document you are reading), which can be used in conjunction with any traditional compiler textbook [2, 22, 8, 19].

1.1 Project Goals

We had five goals in mind when designing the Bantam language and compiler.

Well-known language. Our first goal was to design a language that was a subset of Java, since Java is often used in introductory programming courses. We wanted to leverage the student's experience in earlier programming courses. The student spends less time learning the syntax and semantics of the language, and more time focusing on the implementation of the language, perhaps relearning some advanced language features, which they did not fully grasp the first time around. The student can also apply what they learn directly to Java. Finally, by using a Java-like language, the project as a whole is more compelling to the student.

Carefully-selected features. Of course, we could not include all or even most of the Java features in our language. Building a compiler for such a language would require more than a single semester. Our second goal was to select an appropriate subset of Java so that the project retained the character of Java yet resulted in a project that can be realized in a semester-long course.

We chose to emphasize object-oriented programming in our language leaving out such features as static methods and variables, interfaces, abstract classes, nested classes, packages, exceptions, arrays (at least in the base Bantam Java language), and most primitives (besides int and boolean). While some of these features add significant intellectual content to the compiler implementation, they all require a significant amount of time to implement.

For more advanced compiler courses, we provide an extended Bantam language and compiler infrastructure, which includes all of the features of the base project plus some additional useful features. In particular, the extended language includes (single-dimensional) arrays and some language constructs, which are not strictly required but are convenient such as for loops and break statements.

Well-engineered infrastructure. Our third goal was to build a well-engineered compiler infrastructure. We wanted instructors and students to be able to easily extend our infrastructure. By carefully designing our infrastructure, students can also complete more sophisticated assignments in less time. The compiler project also serves as an example of quality software engineering for students.

Comprehensive documentation. Our fourth goal was to provide a comprehensive, free manual (the document you are reading) describing all aspects of the project. This manual will work well

3

in tandem with any traditional compiler textbook [2, 22, 8, 19].

Customizable project. Our fifth and final goal was to give instructors the flexibility to tailor the project to their own classroom needs. For instance, instructors can choose between a base source language with a small set of core features or an extended language with some additional useful features. The base language includes many object-oriented features of Java, but in order to ensure that the project is doable in a one semester course, it does not include arrays as well as some useful language constructs (*e.g.*, for loops). In particular, the following features were added to the extended language:

- Single-dimensional arrays
- For loops
- Unlabeled break statements
- **++** and **--** increment/decrement operators (prefix and postfix)
- Arbitrary placement of return statements

In addition, there are two other ways that instructors can customize this project. First, instructors can choose between two sets of popular lexer and parser generators. Our infrastructure supports using either the JLex scanner generator [7] and the Java Cup parser generator (an LALR parser generator) [14], or the JavaCC scanner and parser generator (an LL(k) parser generator) [6].

There are also two installation packages to choose from: one that uses the **make** utility and another that uses the Apache **ant** tool. The **make** installation package is intended for a *nix system. The **ant** installation package will potentially install on any machine with pre-installed versions of Java and ant. In fact, this package has been installed on a Windows XP machine.

There are some optional assignments, which instructors can choose to incorporate into their courses. In one assignment, students build an optimizer for their compiler. In a second, students build a Bantam Java interpreter. In the future, we will add other optional projects such as a project in garbage collection.

Finally, the infrastructure supports three target machines: MIPS, x86, and the Java Virtual Machine (JVM). The MIPS target uses the SPIM emulator [15], which can be downloaded and installed on any x86 machine. The x86 target is a native target. It will work on any 32-bit, Intel x86 machine running the Linux operating system (although other x86 processors and Unix-based operating systems may work). The JVM target generates textual Java bytecode (or Java assembly code), which can be assembled into class files via the open source Jasmin Java assembler [18,

17]. In addition to these three targets, there is also a fourth target available to instructors only. Instructors can email the first author to obtain a version of the compiler targeting the educational architecture Larc [10, 9], which can be as a resource in an introductory computer architecture course.

Summary. We made the Bantam language a subset of an existing, ubiquitous programming language, Java, so that students would already be familiar with it, understand the relevance of the course projects, and be able to apply existing intuitions. Because the Bantam compiler had to be fairly simple, we could only choose a small subset of the Java features to include in Bantam. We chose to include those features that emphasize object-oriented programming. We use many tools from software engineering to make our infrastructure easy to understand and extend. We have written a comprehensive manual (this manual), which can be used in concert with a traditional compiler textbook [2, 22, 8, 19]. Finally, this project has some customizable components (language features, lexer/parser generator, target machine) that allow the instructor to tailor the project to their own classroom needs.

1.2   Resources

There are several resources for helping both students and instructors install and work with this toolset. They are all freely available online at *http://www.bantamjava.com*. These include:

- **Student manual.** A comprehensive student manual (this manual), which documents all aspects of the Bantam Java compiler project (a PDF document).

- **SIGCSE paper.** A condensed description of this work (a PDF document) published in SIGCSE 2008 [11]. The slides from the SIGCSE talk are also available (a PDF or Power-Point document).

- **Bulletin board.** An online bulletin board for submitting questions, comments, and requests.

- **Compiler API.** The application programming interface (API) of the compiler generated from comments in the source code using the javadoc utility.

In addition, instructors (and instructors only) can obtain solution code for the project as well as a version of the compiler targeting the educational architecture Larc [10, 9] by emailing the first author of this manual.

1.3   About This Manual

This manual along with the corresponding toolset provide support code and documentation for course projects in a compilers course. This manual should be supplemented with a traditional compiler textbook [2, 22, 8, 19]. For example, this manual does not discuss the theory of lexing and parsing, nor does it discuss the compiler issues that arise when compiling languages besides Bantam. It also assumes that the student is following the design template laid out in the subsequent sections. It will offer limited help for those following a different design recipe.

In addition, this compiler project and this manual are intended for students who already know some Java. It does not teach programming or programming in object-oriented languages, although it does discuss some of the concepts when describing the Bantam Java language. Students should be familiar with Java or at the very least familiar with a similar object-oriented language (*e.g.*, C++, C#). For those students who are less familiar with Java, they may need to consult a Java reference book [5], while reading this manual.

There are two versions of this lab manual: one that uses the Bantam Java base language and one that uses the Bantam Java extended language (which includes additional features such as arrays). **You are reading the extended language version.** If you would like the base language version, it is available online at *http://www.bantamjava.com*. There are also two versions of the toolset: a base version and an extended version. You should use the same version of the toolset as the manual. Either version of the toolset can be downloaded at *http://www.bantamjava.com*.

1.4   Related Work

The Bantam project is not the first classroom compiler project, but it does have some advantages over its predecessors. There are a number of compiler projects [4, 21], which do not use object-oriented source languages. The major drawback with these projects is that the dominant programming paradigm today is object-oriented programming. Most of the recent successful commercial languages are object-oriented (*e.g.*, Java, C#, C++, Python).

The COOL (Classroom Object-Oriented Language) project by Alex Aikens [3] is one example of an object-oriented educational compiler project. This project uses a source language called COOL, a simple object-oriented language, which the authors are able to formally prove is type safe. While this project is appealing for a compiler course that emphasizes type theory, the language itself is syntactically and semantically different from any other language that students are familiar with. Students must spend significant time learning the COOL language before beginning the compiler project. Students may also have difficulty applying what they have learned to practical languages like Java.

The MiniJava project [22] is another educational, Java-based compiler project. The MiniJava project is integrated into a particular compiler textbook [22], reducing the project management for the instructor, but this tight integration makes the project an impractical option for instructors wishing to employ a different text. In addition, the object-oriented aspects of the project appear in the "Advanced Topics" section of the text, relegating object-orientation to an optional add-on. Finally, the publisher-provided instructor code is incomplete. The Bantam Java project can be used with any textbook, object-orientation is fundamental, a complete implementation is available to instructors, and it has a complete and clear student manual.

Although Bantam Java is significantly different from Cool and MiniJava, and has the virtues discussed above, many ideas were borrowed from both projects. First, the context-free grammar for Bantam Java is specified in a similar syntax to the Cool context-free grammar [3]. The encoding of objects, object construction, and parts of the runtime system also work in a similar way as in Cool [3]. The Bantam Java compiler makes use of the visitor design pattern as in MiniJava [22].

1.5   Outline

The remainder of this manual is organized as follows. Section 2 describes the Bantam Java toolset and provides instructions for installing the toolset. Section 3 describes in detail the Bantam Java language. Section 4 describes the (incomplete) Bantam Java compiler (minus the optimization component), which students will complete in course projects. Section 5 describes the Bantam Java runtime system, which will be used by the generated code. Section 6 describes the (incomplete) optimization component in the Bantam Java component, which students can complete in an optional course project. Section 7-10 contain the core project assignments and descriptions for building the lexer, parser, semantic analyzer, and code generator. Section 11-12 contains some

7

optional project assignments. Section 11 contains a project assignment for building an optimizer. Section 12 contains a project assignment for building a Bantam Java interpreter.

2   Bantam Java Toolset

This section gives an overview of the Bantam Java toolset as well as instructions for installing and running the compiler. The toolset can be built using either the Unix **make** utility or the Apache **ant** tool. The next two subsections describe how to install the toolset using **make** and **ant**, respectively. You can safely skip the subsection corresponding to the building tool that you are not using. However, even if your toolset is already pre-installed, you should still read (or at least skim) one of these subsections, as well as the following subsections, as they describe how the project infrastructure is layed out and how to run the Bantam Java compiler.

2.1   Overview and Installation Instructions using Make

The Bantam Java toolset will compile and should run on any linux/x86 machine, although it is primarily tested on an Intel 32-bit x86 machine running Suse, Red Hat, and Ubuntu Linux. It also may run on other *nix platforms and other architectures. It is available as a compressed tarball (bantamjava.tar.gz) at the website: *http://www.bantamjava.com/*. Note: there are two versions of the toolset: one for building a compiler for the base Bantam Java language and one for building a compiler for the extended language. Make sure that you have downloaded the correct toolset. The version of the toolset that you use should be the same as the version of this manual (you are reading the **extended** version).

Once you have downloaded the tarball, move it to the appropriate location. To uncompress and untar the tarball, use the following:

    bash$ tar xvfz bantamjava.tar.gz

This command will create a new directory, **bantamjava/** (the *root directory*), in the current location. The **bantamjava/** directory contains the following files and directories (note: directories end with '/', files do not):

    README   Makefile   api/   bin/   lib/   man/   src/   tests/   tools/

Each of these files and directories are described below (although not in order).

README. The **README** file provides installation documentation as described in this chapter. It also indicates which version of the toolset you have downloaded (base or extended).

bin. The **bin/** directory contains executable files for running the compiler and other auxiliary tools. This directory is initially empty. Executable files are automatically added to **bin/** when the toolset

9

is built.

lib. The **lib/** directory contains library files needed for building the compiler and also running auxiliary tools (*e.g.*, Spim). This directory is initially empty. Library files and directories are automatically added when the toolset is built.

man. The **man/** directory contains man pages for the Bantam compiler as well as some of the auxiliary tools (*e.g.*, Spim), although the man pages for the auxiliary tools are added only after the toolset is built.

Makefile. The **Makefile** will compile and install all of the components of the toolset. There are basically five parts of the toolset: the Javadoc web documentation describing the Java class files of the compiler (**api/**), the source files for the (incomplete) compiler (**src/**), the bantam test programs for debugging the compiler (**tests/**), and various tools used by the compiler (**tools/**). These components are described in more detail below.

To use the **Makefile**, you will need to edit one line of the **Makefile**. The definition of the variable **PROJPATH** (near the top of the **Makefile**) will probably be incorrect. Set this variable to the path of the installation directory. You may also want to edit a line just below the definition of **PROJPATH** that sets whether the compiler will be installed using the JLex and Java Cup lexer and parser generators or the JavaCC lexer and parser generator. The variable **VERSION** indicates which version will be installed. If it is set to "bantam" then JLex/Java Cup is used (the default), if it set to "bantam-jj" then JavaCC is used.

To make the entire toolset, type either "make" or "make all" in the root directory. You can also make an individual component as discussed below. To uninstall the toolset, type "make clean" in the root directory or in any of the subdirectories, which contain a **Makefile**. Note: if the compiler is unimplemented (as it is when first installed) "make" will result in an error since the last thing the **Makefile** attempts to do is compile the test programs using the Bantam Java compiler. However, if installed correctly, this error should occur only at the point that the **Makefile** attempts to compile a test program, and not at any point before.

api. The **api/** directory houses the Javadoc application programming interface documentation. After uncompressing and untarring the toolset, the directory will contain only a **Makefile**. The documentation is created from comments in the source code using the **javadoc** command provided with the Sun JDK [20]. It is placed in a directory **html/** within the **api/** directory. After generating the API (from the **Makefile**), it can be viewed online by copying this directory to a website. For

10

example, if **/var/www/** is a website directory viewable at *http://www.foo.com/*, then the following commands would make the API viewable at *http://www.foo.com/bantam-api/*.

```
bash$ make api
bash$ cp -r api/html/ /var/www/bantam-api
```

When making changes to the compiler, Javadoc comments should be maintained so that any changes will be reflected in the API documentation. Sun Microsystems provides a free tutorial describing Javadoc commenting [20]. The API is also available online at *http://www.bantamjava.com/*.

tools. The tools directory houses external software required by the Bantam Java toolset. These tools include JLex, a lexer generator [7]; Java Cup, a parser generator [14]; JavaCC, a lexer and parser generator [6]; and SPIM, a MIPS emulator [15]. These tools do not come with the toolset, but must be downloaded and moved to the **tools/** directory. Links to the respective websites are included online at *http://www.bantamjava.com*. These tools are built by typing "make tools" in the root directory or "make" in the **tools/** directory. This manual does not describe these tools (at least the auxiliary tools) in much detail. For more information on these tools, check out their respective manuals [7, 14, 6, 15].

Although the **tools/** directory does not initially contain the external tools, it does contain some code. First, it contains a directory **lib/** that holds two Jar files, which contain reference Bantam Java compilers for each version of the compiler: one that uses JLex and Java Cup (**bantamc-ref.jar**) and one that uses JavaCC (**bantamc-jj-ref.jar**). The **tools/lib/** directory also holds a runtime library file (containing assembly code) for two of the supported targets: MIPS/SPIM (exceptions.s) and x86/Linux (x86-runtime.s). In addition, it holds Java class files (**TextIO.class** and **Sys.class**), which serve as the runtime library when targeting the Java Virtual Machine (JVM). (Note: other built-in objects like **Object** and **String** as well as support for garbage collection need not be included in the JVM runtime library as they are already included in the JVM, itself.) In addition to **lib/**, the **tools/** directory also contains a **Makefile** that will install each tool once they are downloaded, untarred and uncompressed, and moved to the **tools/** directory.

In order for the **Makefile** to work properly, you must use the following names for each tool: jlex, javacup, javacc, and spim. For JLex, JavaCC, and Spim you must simply unpack the source files in **tools/** and rename each with the appropriate name (*i.e.*, jlex, javacc, spim). For Java Cup, the latest jar file should be downloaded and put in the directory **tools/javacup**. The jar file should be called **java-cup.jar**.

11

Like the root directory **Makefile**, to use the **Makefile** in **tools/** you will need to edit one line of the **Makefile**. The definition of the variable, **PROJPATH**, on line 2 will probably be incorrect. Set this variable to the path of the installation directory. Once the tools have been downloaded and moved to the correct location, type 'make' to install them (Note: if compiled from the root directory the **PROJPATH** variable does not have to be changed. This change is only necessary when compiling from the sub-directory.)

Note that the tools do not have to necessarily be installed within the **tools/** directory. They can be installed in other locations. However, a shell script or binary for running each tool must be installed in the **bin/** directory within the root directory. In addition, the **lib/** directory within the root directory must contain a link called **JLex** to the JLex directory within the installation of JLex. Also, the **lib/** directory must contain a Java Cup jar file called **java-cup.jar** for running Java Cup. Finally, some files within **tools/lib/** must be copied to **lib/**. These include the reference compilers (**bantamc-ref.jar** and **bantamc-jj-ref.jar**) and the runtime library assembly files for MIPS/SPIM (**exceptions.s**), x86/Linux (**x86-runtime.s**), and JVM (**TextIO.class** and **Sys.class**).

src. The source files for the compiler are contained within **src/**. The Bantam compiler is written entirely in Java. Of course, the compiler is incomplete and will need to be completed by the student (although the code will compile). This directory contains two directories, **bantam/** and **bantam-jj/**. The **bantam/** directory contains the source files that use the JLex lexer generator [7] and the Java Cup parser generator [14]. The **bantam-jj/** directory contains the source files that use the JavaCC lexer and parser generator [6]. The **bantam-jj/** directory uses some symbolic links to the bantam directory to prevent duplication of code. Students should implement the compiler within one of these directories (but probably not both) as specified by the course instructor.

Within each directory (**bantam/** and **bantam-jj/**) are the following files and directories:

    Makefile   Main.java   ast/   cfg/   codegenjvm/   codegenmips/   codegenx86/   interp/   lexer/
    opt/   parser/   semant/   util/   visitor/

Main.java contains the main class file for running the compiler. The **lexer/**, **parser/**, **semant/**, **opt/** directories contain Java packages for performing lexical analysis, syntactic analysis, semantic analysis, and optimization, respectively. The **codegenmips/**, **codegenx86/**, and **codegenjvm** directories contain packages for performing code generation to the MIPS architecture, x86 (32-bit) architecture, and Java Virtual Machine (JVM), respectively. The **ast/**, **cfg/**, **util/**, and **visitor/** contain auxiliary class files required by the other packages. Finally, **interp/** contains a package for interpreting

Bantam Java programs. These packages are described in more detail in Section 4.

The **bantam/** and **bantam-jj/** directories share the ast, cfg, codegenjvm, codegenmips, code-genx86, interp, opt, semant, util, and visitor packages. The actual files are contained within **bantam/**. The **bantam-jj/** directory contains symbolic links to the shared directories in **bantam/**.

The **Makefile** within each of these directories will build the compiler. Each **Makefile** depends on auxiliary tools (in the **tools/** directory within the root directory) and so the tools must be built before building the source code. Like the root directory **Makefile**, to use the **Makefile** in either **bantam/** or **bantam-jj/** you will need to edit one line of the **Makefile**. The definition of the variable, **PROJPATH**, on line 2 will probably be incorrect. Set this variable to the path of the installation directory. To make the source, type "make" in either of the **src/** subdirectories: **src/bantam/** or **src/bantam-jj/**. Alternatively, you can type "make src" in the root directory. This will build the version of the compiler that is specified in the **VERSION** variable within the root directory **Makefile**. (Note: if compiled from the root directory the **PROJPATH** variable in the source makefiles does not have to be changed. This change is only necessary when compiling from the sub-directory.)

As stated above, the Java code for the compiler is incomplete. In particular, the lexer, parser, semantic analyzer, and code generator packages contain only skeleton code. However, within each of these packages are working class files that will allow students to test their phases of the compiler in isolation. These files are backed up in files ending with ".ref" so if they are over-written or removed for some reason they can be recovered. To build a compiler using the reference class files, you can type "make ref" within the **src/bantam/** or **src/bantam-jj/** directories.

tests. The tests directory contains several Bantam programs that can be used to test the Bantam compiler. It must be built after building the compiler. Bantam Java programs end with the suffix ".btm". The **tests/** directory contains several ".btm" files and a **Makefile**.

The source code for the compiler (in **src/**) must be built before building the test files. In addition, compiling the test programs will fail, of course, if the compiler is not fully implemented. Alternatively, the reference class files within the source code directory (see above) can be used to compile the test programs before the compiler is fully implemented. Once a working version of the compiler is built and the program file **bin/bantamc** is built, then the test files can be compiled. To compile the Bantam Java test programs, you can either type "make tests" in the root directory or "make" in the **tests/** directory. This will create several MIPS assembly files (the default target) with extension ".asm", which can run via the **SPIM** simulator. Note: MIPS assembly files use extension

".asm" and x86 assembly files use extension ".s" to avoid confusion.

The SPIM simulator [15] can be used to run any one of the compiled programs, *e.g.*, **hello-world.asm** (assuming we are in **tests**/):

    bash$ ../bin/spim hello-world.asm

The **Makefile** will also allow the user to compile for one of the other two targets: x86 or JVM. The target can be changed by editing the **TARG** variable in the **Makefile**. For example, the following will compile the tests for the x86 architecture (assuming we are in **tests**/):

    bash$ make TARG=x86

This will run the Bantam Java compiler on each source program producing an assembly file (ending in ".s". The **Makefile** then uses **gcc** to assemble these into x86 binaries. Each binary has extension ".bin", which can of course be run natively. For example, the following will run **hello-world.bin**:

    bash$ ./hello-world.bin

The following will compile for the JVM target:

    bash$ make TARG=jvm

This will eventually create a jar file with extension ".jar" for each source program. However, this build is done in several steps. For each source program, the **Makefile** first uses the Bantam Java compiler to generate Jasmin [18, 17] input files for each class defined in the program. It then uses Jasmin to assemble these into class files, which are finally bundled into a single jar file. The jar files can be run via the Java Virtual Machine. The following runs **hello-world.jar** using Sun's JVM:

    bash$ java -jar hello-world.jar

In addition, you may also want to set some of the other Bantam Java compiler arguments. These can set by editing the **FLAGS Makefile** variable (the flags are discussed in more detail below). For instance, the following enables garbage collection (which is disabled by default) and debugging during parsing:

    bash$ make FLAGS="-gc -dp"

Make automatically builds all of the test programs. One can also build a particular test program. The following will make only the x86 assembly file **hello-world.s** from the Bantam program **HelloWorld.btm**.

    bash$ make TARG=x86 hello-world

The base name for the output file (in the case above **hello-world**) can be found by looking at the **TESTS** variable inside the **Makefile**.

14

The Bantam toolset includes a reference compiler for testing purposes. After implementing all of the compiler phases, the student's compiler should work like the reference compiler. The reference compiler is built as part of the auxiliary tools (described above). After building the tools, the reference compiler is placed in **bin/** (either **bantamc-ref** or **bantamc-jj-ref**). To run the reference compiler we could use the following:

        bash$ ../bin/bantamc-ref HelloWorld.btm

We can also use the **Makefile** in **tests/** to compile the Bantam programs with the reference compiler by editing the **BTMC** variable in the **Makefile**. For example, the following will compile all of the Bantam programs with the reference compiler:

        bash$ make BTMC=../bin/bantamc-ref

2.2   Overview and Installation Instructions using Ant

The Bantam Java toolset will compile and should run on any Microsoft Windows or linux/x86 machine, although it is primarily tested on an Intel 32-bit x86 machine running Suse, Red Hat, and Ubuntu Linux. It may also run on other *nix platforms and other architectures. It is available as a ZIP file at the website: *http://www.bantamjava.com*. Note: there are two versions of the toolset: one for building a compiler for the base Bantam Java language and one for building a compiler for the extended language. Make sure that you have downloaded the correct toolset. The version of the toolset that you use should be the same as the version of this manual (you are reading the **extended** version).

Once you have downloaded the ZIP file, move it to the appropriate location. Then extract its contents using the following command. Note: we show commands for *nix. For other operating systems, you will need to use the corresponding commands.

    bash$ unzip bantamjava.zip

Note: in some operations systems, like Windows, you will just need to click on the ZIP file to unpack it. This will create a new directory, **bantamjava/** (the *root directory*), in the current location. The **bantamjava/** directory contains the following files and directories (note: directories end with '/', files do not):

    README   build.xml   api/   bin/   lib/   man/   src/   tests/   tools/

Each of these files and directories is described below (although not in order).

README. The **README** file provides installation documentation as described in this chapter. It

also indicates which version of the toolset you have downloaded (base or extended).

bin. The **bin/** directory contains executable files for running the compiler and other auxiliary tools. This directory is initially empty. Executable files are automatically added to **bin/** when the toolset is built.

lib. The **lib/** directory contains library files needed for building the compiler and also running auxiliary tools (*e.g.*, JLex). This directory is initially empty. Library files and directories are automatically added when the toolset is built.

man. The **man/** directory contains man pages (for use on *nix systems) for the Bantam compiler as well as some of the auxiliary tools (*e.g.*, Spim), although the man pages for the auxiliary tools are added only after the toolset is built.

build.xml. The **build.xml** ant script will compile and install all of the components of the toolset.

There are basically five parts in the toolset: the Javadoc web documentation describing the Java class files of the compiler (**api/**), the source files for the (incomplete) compiler (**src/**), the bantam test programs for debugging the compiler (**tests/**), and various tools used by the compiler (**tools/**). These components are described in more detail below.

Before you start using the ant script, you may want to edit the value of the first two properties in **build.xml**. The first property, called "VERSION," determines which version of the compiler will be installed: use the value "bantam" if you are using the JLex and Java Cup lexer and parser generators; use the value "bantam-jj" if you are using the JavaCC lexer and parser generator.

The second property, called "TARGET," determines which architecture your Bantam Java compiler will target. Use the value "mips" if your compiler's target is the Spim simulator; use the value "jvm" if your compiler's target is the Java Virtual Machine; otherwise, use the value "x86."

To build the entire toolset, type either "ant" or "ant all" in the root directory. You can also make an individual component as discussed below. To uninstall the toolset, type "ant cleanall" in the root directory or in any of the subdirectories that contain a **build.xml** file. Note: if the compiler is unimplemented (as it is when first installed) "ant" will result in an error initially since the last thing the **build.xml** script attempts to do is compile the test programs using the Bantam Java compiler. However, if installed correctly, no error should occur until the point where the script attempts to compile a test program.

api. The **api/** directory houses the Javadoc application programming interface documentation.

After uncompressing and untarring the toolset, the directory will contain only a **build.xml** ant script. The documentation is created from comments in the source code using the **javadoc** command provided with the Sun JDK [20]. It is placed in a directory **html/** within the **api/** directory. After generating the API (from the **build.xml** script), it can be viewed online by copying this directory to a website. For example, if **/var/www/** is a website directory viewable at *http://www.foo.com/*, then the following commands would make the API viewable at *http://www.foo.com/bantam-api/*.

```
bash$ ant api
bash$ cp -r api/html/ /var/www/bantam-api
```

When making changes to the compiler, Javadoc comments should be maintained so that any changes will be reflected in the API documentation. Sun Microsystems provides a free tutorial describing Javadoc commenting [20]. The API is also available online at *http://www.bantamjava.com/*.

tools. The **tools/** directory houses external software required by the Bantam Java toolset. These tools include JLex, a lexer generator [7]; Java Cup, a parser generator [14]; JavaCC, a lexer and parser generator [6]; and SPIM, a MIPS emulator [15]. These tools do not come with the toolset, but must be downloaded and moved to the **tools/** directory. Links to the respective websites are included online at *http://www.bantamjava.com*. These tools are built by typing "ant tools" in the root directory or "ant" in the **tools/** directory. This manual does not describe these tools (at least the auxiliary tools) in much detail. For more information on these tools, check out their respective manuals [7, 14, 6, 15].

The **tools/** directory initially contains the JLex, Java Cup, and JavaCC tools. It does not contain the SPIM simulator, which will need to be downloaded separately. The directory **tools/lib** holds two Jar files, which contain reference Bantam Java compilers for each version of the compiler: one that uses JLex and Java Cup (**bantamc-ref.jar**) and one that uses JavaCC (**bantamc-jj-ref.jar**). The **tools/lib/** directory also holds a runtime library file (containing assembly code) for two of the supported targets: MIPS/SPIM (exceptions.s) and x86/Linux (x86-runtime.s). In addition, it holds Java class files (**TextIO.class** and **Sys.class**), which serve as the runtime library when targeting the Java Virtual Machine (JVM). (Note: other built-in classes like **Object** and **String** as well as support for garbage collection need not be included in the JVM runtime library as they are already included in the JVM itself.) In addition to **lib/**, the **tools/** directory also contains a **build.xml** script that these tools and library files.

src. The source files for the compiler are contained within **src/**. The Bantam compiler is written

entirely in Java. Of course, the compiler is incomplete and will need to be completed by the student (although the code will compile). This directory contains two directories, **bantam/** and **bantam-jj/**. The **bantam/** directory contains the source files that use the JLex lexer generator [7] and the Java Cup parser generator [14]. The **bantam-jj/** directory contains the source files that use the JavaCC lexer and parser generator [6]. The **bantam-jj/** directory uses some symbolic links to the bantam directory to prevent duplication of code. Students should implement the compiler within one of these directories (but probably not both) as specified by the course instructor.

Within each directory (**bantam/** and **bantam-jj/**) are the following files and directories:

Makefile build.xml Main.java ast/ cfg/ codegenjvm/ codegenmips/ codegenx86/ interp/
lexer/ opt/ parser/ semant/ util/ visitor/

Main.java contains the main class file for running the compiler. The **lexer/**, **parser/**, **semant/**, and **opt/** directories contain Java packages for performing lexical analysis, syntactic analysis, semantic analysis, and optimization, respectively. The **codegenmips/**, **codegenx86/**, and **codegenjvm** directories contain packages for performing code generation to the MIPS architecture, x86 (32-bit) architecture, and Java Virtual Machine (JVM), respectively. The **ast/**, **cfg/**, **util/**, and **visitor/** directories contain auxiliary class files required by the other packages. Finally, **interp/** contains a package for interpreting Bantam Java programs. These packages are described in more detail in Section 4.

The **bantam/** and **bantam-jj/** directories share the ast, cfg, codegenjvm, codegenmips, codegenx86, interp, opt, semant, util, and visitor packages. The actual files are contained within **bantam/**. The **bantam-jj/** directory contains symbolic links to the shared directories in **bantam/**.

The **build.xml** script within each of these directories will build the compiler. Each ant script depends on auxiliary tools (in the **tools/** directory within the root directory) and so the tools must be built before building the source code. To compile the source, type "ant" in either of the **src/** subdirectories: **src/bantam/** or **src/bantam-jj/**. Alternatively, you can type "ant src" in the root directory. This will build the version of the compiler that is specified in the "VERSION" property within the root directory **build.xml** script.

As stated above, the Java code for the compiler is incomplete. In particular, the lexer, parser, semantic analyzer, and code generator packages contain only skeleton code. However, within each of these packages are working class files that will allow students to test their phases of the compiler in isolation. These files are backed up in files ending with ".ref" so if they are over-written or removed for some reason they can be recovered. To build a compiler using the reference class files, you can type "ant ref" within the **src/bantam/** or **src/bantam-jj/** directories.

tests. The **tests/** directory contains several Bantam programs that can be used to test the Bantam compiler. It must be built after building the compiler. Bantam Java programs end with the suffix ".btm". The **tests/** directory contains several ".btm" files and a **build.xml** script.

The source code for the compiler (in **src/**) must be built before building the test files. In addition, compiling the test programs will fail, of course, if the compiler is not fully implemented. Alternatively, the reference class files within the source code directory (see above) can be used to compile the test programs before the compiler is fully implemented. Once a working version of the compiler is built and the program file **bin/bantamc** (or **bin/bantamc.bat** on Windows machines) is built, then the test files can be compiled. To compile the Bantam Java test programs, you can either type "ant tests" in the root directory or "ant" in the **tests/** directory. This will create several MIPS assembly files (the default target) with extension ".asm", which can run via the **SPIM** simulator. Note: MIPS assembly files use extension ".asm" and x86 assembly files use extension ".s" to avoid confusion.

The SPIM simulator [15] can be used to run any one of the compiled programs, *e.g.*, **hello-world.asm** (assuming we are in **tests/**):

    bash$ ../bin/spim hello-world.asm

The **build.xml** script will also allow the user to compile for one of the other two targets: x86 or JVM. The target can be changed by editing the "TARGET" property in the script. This change can be performed once and for all in the script file itself. Alternatively, a different target may be specified on the command line. For example, the following will compile the tests for the x86 architecture (assuming we are in **tests/**):

    bash$ ant -DTARGET=x86

(Note the "-D" in front of the "TARGET" property name). This will run the Bantam Java compiler on each source program producing an assembly file (ending in ".s". The **build.xml** script then uses **gcc** to assemble these into x86 binaries. Each binary has extension ".bin," which can of course be run natively. For example, the following will run **hello-world.bin**:

    bash$ ./hello-world.bin

The following will compile for the JVM target:

    bash$ ant -DTARGET=jvm

This will eventually create a jar file with extension ".jar" for each source program. However, this build is done in several steps. For each source program, the ant script first uses the Bantam Java compiler to generate Jasmin [18, 17] input files for each class defined in the program. It then

19

uses Jasmin to assemble these into class files, which are finally bundled into a single jar file. The jar files can be run via the Java Virtual Machine. The following runs **hello-world.jar** using Sun's JVM:

    bash$ java -jar hello-world.jar

In addition, you may also want to set some of the other Bantam Java compiler arguments. These can be set by editing the "FLAGS" property in the **build.xml** script (the flags are discussed in more detail below). For instance, the following enables garbage collection (which is disabled by default) and debugging during parsing:

    bash$ ant -DFLAGS="-gc -dp"

(Again, note the "-D" in front of the "FLAGS" property name). The ant script automatically builds all of the test programs. One can also build a particular test program. The following will make only the x86 assembly file **hello-world.s** from the Bantam program **HelloWorld.btm**.

    bash$ ant hello-world -DTARGET=x86

The base name for the target file (in the case above, **hello-world**) can be found by looking at the **compile-all** target inside the **build.xml** file.

The Bantam Java toolset includes a reference compiler for testing purposes. After implementing all of the compiler phases, the student's compiler should work like the reference compiler. The reference compiler is built as part of the auxiliary tools (described above). After building the tools, the reference compiler is placed in **bin/** (either **bantamc-ref** or **bantamc-jj-ref**). To run the reference compiler we could use the following:

    bash$ ../bin/bantamc-ref HelloWorld.btm

We can also use the **build.xml** script in **tests/** to compile the Bantam Java programs with the reference compiler by editing the **BTMC** variable in the build script. For example, the following will compile all of the Bantam programs with the reference compiler:

    bash$ ant -DBTMC=../bin/bantamc-ref

2.3   Using the Bantam Compiler

This subsection describes the command-line usage of the Bantam compiler. This material is also described in the Bantam Java compiler man page.

Although parts of the compiler are unimplemented (and will be implemented by the student), the main class of the compiler is implemented, which includes command-line, flag handling. So the

unimplemented compiler will accept all options, although some may have no effect (*e.g.*, enabling debugging for the code generator).

The Bantam Java compiler accepts the following options:

bantamc [-o <output_file>] [-t <architecture>] [-gc] [-int] [-opt <level>]
[-dl] [-dp] [-ds] [-di] [-do] [-dc] [-sl] [-sp] [-ss] [-so] <input_files>

Options. Each of the compiler options are enumerated and discussed below.

-o output_file

Specify the output file to use; the default is out.s. Note: with the JVM target the output file is ignored (instead an output is produced for each user-defined class).

-t architecture

Specify the target architecture to generate code for; mips, x86, and jvm (must be all lower-case) are the currently supported target architectures. mips is the default.

-gc

Enable the garbage collector. The Bantam compiler uses a simple, mark-and-sweep garbage collector. When the JVM is the target, garbage collection cannot be toggled (*i.e.*, it is always enabled) and thus this flag is ignored. By default the garbage collector is disabled.

-int

Enable interpretation of the Bantam Java code. This flag allows the compiler to act as an interpreter; it will execute the program rather than compiling it. Note: optimization is not supported while in interpreted so the "-opt", "-so", and "-do" flags are ignored in interpreter mode. Code generation specific flags are also ignored such as "-t", "-dc", and "-gc". By default interpreter mode is disabled.

-opt level

Set the optimization level. The level can be set to an integer between 0 and 4. The higher the level, the more aggressive the optimization. A level of 0 means optimization is disabled. By default the level is 0.

-dl

Enable lexical analysis debugging. Note: students must add the debugging code that uses this flag. By default this flag is off.

-dp

Enable syntactic analysis debugging. Note: students must add the debugging code that uses this flag (unless using the Java Cup implementation, in which case, debugging code is automatically inserted by the parser generator). By default this flag is off.

-ds

Enable semantic analysis debugging. Note: students must add the debugging code that uses this flag. By default this flag is off.

-di

Enable interpreter debugging. This flag is ignored unless interpretation has been enabled via the **-int** flag. Note: students must add the debugging code that uses this flag. By default this flag is off.

-do

Enable optimization debugging. This flag is ignored unless the optimization level has been set to a value greater than 0. Note: students must add the debugging code that uses this flag. By default this flag is off.

-dc

Enable code generation debugging. Note: students must add the debugging code that uses this flag. By default this flag is off.

-sl

Stop the compiler after lexing. If this flag is set, the compiler halts after performing lexical analysis. If lexical errors are discovered, these are printed to standard error, otherwise the scanned tokens are printed to standard output. By default this flag is off.

-sp

Stop the compiler after parsing. If this flag is set, the compiler halts after performing syntactic analysis. If lexical or syntactic errors are discovered these are printed to standard error, otherwise the parsed program is printed to standard output (as a Bantam source program). If -sl is specified, then this flag is ignored. By default this flag is off.

-ss

Stop the compiler after semantic analysis. If this flag is set, the compiler halts after perform-

ing semantic analysis. If errors are discovered these are printed to standard error, otherwise the annotated program (annotations indicate type information) is printed to standard output (as a Bantam source program with annotations in comments). If -sl or -sp is specified, then this flag is ignored. By default this flag is off.

-so

Stop the compiler after optimization. If this flag is set, the compiler halts after performing optimization and prints the optimized program (as a set of control flow graphs) to standard output. If -sl, -sp, or -ss is specified, then this flag is ignored. By default this flag is off.

Input files. The input files to the compiler must be Bantam files. These files should contain Bantam code (as discussed in Section 3) and end with the suffix ".btm". If an input file is specified that does not end with ".btm" then the compiler will immediately print an error and exit. If the input files contain errors, then the compiler will not produce an output file, but instead print error messages and exit (unless the student has implemented a buggy compiler).

Output file. The output file is either a MIPS assembly file or x86 assembly file depending on the particular target machine. Although not required, by convention this file should end with ".s". By default, the compiler uses the file "out.s".

The output file is ignored when targeting the JVM. In this case, one Jasmin [18, 17] input file is created for each Bantam Java class in the program. These files have a ".j" extension. For example, if the program includes the classes **Main** and **Foo** then the compiler will create Jasmin assembly files **Main.j** and **Foo.j**. Jasmin can then be used to translate these into Java class files.

Examples. Here are some example uses of the Bantam compiler. These examples assume we are in the **tests**/ directory.

```
bash$ ../bin/bantamc -o foo.s Foo.btm
bash$ ../bin/spim foo.s
```

The first command above will compile the Bantam program **Foo.btm** and create a MIPS assembly file (to be used with the Spim emulator [15]) called **foo.s**. The second command above runs the program using the Spim emulator for MIPS.

```
bash$ ../bin/bantamc -t x86 -o foo.s Foo.btm
bash$ gcc -o foo foo.s
bash$ ./foo
```

The commands above compile and run Foo.btm on an x86, 32-bit machine. The first command compiles **Foo.btm** to **foo.s**, specifying the target as x86. The second command uses the gcc assembler [1] to translate the assembly file to an executable called **foo**. Finally, the third command runs the program. The x86 target uses the AT&T x86 assembly format, and must be run with an AT&T assembler, such as gas, the GNU assembler that is bundled with gcc. Currently, the x86 target can only be used on 32-bit x86 machines. In the future, we will add a target for 64-bit x86 machines.

```
bash$ ../bin/bantamc -gc -t x86 -o foo.s Foo.btm
bash$ gcc -o foo foo.s
bash$ ./foo
```

The commands above are identical to the previous set of commands (they compile and run **Foo.btm** on an x86 machine) except in this case the garbage collector is used.

```
bash$ ../bin/bantamc -t jvm Foo.btm
bash$ jasmin Main.j
bash$ jasmin Foo.j
bash$ java -cp ../../lib:. Main
```

The commands above compile and run Foo.btm on the Java Virtual Machine (assume here that **Foo.btm** contains two classes: **Main** and **Foo**). The first command compiles **Foo.btm** to the Jasmin input files **Main.j** and **Foo.j**. The next two commands convert these text-based bytecode files into actual Java class files: **Main.class** and **Foo.class**. Finally, the last command runs the program by running the **Main** class. Note: it also includes the runtime library class files **TextIO** and **Sys** from **../../lib** in the classpath. One could also create a JAR file and run that as well as follows:

```
bash$ cp ../../lib/*.class .
bash$ jar cvfm foo.jar manifest Main.class Foo.class TextIO.class Sys.class
bash$ java -jar foo.jar
```

In this case, the class files from the program as well as the built-in class files **TextIO** and **Sys** are bundled together (the other built-in class files are already a part of the JVM). The file **manifest** is a simple JAR manifest file specifying that the main class is **Main**.

The "-opt" flag can be used to optimize the compiled program:

```
bash$ ../bin/bantamc -opt 4 Foo.btm
```

The command above turns on optimization and sets the optimization level to 4 (the most aggressive optimization).

```
bash$ ../bin/bantamc -sl Foo.btm
```

The command above performs only lexical analysis on **Foo.btm** and then prints out the scanned

24

tokens to standard output, unless errors are discovered, in which case the errors are printed to standard error.

    bash$ ../bin/bantamc -dp -sp Foo.btm

The command above performs lexical and syntactical analysis on **Foo.btm** and then prints out the parsed program to standard output, unless errors are discovered, in which case the errors are printed to standard error. In addition, debugging is enabled during parsing.

    bash$ ../bin/bantamc -int Foo.btm

The command above interprets rather than compiles the source program **Foo.btm**.

2.4   Last Words

You have now seen how to install and use the Bantam Java compiler. The upcoming sections will look in more depth at both the Bantam Java language and the Bantam Java compiler, starting first with the language.

3   Bantam Java Language

This section presents the Bantam Java programming language, the source language for the compiler. Bantam Java is a subset of the Java language. This section discusses those Java features that were included and excluded from the Bantam language. This section also describes in some detail the included features, although for a more detailed explanation one should consult a Java reference [5]. This section and this manual are intended for students who are already familiar with the Java programming language and its features. Note that you are reading the **extended** language version of this manual.

3.1   Overview

Bantam Java, like Java, is an object-oriented programming language. As described in numerous textbooks [12, 5] (too many to cite all of them), an object is a self-sufficient component that has some data and behavior associated with it, which is meant for modeling a real world entity such as a player, animal, room, *etc.* To construct an object in Bantam or Java, one defines a class, which is effectively the blue print for all objects generated from that class. In other words, an object is an *instance* of a class. The class defines both the data elements and behavior of any object constructed from that class.

A class in Bantam Java and Java is essentially a *type*. A type specifies how a value (in this case an object) of that type can be used. Bantam Java, like Java, is a *strongly-typed language*, meaning that the compiler, either during compilation or at runtime, prevents a value of one type being used where a value from a different type is required (there is an exception to this rule, which will be discussed in the next subsection). In other words, the compiler effectively prevents values from being used in an improper way.

A Bantam program consists of one or more files (with suffix ".btm"), each of which contains one or more class definitions. All code must reside within a class definition in Bantam. For this reason, classes and objects are critical and necessary components of any Bantam program.

Each class in Bantam, contains zero or more *members*. A member is either a *field* (part of the object's data) or a *method* (part of the object's behavior). A field is one kind of *variable*, which can hold various data values throughout the course of the owner object's lifetime. A field has a type, which specifies the type of data it can store. A method, on the other hand, is effectively a

26

name for a block of code, which can be executed on behalf of the owner object. The method can be executed by referring to its name, which is known as a method *call*. We will show later how fields and methods can be declared and used.

Unlike in Java, all fields and methods must be instance members, that is to say they cannot be declared **static**. In other words, each member is specific to a particular object; there are no class-wide members in Bantam.

When an object is created, the syntax and semantics of which are discussed later in the section, it is allocated its own copy of each field. The object's field values may differ from the field values of other objects created from the same class. When an object's method is called, it may use the object's fields, therefore the outcome of a method depends largely on the owner object. Before an object is created, it is set to the special value **null**, which loosely speaking means it is an empty object.

One of the classes defined in a Bantam program must be called **Main**. This class must contain a method called **main** (among possibly other members). When a Bantam program is started, a **Main** object is created and the **main** method is called. The program terminates when the **main** method finishes. The **main** method does all of the work of the program or calls methods that do the work on its behalf. Note that this starting point is different than in Java. In Java, which has static members, a static **main** method within a user-specified class is initially called.

There are also some built-in classes that allow the programmer to perform input and output to the screen or to a file (**TextIO**), exit the program (**Sys**), and store and manipulate sequences of characters (**String**). These are discussed in Section 3.7.

In addition to classes, there are two primitive types of data in Bantam: **int** and **boolean**. A primitive, unlike an object, holds a simple value. An **int** holds a 32-bit signed integer. A **boolean** holds either **true** or **false**. Other primitives from Java, such as **double** and **long** were not included in Bantam. These additional types add little intellectual challenge to the design and implementation of the compiler and for this reason were excluded from the language. As in Java, primitive types can be used in similar places as object types (*e.g.*, a field can have type **int**).

One important distinction to make is between primitive variables and object variables. A primitive variable (*e.g.*, a field) holds the value (*e.g.*, 42). However, an object variable does not hold an object itself, but rather a *reference*, or more precisely a memory pointer, to the object. If two or more object variables refer to the same object, then we say that they *alias*.

Bantam Java (the extended language) also includes arrays, which are a data type for storing

27

groups of indexed elements. Each element in the array is accessed via its index (an **int**). The first element is accessed at index 0, the second element is accessed at element 1, and so on. An array can hold any type of elements, besides arrays, however, the type of each element, which is called the base type, must be the same. The base type can be any primitive or class. Unlike in Java, the base type cannot, itself, be an array type. In other words, Bantam Java does not support multi-dimensional arrays.

As in Java, arrays in Bantam Java are implemented as objects. Array types are not explicitly declared by the programmer. Instead if a programmer declares or creates an array within the program, then that array type (assuming it is legal) automatically exists. Each array type regardless of base type contains two members: a **length** field that holds the length of the array and a **clone** method for cloning the array. To distinguish between array types and other types, array types are written as <**base type**>**[]** where <**base type**> is replaced with the base type (*e.g.*, **int**). Some example array types include **int[]** and **String[]**. Because an array is an object, an array can be set to **null**.

Memory management in Bantam Java is similar to Java. Memory allocation occurs when objects (including arrays) are constructed (using **new**). Memory deallocation is handled implicitly, *i.e.*, a garbage collector is responsible for deallocating memory. Unlike in Java, the garbage collector is disabled by default and can be enabled using a compiler flag (see Section 2).

There are two types of comments in Bantam Java. The first is a single line comment, which begins with **//**. The text following **//** up to and including the end of the current line is ignored by the compiler (including **//**). The second type of comment is a multi-line comment that begins with **/*** and ends with ***/**. As the name suggests, these may span multiple lines. The compiler ignores all text in between **/*** and ***/** (including **/*** and ***/**). Multi-line comments may not be nested.

In the remainder of the section, we will look in more detail at defining classes and class members. But before doing that, we discuss a core feature of Bantam and most object-oriented languages, namely *inheritance*.

3.2 Inheritance

As in Java, Bantam allows classes to extend other classes. If class **A** extends class **B** then class **A** inherits all of class **B**'s members including all the members that **B** inherits from other classes. We say that **A** is a child of **B** and **B** is the parent of **A**. If a class does not explicitly extend another

Figure 1: Example class hierarchy tree in Bantam.

class then it automatically extends the built-in class **Object**, which is similar to the **Object** class from Java (albeit with less functionality). Bantam and Java support only single inheritance: a class inherits from exactly one other class except the built-in **Object** class, which does not extend any other class.

To visualize the class dependencies, we can build a class hierarchy graph where each node in the graph represents a class. A directed edge is drawn from node **A** to node **B** if and only if class **B** is the parent of class **A**. Because Bantam only supports single inheritance, the class hierarchy graph of a well-defined Bantam program is actually a tree with the class **Object** as the root.

Figure 1 shows an example class hierarchy tree in Bantam. Notice in the figure that the class **Tiger** is a descendant of the class **Object** (it extends **Animal**, which extends **Object**). We say that **Tiger** is a *subclass* of **Object** and that **Object** is a *superclass* of **Tiger**. More precisely, a class **A** is a subclass of a class **B** (or equivalently, **B** is a superclass of **A**) if we must pass through node **B** to get to the root of the class hierarchy tree from node **A**. Because a class is actually a type, we also say that **Tiger** is a *subtype* of **Object** and that **Object** is a *supertype* of **Tiger**. Each class is defined to be a subclass (or subtype) and a superclass (or supertype) of itself. The class **Object** is a superclass of all classes and is only a subclass of itself.

Because arrays are objects, they also inherit from other types. Every array type is a subtype of the generic type **Object**. In fact, the array method **clone** is inherited from **Object** (the field **length** is not). In addition, if the base type of an array type is a subtype of the base type of another array type, then the first array type is a subtype of the second array type.

The class hierarchy tree must be modified slightly to accomodate arrays. In the updated graph, a node represents a class or an array type, assuming the array type is used in the program. An edge exists from each child node to its parent node. Figure 2 shows an example class hierarchy tree, assuming that the array types **int[]**, **String[]**, **Animal[]**, and **Tiger[]** are used in the program. Because
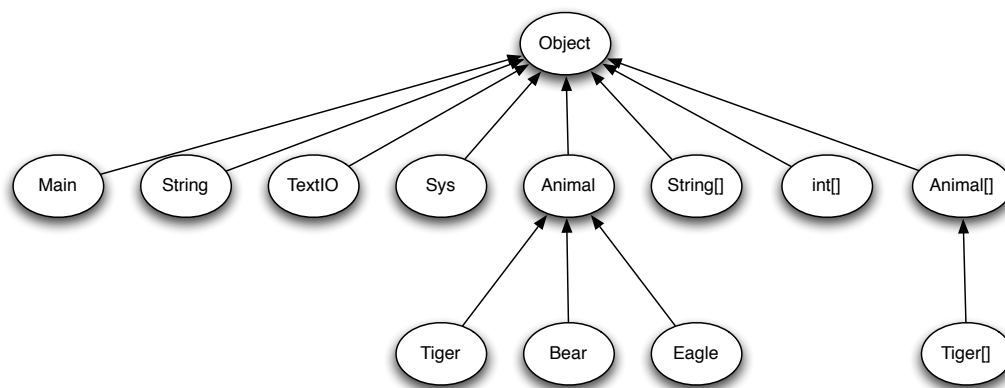
Figure 2: Example class hierarchy tree in Bantam, which includes array types.

Animal is a supertype of **Tiger**, **Animal[]** is a supertype of **Tiger[]**. Array types with a primitive base type (*e.g.*, **int[]**) are children of **Object**. The array types **String[]** and **Animal[]** are also children of **Object**. If the program had used the array type **Object[]** then **String[]** and **Animal[]** would have been children of **Object[]**. But because the program did not use **Object[]**, **String[]** and **Animal[]** are instead children of **Object**.

When a class **A** is a subclass of a class **B**, we also say that **A** *conforms* to **B** since **A** must contain all of the functionality of **B** (likewise for array types that are subtypes of other types). Therefore, anywhere that we can use an object of type **B**, we can also use an object of type **A**. The opposite is not true. Because **A** may have more functionality than **B**, an object of type **B** cannot be used in place of an object of type **A**.

One implication of conformance is that the type of an object variable at compile time may differ from the type at runtime. For example, we can construct a new **Tiger** object and store it in a variable of type **Object** since **Tiger** conforms to **Object**. The compile time type of the variable, which is often called the *static type*, is **Object**, while the runtime type, which is often called the *dynamic type*, is **Tiger**. The dynamic type is often unknown until the program is run.

We will have more to say about inheritance and conformance in our discussion of class members (3.4) and casting (3.6).

3.3 Class Definitions

The format of a class definition is:

```
class <name> [ extends <parent> ] {
    <members>
}
```

where <**name**> is replaced with the name of the class, <**parent**> is replaced with the name of the parent class, and <**members**> is replaced with the class members (fields and/or methods). Note: [and] are not part of the language, but are instead used to indicate that **extends** <**parent**> is optional. (Note: '[' and ']' are a part of the language.) A class can explicitly extend another class by using **extends**. Alternatively, **extends** can be omitted, in which case the class automatically extends the built-in class **Object**.

Both **class** and **extends** are keywords in Bantam. Neither can be used as a name of a class, method, or variable. Bantam is case-sensitive, so **class** and **extends** are keywords, while **Class** and **EXTENDS** are not. The class name (<**name**>) and parent name (<**parent**>) are also case sensitive (in fact, all tokens in Bantam are case sensitive). Therefore, a class named **Animal** is not the same as a class named **aNimal**. The names may contain any sequence of upper or lowercase letters, digits, or the symbol '_', although the name must start with a letter. By convention (taken from Java), class names generally start with an uppercase letter. If a class name contains multiple words combined into one name the start of each new word begins with an uppercase letter. For example, if we wanted a class name to include both "Bantam" and "Java" we would name the class **BantamJava**.

Unlike in Java, class definitions cannot use modifiers such as **public** and **package**. Since there are no packages in Bantam these modifiers are not necessary and using one is illegal. In addition, Bantam does not include abstract and static classes so the **abstract** and **static** modifiers are also illegal. Bantam also does not include interfaces so a class cannot use the keyword **implements** nor can an interface be defined. Finally, classes cannot be nested within other classes.

To define a (trivial) class **Animal**, we write the following:

```
class Animal {
}
```

This class implicitly extends **Object**. It defines no data elements and no behavior. To do this, we must define some class members.

3.4 Member Definitions

As stated above there are two kinds of class members: fields and methods. We look in turn at defining each, below.

Fields. The format for defining a field (*i.e.*, instance variable) is:

$$<\text{type}> <\text{name}> \Big[= <\text{expression}> \Big] \;;$$

where <**type**> is replaced with the field's type, <**name**> is replaced with the field's name, and <**expression**> is replaced with the field's initialization expression. Note the '=' and initialization expression are optional. The <**type**> must be either a primitive type (*i.e.*, **int** or **boolean**), a class name (*e.g.*, **Object**), or an array type (*e.g.*, **Object[]**). The <**name**> is a non-empty sequence of upper and lowercase letters, digits, and the special symbol '_'. Each field defined within a particular class must have a distinct name.

By convention, field names generally start with a lowercase letter. Like class names, when multiple words are combined into one variable name, the start of each word (except the first word) begins with an uppercase letter. For example, if we wanted a field name to contain "Bantam" and "Java" we would call it **bantamJava**. The beginning lowercase letter indicates it is not a class name and the subsequent uppercase letters indicate the separation of words embedded within the name. If the field is meant to be constant (*i.e.*, never changed) then the convention is to use all uppercase and use '_' to indicate separation of words (*e.g.*, **BANTAM_JAVA**).

Below is another (slightly more interesting) definition of class **Animal**, which includes three field definitions:

```
class Animal {
    int numLegs = 4;
    boolean canFly;
    Animal mate;
}
```

The first field, **numLegs**, is a primitive variable with type **int**, which initially holds the value 4, although this can be changed later. The second field, **canFly**, is a also a primitive variable, but with type **boolean**. It defines no initialization expression. When a field definition omits the initialization expression, it is automatically assigned a value based on its type. Fields with type

int are automatically assigned 0, fields with type **boolean** are automatically assigned **false**, and fields with object types (including array types) are automatically assigned **null**. Therefore, **canFly** is initially set to **false**. The third field is an object of type **Animal**. Its initial value is **null**, since it does not define an initialization expression. It has the same type as the class that is being defined, which is okay so long as we do not attempt to initially construct the object (object construction is discussed later in Subsection 3.6). If we try to construct a new **Animal**, we will end up in an endless cycle of **Animal** object constructions. This problem only occurs for objects with types that are a subclass of the defined class.

We could have also written the following equivalent code:

```
class Animal {
    int numLegs = 4;
    boolean canFly = false;
    Animal mate = null;
}
```

It would not be equivalent to omit the assignment of **numLegs** to the value **4** since **numLegs** would have been automatically assigned **0**.

The initialization expression can be arbitrarily complicated as discussed in 3.6. For example, we could have the following:

```
class Animal {
    int numLegs = 4;
    boolean canFly;
    Animal mate;
    int strength = (numLegs * 10) + 50;
}
```

The (contrived) initialization value for the fourth field, **strength**, is the initialization value of **numLegs** multiplied by 10, and added to 50 (*i.e.*, 90). Note that the initialization expression uses another field defined within the class. This is only legal if the referenced field is defined above the current field, since field initialization expressions are evaluated in the order in which they are defined in the class. Therefore the ordering of fields within the class matters. Also, the initialization expression is only evaluated once when the object is first created and used to assign a starting value to the field. If **numLegs** is changed after the object is created, the initialization expression for **strength** is not re-evaluated.

33

We could also define an array field variable. For example, some animals have multiple mates throughout the course of their life, so we could use an array rather than a single object variable to hold the list of mates:

```
class Animal {
    int numLegs = 4;
    boolean canFly;
    Animal[] mates; // field array
    int strength = (numLegs * 10) + 50;
}
```

As the variable **mates** is an object (because arrays are objects) it is initially set to **null** since there is no initialization expression.

If a class extends another class, it can also reference any inherited fields:

```
class Tiger extends Animal {
    int tigerStrength = strength * 2;
}
```

Inherited fields are always evaluated before non-inherited fields, with the fields from the higher inherited classes (*i.e.*, closer to the root of the class hierarchy tree) evaluated earlier than fields from lower inherited classes (*i.e.*, farther from the root of the class hierarchy tree). When a **Tiger** object is created, the fields of **Object** are evaluated first, followed by the fields of **Animal**, and finally followed by the fields of **Tiger**.

Inherited fields in an extended class can be redefined. For example, we can do the following:

```
class Tiger extends Animal {
    int tigerStrength = strength * 2;
    Tiger mate;
}
```

The new field, **mate**, is treated as a separate instance variable, which essentially hides the inherited field within class **Tiger**. Notice that **mate** does not have to have the same type as the hidden inherited variable (or actually even a conforming type).

Every class implicitly has two special fields. The first variable, **this**, allows the programmer to refer to the current owner object. **this** can be used to pass the owner object to other methods. It can also be used to access members. The second, **super**, is similar to **this**, however, it only allows

access to inherited members, not members defined within the class. In particular, **super** allows programmers to access hidden members. Neither **this** or **super** can be altered (it is illegal to assign to **this** or **super**), they are implicitly set when an object is created. In addition, **super** can not be used as a stand-alone variable. It can only be used to reference hidden inherited members.

Unlike in Java, Bantam fields have no modifiers (*e.g.*, **public**, **static**). All fields are instance variables and are **protected**, and therefore field modifiers are not necessary and are actually illegal. Fields can only be accessed from within the class or a subclass of the class. In Java, **protected** fields can be accessed from classes within the same package. Since Bantam does not include packages, **protected** fields cannot be directly accessed by any class that is not a subclass of the class where the field is defined. Accessing fields from other classes can only be done via methods (which are **public**). There is also no **final** modifier for indicating to the compiler that a field is constant. It is the responsibility of the programmer to guarantee that a constant field is never changed.

The use of protected variables is more restricted in Bantam Java than in Java. In Java, a protected variable can be accessed via an object variable (*e.g.*, a field variable) so long as that variable's type is the same or a super type of the class containing the object reference. In order to simplify the implementation of the compiler, Bantam Java disallows this usage. Therefore, the following is illegal in Bantam Java:

```
class Tiger extends Animal {
    int tigerStrength = strength * 2;
    Tiger mate = ... // initialization not shown
    // LINE BELOW IS ILLEGAL IN BANTAM JAVA
    int mateStrength = mate.tigerStrength;
}
```

In the code above, accessing **tigerStrength** using the object reference **mate** is illegal. In Bantam Java, only **this** and **super** can be used as an object reference in accessing a protected field. As discussed below, there are no restrictions on performing method calls as methods are public. In the code above, we could create a public method for accessing **tigerStrength** and use this method in the assignment to **mateStrength**.

There is one exception to the rules outlined above. The field **length** within an array is **public** and **final**, meaning it is accessible from outside the class but it cannot be modified. If we have an array variable, called **array**, then we can access its length via **array.length**. It is illegal to assign **array.length** to a value.

For reasons having to do with garbage collection, the maximum number of combined inherited and defined fields cannot exceed 1500. This restriction is not a part of the language (and could change in the future, although the maximum would only be made larger not smaller), but rather a part of the particular implementation of the language.

Methods. The format of a method definition is:

```
<type> <name> ( <parameters> ) {
    <statements>
}
```

where <**type**> is the return type of the method (*e.g.*, **int**, **Object[]**), <**name**> is the name of the method, <**parameters**> is a list of comma-separated formal parameters to the method (formal parameters are defined below), and <**statements**> is a list of statements (possibly empty) terminated by semi-colons, which will be executed when the method is called. Like fields, each method defined within a single class must have a unique name (as discussed below, method overloading is not allowed). Method names follow the same conventions as field names.

Below is an example of method **getStrength** in class **Animal**:

```
class Animal {
    <field definitions (includes strength)>

    int getStrength() {
        return strength;
    }
}
```

The method **getStrength** takes no formal parameters and returns the value of the field **strength** (since the expression in the return statement is **strength**). This method is not contrived since fields are implicitly **protected** and can only be accessed from other classes via a **public** method. The method contains a single return statement. A return statement immediately returns from the method, perhaps returning a value (return statements will be discussed in more detail in Subsection 3.5). In this case, the return statement returns the value stored in **strength**. Because **strength** is an **int**, the return type of **getStrength** is **int**.

Methods can also return nothing by declaring the return type as **void**. In this case, the return statement should not return an expression.

Here is another method in **Animal**, which is **void**:

```
class Animal {
    <member definitions>

    void fight(int amount, boolean isWinner) {
        if (isWinner)
            strength = strength + amount * 5;
        else
            strength = 0;
    }
}
```

Notice that the method **fight** has two formal parameters: **amount** and **isWinner**. A *formal parameter* is a variable declaration (like a field declaration without the initialization expression). The first formal parameter in **fight**, **amount**, is an **int** variable, indicating the amount to fight. The second formal parameter, **isWinner** is a **boolean** variable indicating whether the animal is the winner of the fight. During a call to the method, each formal parameter is assigned a value called an *actual parameter* passed from the call site. Because each call site can specify **amount** and **isWinner**, **fight** can be used to fight any numeric amount and it can be used both when the animal wins or loses the fight. Parameters effectively allow programmers to write general-purpose methods, which can be called in different contexts. In Section 3.6 we will look at calling a method.

The method **fight** uses the two parameters to update the **strength** field in the animal object. The statement list in **fight** contains an if statement, which will be discussed in Subsection 3.5, but for now it is enough to know that the statement increments the **strength** by 5 times the parameter **amount** if the animal is the winner, otherwise, it sets **strength** to 0. No value is returned so the method must be declared **void**.

For **void** methods, no return statement is necessary, although one can be used if the programmer wishes, so long as it does not return an expression (*i.e.*, "return;"). For non-void methods, the last statement in the method must be a return statement, which cannot be embedded within another statement (*e.g.*, an if statement). This requirement does not exist in Java, but without it, the compiler would have to prove that there is no path through the method that does not end with a return statement. This check requires a sophisticated analysis. By requiring a return statement at the end of the method (which must be executed on any path through the method), this analysis is not necessary.

As in Java, return statements can be used at earlier points within the method besides the last

37

statement. However, unlike in Java, statements are allowed to follow a return statement even though they will not be executed. In Java, a statement must be *reachable*, meaning control must reach the statement via at least one path through the method. But to find unreachable statements requires a sophisticated analysis like the one needed to prove a return statement is executed on every path through the method. Again, to simplify the design of the compiler, statements in Bantam Java do not have to be reachable and can follow return statements.

In Java, the *signature* of a method is the method name combined with the ordered list of parameter types. In the definitions above, the signature of **getStrength** is **getStrength()** and the signature of **fight** is **fight(int,boolean)**. Methods within the same class, in Java, need only have a unique signature and not necessarily a unique name. This feature is called *method overloading*. In Bantam, method overloading is illegal and each method must have a unique name. However, fields and methods can share names. It is up to the compiler to determine, which is being referenced (the field or method) based on the context in which the name is used.

In the definitions of **getStrength** and **fight**, the type in the return statement matches the declared return type (**int** and **void**, respectively). However, that is only necessary when the declared return type is primitive or **void**. If the declared return type is instead an object type (*i.e.*, a class name or an array type), then the type of the return expression only has to conform to the declared return type.

```
class Animal {
    <member definitions>

    Animal mate;

    Object getMate () {
        return mate;
    }
}
```

The variable **mate** was defined with type **Animal**, however, the method **getMate** is declared to return type **Object**. This method definition is legal since **Animal** is a subclass of **Object** (although in this case we would probably want to specify **Animal** as the return type).

We can also use array types as a return type or a parameter type in a method. For example, assume we use an array to hold a list of animal mates. Then **getMates** would look as follows:

```
class Animal {
    <member definitions>

    Animal[] mates;

    Animal[] getMates() {
        return mates;
    }
}
```

The method **getMates** returns the **mates** array. As with non-array object types, the type returned in the return statement only has to conform to the declared return type. Therefore, we could also have declared **getMates** with return type **Object** or **Object[]**.

There is one required method in any Bantam Java program: the **main** method within the **Main** class. This method is the starting point of a Bantam Java program. The **main** method must be declared void and define no parameters. Omitting the **main** method or defining **main** in any other way is an error.

Unlike fields, there are no restrictions on the number of inherited and defined methods within a particular class. Like fields, methods do not have modifiers. All methods are **public**, meaning they can be called from within any class (given a dispatch object as discussed in 3.6). Like fields, methods can be redefined, however, a redefined method must have the same signature and return type as the inherited method.

Let's now take a closer look at the statements that we can use within a method.

3.5 Statements

There are eight types of statements in Bantam Java. We discuss each below.

Block statement. The first statement, a block statement, allows us to insert multiple statements where exactly one is needed. The format of a block statement is:

$$\{ \ <statements> \ \}$$

The block statement begins and ends with '{' and '}', respectively. Inside the braces are a list of statements. These statements can include potentially another block statement, *i.e.*, a block statement nested within another block statement. The utility of the block statement will become clear when we look at the other types of statements.

Declaration statement. A declaration statement declares a new variable and assigns it an initial value. This new variable is called a *local* variable since it may be accessed only within the method (there are other rules for accessing the variable, which are discussed below). The format of a declaration statement is:

<center><type> <name> = <expression> ;</center>

where <**type**> and <**name**> are the type and name of the declared variable, respectively, and <**expression**> is the initialization expression of the variable. Declaration statements must be terminated with a semi-colon. The rules and conventions for declaration statements are similar to fields except that the initialization expression is not optional. The initialization expression is not required in Java, however, Java requires that all local variables are assigned before they are used. Determining whether the variable is assigned before its first use requires some advanced compiler techniques, which we wanted to avoid in the Bantam compiler. We get around this by requiring all locally declared variables to be initially assigned.

Below is an example of a declaration statement:

```
class Main {
    void main () {
        int x = 3;
        ...
        return;
    }
}
```

In this case, we have declared a new variable **x**, which is local to the method **main**. The variable **x** cannot be used outside of the method **main**. In fact, it can only be used from the point it was defined until the end of the innermost enclosing block statement or the end of the method, whichever comes first. This range is called the *scope* of **x**. In this way, methods and block statements define new scoping levels. As in Java, a local variable can be defined using the same name as a field declared in the same class or an inherited class. While in the local variable's scope, the field is hidden, although it can still be accessed via **this**. Two local variables defined within the same method can also use the same name as long as they do not have overlapping scopes. It is a compiler error to declare two local variables with overlapping scopes using the same name.

The code below shows an example in Bantam Java of some variable declarations. Note: the

<center>40</center>

subscripts are not a part of the program, but instead, used to help the reader differentiate between block statements and variables with the same name.

```
class Main {
    int x_0 = 0;
    void main () {
        int a = x_0;
        {b1
            int b = x_0;
            int x_1 = 1;
            int c = x_1;
            {b2
                // this would be illegal:
                // int x_2 = 2;
                int d = x_1;
            }b2
        }b1
        {b3
            int e = x_0;
            int x_3 = 3;
            int f = x_3;
        }3
    }
    return;
}
```

The variable x_0 is a field and its scope is the entire class **Main**. The assignment to **a** will use x_0 (it will be assigned to the value 0) since no local variable **x** has yet been declared within **main**. The variable x_1 is a local variable and its scope is from its declaration until the end of the first block statement (block statement labeled with subscript **b1**) within **main**. The assignment of **b** comes before the declaration of x_1 so **b** is assigned the value of x_0 (*i.e.*, 0). The assignment of **c** comes after the declaration and within the first block statement so **c** is assigned the value of x_1 (*i.e.*, 1). Within the first block statement, there is a second block statement (labeled with subscript **b2**). Within this second block statement, it would be illegal to declare another **x** (x_2) since the two variables would have overlapping scopes. The variable **d** is assigned the value in x_1 since it is still within the first block statement and it appears after the declaration of x_1. The variable x_3 is another local variable like x_1 but it has a disjoint scope. Its scope is from the declaration until the end of the third block statement (labeled with subscript **b3**). The assignment of **e** comes before the declaration of x_3 so **e** is assigned the value of x_0 (*i.e.*, 0). The assignment of **f** comes after the declaration of x_3 and within the third block statement so **f** is assigned the value in x_3 (*i.e.*, 3).

Notice that the declaration of local variables can hide field variables. However, these can still

be accessed via **this**.

If statement. As in Java, an if statement allows the programmer to branch between statements based on some predicate. The format of an if statement is:

$$
\begin{array}{l}
\text{if (<predicate>)} \\
\qquad \text{<then statement>} \\
\left[\begin{array}{l} \text{else} \\ \qquad \text{<else statement>} \end{array}\right]
\end{array}
$$

where **if** and **else** are keywords, the <**predicate**> is a **boolean** expression, the <**then statement**> is a statement that is evaluated if the predicate evaluates to **true**, and the <**else statement**> is a statement that is evaluated if the predicate evaluates to **false**. As in Java, the else statement is optional. Here is an example if statement:

```
class Main {
    void main () {
        ...
        if (animal.getStrength() > 100)
            animal.fight(100, true);
        else
            animal.fight(100, false);
        ...
    }
}
```

Only one statement can be used as the then or else statement. However, if we need to execute more than one statement in either of these (or both) we can use a block statement:

```
class Main {
    void main () {
        ...
        if (animal.getStrength() > 100) {
            animal.fight(100, true);
            strength = animal.getStrength();
        }
        ...
    }
}
```

If more than one predicate needs to be tested, if statements can be cascaded as they are in Java. For example:

42

```
class Main {
    void main () {
        ...
        if (animal.getStrength() > 100)
            animal.fight(100, true);
        else if (animal.getStrength() > 50)
            animal.fight(50, false);
        else if (animal.getStrength() > 10)
            animal.fight(10, false);
        else
            animal.fight(0, false);
        ...
    }
}
```

This code is actually made up of three if statements not one. The last two if statements are nested within the previous if statement's else clause.

Unlike Java, Bantam does not have switch statements, however, these can often be implemented using cascaded if statements like those in the code above.

As shown in the example below, a potential ambiguity can arise due to the optional else statement:

```
if (x > y)
    z = x;
if (r > s)
    t = r;
else
    t = s;
```

In the code above, the else could potentially bind to the first or second if statement. As in Java, an else always binds to the innermost, unclosed if. Therefore, the else in the code above binds to the second if (*i.e.*, it is executed if and only if **r <= s**).

While statement. A while statement can be used to repeat a statement some number of times. The format of a while statement is:

```
while ( <predicate> )
    <statement>
```

where <**predicate**> is a boolean expression, and <**statement**> is executed while the predicate evaluates to **true**. The while loop is evaluated by first evaluating the predicate, and if this is **true** then

evaluating the statement, and repeating the process. The loop stops once the predicate evaluates to **false**. Each time the statement within the loop executes is called a loop *iteration*.

Here is an example, while loop:

```
class Main {
    void main () {
        ...
        while (animal.getStrength() < 100)
            animal.fight(100, true);
        ...
    }
}
```

This code repeatedly checks whether the **animal.getStrength()** is less than 100. If it is less than 100, then **animal.fight(100, true)** is executed. If it not less than 100, then the loop terminates.

Like if statements, the statement within a while loop can be replaced with a block statement.

For statement. A for statement is another form of a loop like the while statement. The format of a for statement is:

$$\text{for (} \big[\text{ <initialization> } \big] \text{ ; } \big[\text{ <predicate> } \big] \text{ ; } \big[\text{ <update> } \big] \text{)}$$
$$\text{<statement>}$$

where <**initialization**> is an expression for initializing a variable used within the loop, <**predicate**> is a boolean expression for determining when the loop should continue (**true**) and when it should stop (**false**), <**update**> is an expression for updating a variable within the for loop, and <**statement**> is executed while the predicate evaluates to **true**. The initialization, predicate, and update expressions are optional and be omitted.

The for loop is evaluated by first evaluating the initialization expression, which is not a part of the loop. Then the predicate is evaluated, and if it is **true**, another iteration of the loop is performed, otherwise, the loop terminates. Each iteration of the loop, the statement is executed followed by the update expression. If the predicate expression is omitted, then the predicate expression **true** (which is always **true**) is used.

Although a while loop can always be used in place of a for loop, a for loop is convenient if a counter variable (sometimes called an iterator variable) is required. For example, the loop below uses a counter variable **i**. Notice the initialization of **i** can be put in the initialization expression and

44

the increment of **i** can be put in the update expression.

```
for (i = 0; i < numArray.length; i = i + 1)
    numArray[i] = 100;
```

In the code above, the variable **i** is first set to 0. In each iteration of the loop, the next element of **numArray** is assigned the value 100 and **i** is incremented. The loop stops when **i** is no longer less than **numArray.length**. To summarize, this loop sets each element of an array **numArray** to 100.

Unlike in Java, a variable cannot be declared in the initialization component of the for loop. The initialization component in the Bantam Java for loop must be an expression, and consequently, there is no way to declare a variable within it.

Like if and while statements, the statement within a for loop can be replaced with a block statement.

Break statement. A break statement allows the programmer to exit a loop from within the body of the loop. The format of a Bantam Java break statement is:

```
break ;
```

The result of executing a break statement is that the innermost, encompassing loop is terminated, *i.e.*, control is immediately transferred to the end of that loop. Bantam Java does not support labeled break statements, which allow the programmer to break out of any encompassing loop. The break statement must be used within a loop (possibly within several loops, although it only breaks from the innermost loop). It is illegal to use a break statement in a location that does not reside within some loop.

The following contrived code uses a break statement within a for loop. It uses an auxiliary method **print** (defined somewhere else in the class) to print an **int** to the screen.

```
for (i = 0; i < 10; i = i + 1) {
    if (i == 7)
        break;
    print(i);
}
```

The loop above terminates after **i** reaches 7 as a result of the break statement, and therefore

prints 0 through 7 rather than 0 through 9. Note that **i** is 7 when the loop terminates not 8 since the update expression (**i = i + 1**) is not evaluated after the break statement is executed.

Although not apparent from the example above, break statements can be convenient within large, complex loops.

Expression statement. An expression statement is a statement made up of a single expression (which itself may contain subexpressions) and terminated by a semi-colon. Although we save the discussion of expressions until the next subsection, here is an example expression statement using a simple assignment expression (the variable on the left side of '=' is set to the computed value on the right side):

$$x = y + 2;$$

Not all expressions can be used within an expression statement. In fact, the only expressions that are legal within an expression statement are assignments, increment/decrement operations, method calls, and new object constructions. The code below is not a legal expression statement:

$$foo() + 2;$$

The addition of 2 to **foo()** is a useless operation in this case since it is not assigned to any variable, and for this reason, this is an illegal expression statement.

Return statement. A return statement returns from the method (*i.e.*, it exits the method and transfers control to the call site), possibly returning an expression to the call site. The format of a return expression is:

$$\text{return } \big[\text{ <expression> } \big] ;$$

where <**expression**> is an optional expression. If the method is non-void, then the return statement must return an expression whose type matches (for primitive types) or conforms (for object types) to the declared return type. It is an error if the expression is omitted or the expression is of the wrong type. If the method is void, then it is an error for the return statement to include an expression.

A return statement can be used anywhere within a method. The only requirement, as stated

Operator types	Operators
Arithmetic	**+ - * / %**
Relational	**== !=** $<$ $>$ **<= >=**
Boolean	**! &&** ||
Assignment	**=**
Increment/decrement (prefix and postfix)	**++ --**
Cast	**(type)(expression)**
Instanceof	**instanceof**
Member reference	**.**
Object construction	**new**
Array index	**[]**

Table 1: Bantam Java operators.

earlier, is that the last statement in a non-void method is a return statement, which cannot be embedded within another statement (to simplify the compiler).

Below is a simple method **isEven**, which determines whether an **int n** is even or not. It uses two return statements, which return either **true** (if **n** is even) or **false** (if **n** is odd).

```
boolean isEven(int n) {
    if (n % 2 == 0)
        return true;
    return false;
}
```

Because the declared return type of **isEven** is **boolean**, each return statement must include a return expression and that expression must have the type **boolean**. Note that in Bantam Java, it is not legal to put the "return false;" within an else clause. The last statement in the method must be a return statement, and this statement cannot be embedded within another statement.

Let's now look at the last category of language constructs in Bantam: expressions.

3.6 Expressions

Bantam supports a variety of expressions, which are described below. Figure 1 lists the operators that are supported in Bantam. Bantam supports most Java operators including arithmetic operators (**+**, **-**, *****, **/**, and **%**), relational operators (**==**, **!=**, $<$, **<=**, $>$, and $>=$), boolean operators (**!**, **&&**, \|\|), assignment (**=**), and increment/decrement shortcut operators (**++**, **--** – other shortcut operations are not supported). Bantam also supports object construction, member reference, (object) casting, **instanceof**, and array indexing **[]**. We discuss each of these in more depth below.

Every expression in Bantam computes a value with some data type. The type of the computed value is the type of the expression. Because expressions are used within field declarations, statements, or other expressions, the type of the expression is often important. For example, an expression that computes a value of type **int** cannot be assigned to a **boolean** field. It is the responsibility of the compiler to determine the type of each expression and verify that it is used in the appropriate way. For each expression below, we will discuss both what the expression computes and its resulting data type.

Assignment. An assignment expression has the form:

$$\Big[\ <\text{reference}> . \ \Big] \ <\text{name}> \ \Big[\ [<\text{index}>] \ \Big] \ = <\text{expression}>$$

where <**reference**> is an optional reference object (followed by '.'), <**name**> is the name of the target variable, <**index**> is the optional index expression (if <**name**> refers to an array), and <**expression**> determines the value to assign to the variable. The target must refer to some variable defined in the same scope or an encompassing scope. The reference object can be used only in the assignment of a field. Because fields are either **protected** (and **protected** is more restrictive in Bantam Java than Java) such as user-defined fields or cannot be modified such as **length** in an array type, the reference object, if used, must be either **this** or **super**.

Here is an example assignment:

$$x = y + 2;$$

The expression on the right is evaluated first and then stored in the lefthand variable. The result of the entire expression is the result of the righthand expression.

If the type of the righthand expression or the lefthand variable is primitive then the types must match, otherwise the (object) type on the right must conform to the (object) type of the lefthand variable. The static type of the entire expression is the static type of the righthand expression.

Here is another example assignment, but this time the righthand expression is assigned to an element in an (**int**) array:

$$\text{this.numArray[i-1]} = x * 5;$$

48

Dynamic dispatch. A dynamic dispatch is simply a method call of an instance method. To call a method in Bantam, we need a **reference object** (of an appropriate type) to invoke the method on. The reason for the name *dynamic dispatch* is that in object-oriented languages a method call of an instance method is often thought of as passing a message to the object, hence the word "dispatch" (the "dynamic" part will be discussed shortly). Notice that because the called method can access fields, the outcome of the method depends on the reference object.

The format of a dynamic dispatch is:

$$\Big[\text{ <reference> . } \Big] \text{ <name> (<parameters>)}$$

where <**reference**> is an optional expression that corresponds to the reference object, <**name**> is the name of the method, and <**parameters**> is a list of comma-separated actual parameters, each of which is an expression. If the reference expression is not included, then **this** is used as the reference object.

In order for the dynamic dispatch to be legal, the called method must be declared within the reference object's static type class or within a superclass. The number of actual parameters and declared formal parameters must match and the type of each actual parameter expression must conform to the type of the declared formal parameter or match if either type is a primitive or an array type. Finally, the dynamic dispatch must be used according to its declared return type. For example, a **void** method should not be used within a larger expression since it has no return value. A method that returns a **boolean** cannot be used in an arithmetic expression where an **int** is required.

The following code below shows how to dispatch on methods **fight** and **getStrength** within class **Animal** (defined on page 37):

```
class Main {
    Animal animal = new Animal();

    void main() {
        animal.fight(20*5, true);
        int halfStrength = animal.getStrength() / 2;
        ...
    }
}
```

The dispatch on **fight** passes an **int** (**20*5**) as the first parameter and a **boolean** (**true**) as the second parameter, which match the declared formal parameters of **fight**. **fight** is a **void** method so it must

appear on a line by itself and not within a larger expression. The method **getStrength** takes no parameters and returns an **int**. Therefore, we can use the result of **getStrength** in an expression that requires an **int**.

The actual parameter expressions in a dynamic dispatch are executed from left to right and passed to the method. They are evaluated after the reference expression. Bantam Java, like Java, uses *call by value* meaning that the formal parameter in the method is assigned the value from the corresponding actual parameter expression. In the code above, the formal parameter **amount** is set to **20*5** (*i.e.*, **100**) and the formal parameter **isWinner** is set to **true**.

Below is another example of a (contrived) method call. In the call to **foo** from method **main**, because of call by value, the formal parameter **y** is set to the value in **main**'s variable **x**. But **x** and **y** are separate variables. If **foo** sets **y** to a new value, this will have no effect on **x** in **main**.

```
class Main {
    void main() {
        int x = 3;
        foo(x);
    }

    void foo(int y) {
        ...
    }
}
```

The parameter in the code above has a primitive type (*i.e.*, **int**). Call by value is also used for parameters with object types. However, as variables in Bantam Java store references (or memory pointers) to objects rather than the objects themselves, it is the reference that is copied not the object itself. For example, in the code below, the **animal** reference is passed to method **foo**.

```
class Main {
    void main() {
        Animal animal = new Animal();
        foo(animal);
    }

    void foo(Animal a) {
        ...
    }
}
```

While in **foo**, the variables **a** and **animal** (which is not actually accessible in **foo**) refer to the

same object (*i.e.*, alias). If we modify a field of **a**, then the field will be changed for **animal**. Of course, if we assign **a** to another **Animal** object, then the two variables would no longer alias.

Because of inheritance, the exact method that is called on a dispatch is not always known at compile time, which is the reason for the word "dynamic" in dynamic dispatch. This property of object-oriented languages, *i.e.*, that one call site can call multiple methods during the course of a program, is known as *polymorphism*. For example, the class **Tiger** could define its own version of **fight**:

```
class Tiger {
        <member definitions>

        void fight(int amount, boolean isWinner) {
          if (isWinner)
              strength = strength + amount * 10;
          else
              strength = 0;
        }
}
```

If the method **fight** is called on an **Animal** object, which **fight** is called? It depends on the runtime type of the reference object, which as discussed above may differ from the static type (*i.e.*, **Animal**). If at runtime, the object has type **Animal**, then the **fight** method within the **Animal** class is called. If, on the other hand, the object has dynamic type **Tiger**, then the **fight** method within the **Tiger** class is called. This example is illustrated in the **Main** class shown on page 49. Although we know the original dynamic type of **animal** is **Animal**, the dynamic type of this field may change later, and in general, it is impossible for the compiler to determine the dynamic type of a variable. The compiler must generate code for calling the appropriate method at runtime based on the dynamic type of the reference object. A runtime error is generated if the reference object is **null**.

The reference expression can evaluate to any object type (so long as that type includes the appropriate method) or can refer to the special variables **this** or **super**. Using **this** is not necessary since excluding the reference object has the same effect. However, using **super** is often helpful. It allows the programmer to call hidden methods. For example, in the **Tiger** class, if we want to call the **fight** method from **Animal**, we can only do this using **super** as the reference object.

Object construction. We can use the **new** operator to create a new (non-array – array construction is discussed below) object from a class (called *object construction*). The format of a **new** operation

is:

$$new\ <class\ name>()$$

where **new** is a keyword and <**class name**> is replaced with a class name. For example, if we wanted to create an object from a class **Tiger** and store it in a variable of type **Tiger** we could do the following:

$$Tiger\ t = new\ Tiger()$$

In Java, the programmer can define their own constructors that are called when the object is created. These allow the programmer to customize the object based on the particular context in which it is created. User-defined constructors are not allowed in Bantam. Object construction in Bantam creates a new object and initializes each field based on the initialization expression (or lack of). However, a Bantam programmer can write methods that simulate user-defined Java constructors. Of course, it is up to the programmer to make sure that the pseudo-constructor is called after every object construction. For example, in the **Animal** class, we could define a pseudo-constructor to set the number of legs and the flag indicating whether the animal can fly.

```
class Animal {
    <member definitions>

    // simulates a constructor
    // must always be called immediately after construction
    Animal init(int l, boolean f) {
        numLegs = l;
        canFly = f;
        return this;
    }
}

class Main {
    void main() {
        Animal animal = (new Animal()).init(4, false);
        ...
    }
}
```

Array construction. We can also use the **new** operator to construct a new array. The format of array construction is:

$$\text{new} < \text{base type} > [\ < \text{size} >]$$

where <**base type**> is a primitive type or a class name and <**size**> is an **int** expression, which evaluates to the size of the array. The size must be non-negative, and for reasons related to garbage collection, must be less than or equal to 1500, but only if garbage collection is enabled. A runtime error occurs if the size evaluates to an illegal value (*e.g.*, -1). If garbage collection is disabled the array size is unbounded, although it still must be non-negative. This upper bound is not a part of the language, but rather a part of the particular implementation of the language.

The code below constructs an array of 10 **int** elements and assigns the newly-constructed array to the variable **numArray**:

```
int[] numArray = new int[10];
```

Each element is initially set to a default value. If the base type is **int** (as it is above), then the default value is 0, if the base type is **boolean**, then the default value is **false**, and if the base type is an object type, then the default value is **null**.

Casting. As we have already seen, we can store an object of one type into a variable whose type is a super type. For example, we can store a **Tiger** object into a variable with type **Animal**. This action is called *upcasting*, since we are moving up the class hierarchy tree. Upcasting can be done either implicitly or explicitly. For example, the following two statements are equivalent:

```
Animal animal = new Tiger();
            OR
Animal animal = (Animal)(new Tiger());
```

In the first assignment statement, we implicitly upcast a **Tiger** object into an **Animal** so that it can be stored in **animal**. In the second assignment statement, we do the same except we explicitly show the upcast. The format of an explicit cast is as follows:

$$(\ < \text{class name} >)\ (\ < \text{expression} >)$$

53

where <**class name**> is the name of some existing class and <**expression**> is an expression with an object type (*i.e.*, not a primitive). In Bantam, unlike in Java, the expression to be casted must be enclosed by parentheses (this simplifies parsing).

In the example cast above (*i.e.*, **(Animal)(new Tiger())**) only the static type of **animal** is **Animal**. The runtime type is still **Tiger**. The additional **Tiger** functionality (non-inherited members) is still maintained in the variable **animal**, however, it cannot be utilized. But it can be recovered by performing a *downcast*. For example, we could do the following:

<div align="center">

Tiger tiger = (Tiger)(animal);

</div>

In this case, the cast must be explicit since it needs to be checked at runtime, *i.e.*, the compiler does not know statically whether **animal** is actually a **Tiger** or a subclass of **Tiger**. If the dynamic type of **animal** is not **Tiger** or a subclass of **Tiger** (*e.g.*, **Bear**) then a runtime error occurs.

Sometimes the compiler can statically catch an illegal cast, *i.e.*, a cast that is neither an upcast or a downcast. In particular, any cast from one class to another class, which is neither a subclass or a superclass, is a compiler error. For example, the following is a compiler error:

<div align="center">

Bear bear = (Bear)(new Tiger());

</div>

Upcasting and downcasting are common operations in Bantam (as well as in Java, prior to the addition of generics). Upcasting and downcasting make it possible to create collections of heterogeneous objects like the **Vector** class in Java. The collection is simply an object that stores each element as an **Object**, *i.e.*, objects are upcasted to **Object**. Because **Object** is a superclass of all classes, the collection can house any kind of object (but not a primitive). When an object is retrieved from the collection, it usually must be downcasted from **Object** to its original type.

As in Java, arrays can be casted just like other objects. For example, we could perform the following downcast, which converts an array of **Animal** objects into an array of **Tiger** objects:

<div align="center">

Tiger[] tigers = (Tiger[])(animals);

</div>

Of course, the cast above is legal if and only if the array **animals** holds **Tiger** objects (or objects

whose type is a subtype of **Tiger**). A runtime error occurs if that is not the case. However, an additional problem can arise when casting arrays as shown in the code below:

```
Animal[] animals = (Animal[])(tigers);
animal[0] = new Bear();
```

The problem is that **animals** holds **Tiger** objects. It cannot be used to store **Bear** objects. A runtime error occurs when an array element is assigned a value that does not conform to the dynamic base type of the array.

Bantam Java, unlike Java, does not have support for casting primitives. Therefore, it is a compiler error if the programmer attempts to cast to a primitive type or cast an expression, which has a primitive type. A compiler error also occurs if the target type does not exist.

Instanceof. The **instanceof** operation in Bantam allows the programmer to determine the runtime type of an object. Like casting, **instanceof** is often a useful operation when dealing with collections of heterogeneous objects. If we pull an object off of a collection, such as a **Vector**, we may not know its exact type, but rather that the object is one of several types. The format of **instanceof** is:

<expression> instanceof <type>

where <**expression**> is an expression with an object type (*i.e.*, not a primitive) and <**type**> is some type (*i.e.*, name of a class or an array type). **instanceof** evaluates to **true** if the dynamic type of the expression conforms to the righthand type. Otherwise, it evaluates to **false**.

The following code illustrates a common use of **instanceof**. The code retrieves a generic **Animal** object from a list (via the method **getNextAnimal**). The **instanceof** operator is used to determine the specific type of **Animal** (**Tiger**, **Bear**, or **Eagle**) so that a counter can be updated, which keeps track of the total number of each kind of **Animal**. (Note: if the dynamic type of **animal** is **Animal** then none of the **instanceof** expressions will evaluate to **true** and none of the counters will be incremented.)

```
Animal animal = getNextAnimal();
if (animal instanceof Tiger)
    numTigers = numTigers + 1;
else if (animal instanceof Bear)
    numBear = numBear + 1;
else if (animal instanceof Eagle)
    numEagles = numEagles + 1;
```

We can also use **instanceof** on array types. For example, the following code tests whether an array **animals** holds **Tiger** objects, and if so, casts **animals** to a **Tiger[]** array:

```
if (animals instanceof Tiger[]) {
    Tiger[] tigers = (Tiger[])(animals);
    ...
}
```

A compiler error occurs if the lefthand expression in the **instanceof** expression has a primitive type. An error also occurs if the righthand type in the **instanceof** is primitive or does not exist. As with casting, a compiler error also occurs if the expression's type is neither a subtype nor a supertype of the righthand type.

Arithmetic operators. Bantam supports five binary arithmetic operators (+, -, *, /, %) and one unary arithmetic operator (-). The format of the binary operators is:

$$<\text{left expression}> <\text{operator}> <\text{right expression}>$$

where <**left expression**> and <**right expression**> are both expressions with type **int** and <**operator**> is one of '+', '-', '*', '/', or '%'. The operators are similar to those in Java. **+** computes the sum of the left and right expressions, **-** computes the difference, ***** computes the product, **/** computes the integer division, and **%** computes the remainder when dividing the left expression by the right expression. The resulting type of these expressions is an **int**.

The format of the unary operator is:

$$- <\text{expression}>$$

where <**expression**> is an **int** expression. This operation computes the arithmetic negation of the expression. The resulting type is an **int**.

56

Here are some examples:

```
int x = 0;
int y = 1;
int z = 2;

x = y + z; // x is set to 3
x = y - z; // x is set to -1
x = y * z; // x is set to 2
x = y / z; // x is set to 0
x = y % z; // x is set to 1
x = - y; // x is set to -1
```

Boolean operators. Bantam supports two binary boolean operators (&& and ||) and one unary boolean operator (!). The format of the binary operators is:

<left expression> <operator> <right expression>

where <**left expression**> and <**right expression**> are both expressions with type **boolean** and <**operator**> is one of '&&' or '||'. **&&** computes the logical AND of the left and right expressions. If both the left and right expressions evaluate to **true**, then the result of **&&** will also be **true**. Otherwise, the result is **false**. || computes the logical OR of the left and right expressions. If both the left and right expressions evaluate to **false**, then the result of || will also be **false**. Otherwise, the result is **true**. The resulting type of these expressions is a **boolean**.

The format of the unary operator is:

! <expression>

where <**expression**> is a **boolean** expression. **!** computes the complement of the expression. If the expression evaluates to **false**, the result is **true**. If the expression evaluates to **true**, the result is **false**. The resulting type of these expressions is a **boolean**.

Here are some examples:

```
boolean b1 = false;
boolean b2 = false;
boolean b3 = true;

b1 = b2 && b3; // b1 is set to false
b1 = b2 || b3; // b1 is set to true
b1 = !b2; // b1 is set to true
```

Relational operators. Bantam supports six binary relational operators (==, !=, $<$, $<=$, $>$, $>=$) for comparing two values. The format of the relational operators is:

<center><left expression> <operator> <right expression></center>

If <**operator**> is '==' or '!=' then the two expressions must have the same type if either is primitive or one type must conform to the other type if either is an object. If <**operator**> is '$<$', '$<=$', '$>$', or '$>=$' then the expressions must both have type **int**. The result of a relational operation is a **boolean** indicating whether the test succeeded.

The **==** operator tests if the left and right expressions are equivalent. Note: if the expressions are objects the operator tests if the left and right expression refer to the exact same object. If they refer to different objects then the test fails, even if the objects have the exact same field values. If the left and right expressions are equivalent then the result of **==** is **true**, otherwise it is **false**. The **!=** operator works like **==** except that it tests if the left are right expressions are not equivalent. If the left and right expressions are equivalent then the result of **!=** is **false**, otherwise it is **true**.

The $<$ operator tests if the lefthand **int** is less than the righthand **int**, $<=$ tests if the lefthand **int** is less than or equal to the righthand **int**, $>$ tests if the lefthand **int** is greater than the righthand **int**, and $>=$ tests if the lefthand **int** is greater than or equal to than the righthand **int**. If the test succeeds then the result is **true**, otherwise the result is **false**.

Here are some examples:

```
boolean b = false;
int i1 = 0;
int i2 = 1;
int i3 = 0;
Object o1 = new Object();
Object o2 = new Object();
Object o3 = o1;

b = i1 == i2; // b is set to false
b = i1 != i2; // b is set to true
b = o1 == o2; // b is set to false
b = o1 == o3; // b is set to true
b = i1 < i2; // b is set to true
b = i1 >= i3; // b is set to true
```

Increment/decrement shortcuts. Bantam Java (the extended language) also includes the increment/decrement shortcut operators (**++** and **--**). These are unary operator expressions. As in Java, these can be used as either prefix or postfix operators. The format of the **++** prefix shortcut operator is (the format of the **--** prefix shortcut is the same):

$$++<\text{variable}>$$

where $<$variable$>$ is either an **int** variable or an element in an **int** array. A compiler error occurs if the variable does not have type **int**. This expression is a short cut for:

$$<\text{variable}> = <\text{variable}> + 1$$

The outcome of a (prefix) **++** or **--** is that the variable is incremented or decremented, respectively, and the result is returned. For example, the code below passes the result of incrementing **x** (using prefix **++**) to an auxiliary method **print**, which prints an **int** to the screen.

```
int x = 5;
print(++x);
```

The result of executing the code above is that **x** is set to 6 and 6 is printed to the screen.

In addition, **++** and **--** can be used as postfix operators. The format of the postfix **++** and **--** is the same as the prefix **++** and **--** except that the operators appear after the variable rather than before. The semantics of the postfix operators are slightly different than the corresponding prefix

59

operators. The variable is still incremented/decremented, however, the expression evaluates to the previous value of the variable rather than the incremented/decremented value. For example, the following code prints 5 not 6 (note the change from **++x** to **x++**):

```
int x = 5;
print(x++);
```

But although the code above prints 5, **x** is still incremented. If we added the statement **print(x);** to the bottom of the code above, this would print 6.

The increment/decrement operators are particularly convenient within the for loop. For instance, we can use **++** to iterate over the elements of the array and print them out:

```
for (i = 0; i < numArray.length; i++)
    print(numArray[i]);
```

Note: other shortcut operations such as **+=**, ***=**, and **&&=** are not supported.

Constants. There are three constant types in Bantam: **int**, **boolean**, and **String**. An **int** constant is any sequence of numeric digits (*e.g.*, **0**, **124**). It must be between **0** and **2147483647** ($2^{31} - 1$) since it is stored as a 32-bit signed integer. Integer constants must be written in decimal, other forms such as hexadecimal are not supported. Leading 0's are allowed in an **int** constant.

A **boolean** constant is either **true** or **false**. These are keywords and may not be used as identifiers.

A **String** constant is a sequence of characters between double quotes (*e.g.*, "abc"). Unlike in Java, Bantam Java represents characters using an ASCII encoding rather than an Unicode encoding, which limits the types of characters that can be represented in a Bantam Java **String** constant. As in Java, **String** constants are treated as objects of type **String** (which is discussed in 3.7). **String** constants may not span multiple lines. A backslash is interpreted as an escape character within a **String** constant, which allows the programmer to specify some special characters. These include \n (newline), \t (tab), \" (double quote), \\ (backslash), and \f (form feed). Other special characters are illegal. In particular, a backslash within a **String** constant must be followed by 'n', 't', 'f', '\', or a double quote. For garbage collection reasons, the size of a **String** constant (and a **String** variable) is limited to 5000 characters.

Variables. A variable reference is also an expression. Of course, the variable must have been

., (params), [], ++ (postfix), -- (postfix)
- (unary), !, ++ (prefix), -- (prefix)
new, cast
*, /, %
+, -
<, >, <=, >=, instanceof
==, !=
&&
||
=

Table 2: Operator precedence and associativity. Operators with higher precedence are listed higher in the table than those with lower precedence. Operators in the same row have the same precedence. All binary operators are left-associative except assignment ('='). These rules were taken from the Java book by Arnold, Gosling, and Holmes (Gosling led the team at Sun, which created Java) [5]. Note: (params) refers to the parentheses and actual parameter list in a method call.

previously declared, otherwise, it is an error. If the variable is a field, then it can be preceded by **this** followed by '.'. If the variable is an inherited field (hidden or otherwise) then it can be preceded with **super** followed by '.'. Finally, if the variable is **length**, then it can be preceded by an array variable followed by '.' as this is a public field within each array.

If the variable is an array (as opposed to the reference object), then it can be indexed to access a particular element. The format for accessing an array element is:

$$\Big[<reference> . \Big] \ <name> [<index>]$$

where <**reference**> is the optional reference object (either **this** or **super**), <name> is the name of the array, and <**index**> is an expression that evaluates to the array index.

The code below sets **x** to the third element (index 2) in the field **numArray**:

```
int x = numArray[2];
```

Parentheses. The final type of expression is an expression within parentheses: **(** <**expression**> **)**. Parentheses are useful for grouping operations. Table 2 shows the order of operations and the associativity for each operation discussed in this section. Parentheses should be used to override

precedence and/or associativity. For example, we could do the following to force the addition to occur before multiplication:

int x = (2 + 3) * 5; // x is set to 25

3.7 Built-In Classes

There are four built-in classes in Bantam. These are listed in Table 3.

Object. The **Object** class is similar to the **Object** class in Java except with less functionality. The **Object** class is the only built-in class that can be extended, and in fact, all classes must either directly or indirectly extend **Object**.

   Object has three methods: **clone**, **equals**, and **toString**. The method **clone** takes no parameters and makes a copy of the reference object. If the cloned object contains fields (inherited or otherwise) then the values of these fields are copied to the new object. If these fields are themselves objects, however, the contents of those objects are not copied. In this case, the field in both the original and the copy would refer to the same object. To overcome this limitation, one can always redefine the **clone** method when defining a new class.

   The method **equals** compares two objects. It takes as a parameter an object to compare to the reference object and returns a boolean indicating whether the two objects are equivalent. The method checks to see if two objects alias (*i.e.*, refer to the same address). It does not check the contents of the two objects. However, this method can be overriden by other classes to compare the contents of two objects.

   The method **toString** returns a **String** representation of the reference object. The default representation has the following format: <dynamic type>@<address>. For example, calling **toString** on variable **o** might return **Sys@3201** if the dynamic type of **o** is **Sys** and its address is **3201**. In general, this method should be overriden by other subclasses to return a more informative **String** representation (as in the **String** class below).

Sys. The **Sys** class contains one method for exiting the program. It extends **Object**. The method **exit** takes one parameter, an **int** representing the exit status (0 is generally used for no error, non-zero values are used when an error has occurred). It replaces the static call to **System.exit** in Java. The **Sys** class cannot be extended.

Class	Method signature	Method description
Object	**Object clone()**	copy an object
	boolean equals(Object s)	test if objects are equal (*i.e.*, alias)
	String toString()	return string representation of object
Sys	**void exit(int status)**	exit program with specified status
	int time() *	return UTC time
	int random() *	return random int
String	**int length()**	return string length
	boolean equals(Object s)	test if strings are equivalent
	String toString()	return itself
	String substring(int beginIndex, **int endIndex)**	return substring between the indices
	String concat(String s)	return concatenated string
TextIO	**void readFile(String filename)**	set to read from specified file
	void writeFile(String filename)	set to write to specified file
	void readStdin()	set to read from standard input
	void writeStdout()	set to write to standard output
	void writeStderr()	set to write to standard error
	String getString()	read next string
	int getInt()	read next int
	void putString(String s)	write specified string
	void putInt(int i)	writes specified int

Table 3: Bantam built-in classes. Note: **Sys**, **String**, and **TextIO** cannot be extended. * indicates methods that are not automatically included when targeting MIPS.

The **Sys** class may also contain two other methods depending on your implementation of the Bantam Java compiler. In particular, it can contain a method for obtaining the time and for computing a random integer. The method **time** computes the seconds since epoch (January 1st 1970) in UTC. The method **random** computes a random integer. If targeting MIPS, these methods are not automatically included since they cannot be implemented on the unmodified SPIM simulator [15]. Email the first author to obtain a MIPS runtime system and SPIM simulator that does support these methods. This is not an issue when using the other two targets (JVM or x86).

String. The **String** class is a class for representing sequences of characters. It extends **Object** and cannot be extended. It is similar to the **String** class from Java, but with less functionality. For reasons having to do with garbage collection, the length of a **String** may not exceed 5000 characters. The characters in a **String** are represented using ASCII rather than Unicode, which limits the types of characters that can be stored in a Bantam Java **String**.

String has four methods: **length**, **equals**, **substring**, and **concat**. The method **length** takes no parameters and returns the length (an **int**) of the **String**. For example, if the variable **s** is a **String**

representing "abc", then calling **s.length()** returns 3.

The method **equals** compares two strings. It overrides the **equals** method from the the **Object** class. It takes one parameter: an **Object** to compare to the reference **String**. If the parameter has the runtime type **String** and it represents the same sequence of characters as the reference string then **equals** returns **true**. Otherwise, it returns **false**. A runtime error occurs if the **String** parameter is **null**. Notice that this method behaves differently than **==**. The operator **==** can only determine if the two strings refer to identical objects. It cannot be used to determine if two strings represent the same sequence of characters.

The method **toString** takes no arguments and simply returns itself. It overrides the **toString** method from the **Object** class.

The method **substring** takes two parameters: a beginning index (an **int**) and an end index (an **int**). It returns the **String** that starts at the beginning index and ends at the end index. The character at the beginning index is included in the resulting **String**, but the character at the end index is not included in the **String**. If **s** is a **String** representing "abc", then calling **s.substring(1,3)** returns a **String** representing "bc". A runtime error occurs if the beginning index is less than 0, the end index is greater than the length of the **String**, or the beginning index is greater than the end index.

Finally, the method **concat** takes as parameter a **String** and returns a new string representing the concatenation of the reference **String** object with the parameter **String** object. For example, if **s** is a **String** representing "abc" and **s2** is a **String** representing "def", then **s.concat(s2)** would return a new **String** representing "abcdef". A runtime error occurs if the **String** parameter is **null** or if the resulting **String** has more than 5000 characters.

There is no method for copying strings, but because **String** extends **Object**, the **clone** method can be used to copy a **String**.

TextIO. The **TextIO** class provides methods for performing file input and output. It extends **Object**. It is similar to the **TextIO** class defined by David Eck in his textbook [12], except the class does not use **static** members. Reading and writing are separately handled by this class and so it is possible with one **TextIO** object to read from one file, while writing to another. The **TextIO** class contains a method called **readFile** for setting the read file. This method takes as input a **String** representing the file name. It terminates the program if the file does not exist or the program does not have permission to read from the specified file. There is also a method called **writeFile** for setting the write file. This works similarly to **readFile** except the specified file need not exist. If the file

does not exist, then **writeFile** will create the specified file. The program does, however, have to have permission to write to the specified file. For both **readFile** and **writeFile** the programmer can specify a path along with the file name, *e.g.*, "../file.txt" or "/home/someuser/foo". A runtime error occurs in either method if the input **String** is **null**.

TextIO also contains methods for reading from standard input (**readStdin**) and writing to either standard output (**writeStdin**) or standard error (**writeStderr**). These methods do not take any parameters. By default, **TextIO** reads from standard input and writes to standard output.

None of the **TextIO** methods described above actually read or write any data, instead they set either the read location or the write location. To read or write from some location, the programmer can use **getString**, **getInt**, **putString**, or **putInt**. The method **getString** reads the next line of text (minus the newline character) from the current read location and returns a **String** representing this text. The method **getInt** reads the next line of text (minus the newline character) from the current read location and interprets the text as an **int**. If an **int** was entered then **getInt** returns the entered value otherwise it returns 0. The method **putString** takes a **String** parameter and writes it to the current write location. A runtime error occurs in **putString** if the input **String** is **null**. The method **putInt** takes an **int** parameter and writes it to the current write location. Neither **putString** or **putInt** write a newline character to the output stream, however, this can be achieved by performing **putString("\n")** after writing the **String** or **int**.

Example. Figure 3 shows a somewhat contrived example class that uses all of the built-in classes. There is one **Main** class with three methods: **main**, **getNextLine**, and **error**. The **main** method prompts the user via standard input for a positive integer **n** and then calls **getNextLine n** times to read **n** lines from the file "input.txt". It stops early if the program reads the string "quit". The method **getNextLine** reads the next line from "input.txt" and calls **error** if either there is not another line to read or if the length of the read string is less than 2. The method **error** prints an error message to standard error and exits the program with status 1. After **main** gets the next line from **getNextLine** it concatenates it with the previous read output. At the end of the method, **main** prints the concatenated output to standard output and to the file "output.txt".

3.8   Bantam vs. Java

Although we have described some of the differences between Bantam and Java throughout this section, we lay out the main differences below. We also show how to convert a Bantam program

```
class Main {
    TextIO io = new TextIO();
    String output = "";

    void error() {
        io.writeStderr();
        io.putString("Bad input; exiting\n");
        (new Sys()).exit(1);
    }

    String getNextLine() {
        String s = io.getString();
        if (s == null || s.length() < 2)
            error();
        return s.substring(1, s.length());
    }

    void main() {
        String s = "";
        io.readStdin();
        int n = io.getInt();
        if (n < 1)
            error();

        io.readFile("input.txt");
        int i = 0;
        for (i = 0; i < n && !s.equals("quit"); i++) {
            s = getNextLine();
            output = output.concat(s).concat("\n");
        }

        io.writeStdout();
        io.putString(output.toString());

        io.writeFile("output.txt");
        io.putString(output);

    }
}
```

Figure 3: Example program using all of the built-in classes.

to a Java program (note: converting a Java program to a Bantam program is a more difficult process, especially if the Java program uses a significant number of features that were excluded from Bantam Java).

Differences. The Bantam Java (extended) language excludes the following features (this list may not be complete):

- Primitives besides **int** and **boolean**
- Multi-dimensional arrays

66

- Static members
- Packages
- Modifiers such as **public**, **private**, **protected**, and **package**
- Nested classes
- Abstract classes
- Interfaces
- Final variables or methods
- Enumerations
- Generics
- Loops besides while loop and for loop
- Labeled break statements and continue statements
- Switch statements
- Logic operators ($\&$, $|$, $\hat{}$, $<<$, $>>>$, $>>$)
- Shortcut operations besides **++** and **--** (*e.g.*, **+=**, ***=**)
- Conditional expressions
- Hexadecimal/octal representation of **int** constants
- User-defined constructors
- Method overloading
- Declaration statements without assignments
- Unicode characters in **String** objects (ASCII is used instead)
- Extensive library for graphics, networking, mathematical functions, *etc.*

Converting Bantam programs to Java programs. Because the Bantam language is a subset of the Java language, it is always straightforward to convert Bantam programs into Java programs. The programmer must perform the following steps:

- Add the **protected** keyword in front of all fields
- Add the **public** keyword in front of all methods
- Create classes **TextIO** and **Sys** in Java
- Add the following static **main** method to the **Main** class (without replacing the original **main** method):

```
public static void main(String[] args) {
    (new Main()).main();
}
```

Otherwise, the program will not need to be changed.

3.9 Last Words

We have now looked closely at the Bantam Java language. As was discussed throughout this section, it is the responsibility of the Bantam Java compiler to enforce the lexical, syntactic, and semantic rules of the Bantam Java language layed out in this section. Therefore, this section will be an important resource when implementing the compiler. The next section gives an overview of the Bantam Java compiler, which must be completed by the student.

Figure 4: Bantam Java compiler phases.

4 Bantam Java Compiler

This section describes the Bantam Java compiler. This section first gives an overview of the compiler and then discusses the data structures that students will need to use in course projects. Note: this section does not discuss the compiler components that are specific to optimization (an optional project), which are instead discussed in Section 6.

4.1 Overview

The compiler is written in the Java programming language and will run on any x86/Linux machine (and perhaps other Unix-based machines). As shown in Figure 4, the compiler is split into four phases: lexical analysis, syntactic analysis, semantic analysis, and code generation. This is a standard design template used in many educational compilers [3, 4, 22, 21]. These phases are unimplemented; students must complete them in four course assignments. Note: this design structure will vary slightly if using some of the optional projects such as the optimization project (which is discussed in Section 6).

Compiler phases. The lexer and parser are built using automatic lexer and parser generators (either JLex [7]/Java Cup [14] or JavaCC [6]). Students must implement the lexer and parser specification files. The semantic analysis and code generation assignments are somewhat more challenging. Students must build these almost from scratch given some auxiliary data structures (*e.g.*, symbol table). These two assignments will give students ample opportunity to make important design choices, such as how to build class hierarchy trees and class environments.

Abstract syntax tree. For the most part, the compiler uses a single intermediate form (minus the optimizer, which is an optional project): an abstract syntax tree (AST). Both the semantic analyzer and the code generator must traverse this AST in order to check the semantics of the program and generate code, respectively.

Auxiliary data structures. Several auxiliary data structures are provided for use in each of the compiler phases. These include (among others) a symbol table, an error handler, and AST nodes.

69

lexer	Performs lexical analysis
parser	Performs syntactic analysis
semant	Performs semantic analysis
codegenmips	Generates Mips code
codegenx86	Generates x86 code
util	Utility classes (symbol table, error handler, class tree node, location)
ast	AST node classes
visitor	Visitor classes (generic visitor, print visitor)

Table 4: Bantam compiler packages.

Packaging. For modularity purposes, the compiler relies heavily on Java packaging. Figure 4 lists the packages used in the Bantam Java compiler. Each compiler phase is placed in a separate package: lexer, parser, semant, codegenmips, and codegenx86 (there are two packages for code generation, one for Mips and the other for x86.) In addition, three other packages provide auxiliary support (*e.g.*, symbol table): ast, util, and visitor. The ast package provides classes for representing the nodes of the abstract syntax tree. The util package provides classes for error handling, for representing the class hierarchy tree, and for representing the symbol table. The visitor package provides classes for implementing the visitor design pattern [13]. Finally, there are a few packages for use in optional projects. Packages opt and cfg contain the optimizer and the optimization representation (a control flow graph), respectively.

Visitor pattern. Again for modularity purposes, the source code also utilizes the visitor design pattern [13] for traversing a tree (*i.e.*, the AST) from an outside class. With the visitor pattern, AST traversals for checking the program semantics can reside within the semantic analysis package. Similarly, AST traversals for generating code can reside within the code generation package. We will discuss the visitor design pattern in more detail in 4.2.

Runtime system. The Bantam compiler includes a runtime system, which handles memory allocation and deallocation (*i.e.*, the garbage collector) as well as all built-in object methods. The runtime system is provided to the student as it would take most students more than a single semester to implement on top of the rest of the compiler.

4.2 Data Structures

Below we describe the auxiliary data structures used by the compiler. This documentation is not complete, for example, we do not discuss every public method and field. This information is

already available within the API documentation (see Section 2 for instructions on how to make the API documentation). Instead, we describe at a high-level how to use each data structure.

Error handler. The error handler data structure (**ErrorHandler**) allows for a unified way of handling and printing error messages from within any compiler phase. Each phase has an error handler object, which it uses to report errors. To report an error it must be registered with the error handler object by calling the **register** method. The **register** method does not immediately print the error, it is printed only after a call to the **checkErrors** method. **checkErrors** prints out all errors (if there are any) and exits if there are any errors. Before printing the errors, **checkErrors** sorts the errors based on where the errors were discovered, *i.e.*, the filename and line number. It first sorts the errors based on filename (using alphabetic order) and then sorts errors within the same file based on line number (using numeric order).

Class tree node. In order to compile a Bantam program, we must build a representation of the source program's class hierarchy tree (shown in Figure 1). The class tree node (**ClassTreeNode**) data structure represents each node (*i.e.*, a class) in the class hierarchy tree. Each class tree node contains a parent class tree node, which is **null** by default. It can and eventually should be set by calling **setParent** (only **Object** can have a **null** parent link). A class tree node also contains a list of child class tree nodes, which is empty by default. A child can be added by calling **addChild**. In general, only one of **setParent** and **addChild** needs to be called since both methods automatically set both the parent and child links. The parent class tree node, child class tree nodes, total children, and total descendants (total strict subclasses) can be accessed via public getter methods.

Each class tree node also has a variable and method symbol table associated with it, which can be retrieved by calling **getVarSymbolTable** and **getMethodSymbolTable**, respectively. These are the symbol tables for the particular class that the class tree node represents. Note: two symbol tables are needed since the variables and methods have separate name spaces. The class tree node also contains other meta information about the class, such as the class name and whether the class is extendable.

Location. The location data structure (**Location**) is a simple data structure for representing the location of a variable. It is used exclusively by the code generator (in concert with the symbol table) to record where a variable was allocated. The location data structure supports both variables allocated to memory and in a register. This information is needed during code generation. At a variable declaration, the variable name and location is added to the symbol table. At later

references to the same variable, the location can be retrieved from the symbol table.

Symbol table. The Bantam compiler includes a symbol table (**SymbolTable**) for keeping track of all the identifiers and constants that appear within the source program. Unlike in commercial compilers, the symbol table in the Bantam compiler stores all symbols as strings. This simplifies the manipulation of symbols at the expense of a larger memory footprint.

The symbol table supports scoping, *i.e.*, nested name spaces. Given a symbol table object, the methods **enterScope** and **exitScope** will enter a new scope and exit the current scope, respectively. In Bantam, a subclass of another class has access to all of the inherited classes symbols. Rather than copying these symbols into the subclass, the symbol table has a **setParent** method for setting the parent symbol table. This method effectively nests the child class's scopes inside the parent class's scopes.

The method **add** is used for adding a new name (a **String**) to the current scope. This method also takes a value (which can be any kind of object, *e.g.*, a **Location**), which can be retrieved later. For example, we might add the name "x" with value "int" (a **String**) to indicate that the variable "x" was declared with type "int". To lookup the value, we can use the method **lookup**, which takes as input a name. For example, we could call **lookup("x")**, which would return "int" (as an **Object**). If no entry is found with name "x", then **null** is returned. The method **lookup** looks for **"x"** in all scopes. To look in just the current scope we would use **peek**.

Abstract syntax tree. Except during optimization (an optional project), the Bantam compiler uses an abstract syntax tree [2, 22, 8, 19] (AST) as its intermediate representation. The AST is defined within the ast package. It has many different types of nodes, which are enumerated and discussed below.

ASTNode

An **ASTNode** is an abstract class that represents a generic node in the AST. It contains a line number indicating where the original source code (represented by the node) was scanned. All nodes in the AST extend **ASTNode**.

- **Program**
  Each (correct) Bantam program contains one top-level **Program** node. A **Program** node extends **ASTNode** and contains a list of classes (**ClassList**).

- **Class_**

A **Class** node represents a class from the Bantam source program. It extends **ASTNode** and contains a filename (**String**), a class name (**String**), a parent name (**String**), and a list of class members (**MemberList**).

- **Member**

  A **Member** node is an abstract class that represents a class member in a Bantam source program. It is extended by the following:

  - **Method**

    A **Method** node represents a method from the Bantam source program. It extends **Member** and contains a return type (**String**), a name (**String**), a list of parameters (**FormalList**), and a list of statements (**StmtList**).

  - **Field**

    A **Field** node represents a field from the Bantam source program. It extends **Member** and contains a declaration type (**String**), a name (**String**), and an (optional) initialization expression (**Expr**). For fields without initialization expressions, the initialization expression is **null**.

- **Stmt**

  A **Stmt** node is an abstract class that represents a statement in a Bantam source program. It is extended by the following:

  - **ExprStmt**

    An **ExprStmt** represents an expression statement from the Bantam source program. It extends **Stmt** and contains an expression (**Expr**).

  - **DeclStmt**

    A **DeclStmt** represents a declaration statement from the Bantam source program. It extends **Stmt** and contains a declaration type (**String**), a name (**String**), and a (non-optional) initialization expression (**Expr**).

  - **IfStmt**

    An **IfStmt** represents an if statement from the Bantam source program. It extends **Stmt** and contains a predicate expression (**Expr**), a then statement (**Stmt**), and an optional else statement (**Stmt**, which can be **null**).

73

- **WhileStmt**

  A **WhileStmt** represents a while statement from the Bantam source program. It extends **Stmt** and contains a continuation expression (**Expr**) and a body statement (**Stmt**).

- **ForStmt**

  A **ForStmt** represents a for statement from the Bantam source program. It extends **Stmt** and contains an optional initialization expression (**Expr**, which can be **null**), an optional continuation expression (**Expr**, which can be **null**), an optional update expression (**Expr**, which can be **null**), and a body statement (**Stmt**).

- **BreakStmt**

  A **BreakStmt** represents a break statement from the Bantam source program. It extends **Stmt** and contains no additional fields.

- **ReturnStmt**

  A **ReturnStmt** represents a return statement from the Bantam source program. It extends **Stmt** and contains an optional return expression (**Expr**, which can be **null**).

- **BlockStmt**

  A **BlockStmt** represents a block statement from the Bantam source program. It contains a list of expressions (**StmtList**).

- **Expr**

  An **Expr** node is an abstract class that represents an expression in a Bantam source program. It is extended by the following:

  - **AssignExpr**

    An **AssignExpr** represents an assignment expression (where the assigned variable is not an indexed array) from the Bantam source program. It extends **Expr** and contains an optional reference object (**String**, which can be **null**), a lefthand variable name (**String**), and a righthand expression (**Expr**).

  - **ArrayAssignExpr**

    An **ArrayAssignExpr** represents an array assignment expression from the Bantam source program. It extends **Expr** and contains an optional reference object (**String**, which can be **null**), a lefthand array variable name (**String**), an index expression (**Expr**), and a righthand expression (**Expr**).

74

- **DispatchExpr**

  A **DispatchExpr** represents a dynamic dispatch expression from the Bantam source program. It extends **Expr** and contains an optional reference object (**Expr**, which can be **null**), a method name (**String**), and a list of arguments (**ExprList**).

- **CastExpr**

  A **CastExpr** represents an explicit cast expression from the Bantam source program. It extends **Expr** and contains a target object type (**String**), an expression to cast (**Expr**), and a flag (**boolean**) indicating whether it is an upcast or a downcast.

- **InstanceofExpr**

  A **InstanceofExpr** represents an **instanceof** operation in a Bantam source program. It extends **Expr** and contains an expression to test (**Expr**) and an object type (**String**).

- **BinaryExpr**

  A **BinaryExpr** is an abstract class that represents all binary expressions where the left and right operands are both expressions (*i.e.*, not casts or **instanceof**). It extends **Expr** and contains a lefthand expression (**Expr**) and a righthand expression (**Expr**). It is extended by the following:

  - **BinaryArithExpr**

    A **BinaryArithExpr** is an abstract class that represents binary arithmetic expressions. It extends **BinaryExpr**. It is subclassed by the following:

    - **BinaryArithPlusExpr**

      A **BinaryArithPlusExpr** represents an addition ('**+**') expression.

    - **BinaryArithMinusExpr**

      A **BinaryArithMinusExpr** represents a subtraction ('**-**') expression.

    - **BinaryArithTimesExpr**

      A **BinaryArithTimesExpr** represents a multiplication ('*****') expression.

    - **BinaryArithDivideExpr**

      A **BinaryArithDivideExpr** represents a division ('**/**') expression.

    - **BinaryArithModulusExpr**

      A **BinaryArithModulusExpr** represents a modulus ('**%**') expression.

  - **BinaryCompExpr**

75

A **BinaryCompExpr** is an abstract class that represents binary comparison expressions. It extends **BinaryExpr**. It is subclassed by the following:

- **BinaryCompEqExpr**

  A **BinaryCompEqExpr** represents an equivalence ('**==**') expression.

- **BinaryCompNeExpr**

  A **BinaryCompNeExpr** represents a not equals ('**!=**') expression.

- **BinaryCompLtExpr**

  A **BinaryCompLtExpr** represents a less than ('$<$') expression.

- **BinaryCompLeqExpr**

  A **BinaryCompLeqExpr** represents a less than or equal to ('$<$**=**') expression.

- **BinaryCompGtExpr**

  A **BinaryCompGtExpr** represents a greater than ('$>$') expression.

- **BinaryCompGeqExpr**

  A **BinaryCompGeqExpr** represents a greater than or equal to ('$>$**=**') expression.

- **BinaryLogicExpr**

  A **BinaryLogicExpr** is an abstract class that represents binary boolean logic expressions. It extends **BinaryExpr**. It is subclassed by the following:

  - **BinaryLogicAndExpr**

    A **BinaryLogicAndExpr** represents a logical AND ('**&&**') expression.

  - **BinaryLogicOrExpr**

    A **BinaryLogicOrExpr** represents a logical OR ('| |') expression.

- **UnaryExpr**

  A **UnaryExpr** is an abstract class that represents all unary expressions where the operand is an expression (*i.e.*, not **new**). It extends **Expr** and contains an expression operand (**Expr**). It is extended by the following:

  - **UnaryNegExpr**

    A **UnaryNegExpr** represents a unary integer negation ('**-**') expression.

  - **UnaryNotExpr**

    A **UnaryNotExpr** represents a unary boolean complement ('**!**') expression.

  - **UnaryIncrExpr**

76

A **UnaryIncrExpr** represents a unary integer increment ('**++**') expression and contains a flag indicating whether it a prefix or postfix operation.

- **UnaryDecrExpr**

  A **UnaryDecrExpr** represents a unary integer decrement ('**--**') expression and contains a flag indicating whether it a prefix or postfix operation.

- **ConstExpr**

  A **ConstExpr** is an abstract class that represents all constant expressions. It extends **Expr** and contains a constant value (**String**). It is extended by the following:

  - **ConstIntExpr**

    A **ConstIntExpr** represents an integer constant (*e.g.*, **12**) expression.

  - **ConstBooleanExpr**

    A **ConstBooleanExpr** represent a boolean constant (*e.g.*, **true**, **false**) expression.

  - **ConstStringExpr**

    A **ConstStringExpr** represent a string constant (*e.g.*, **"abc"**) expression.

- **VarExpr**

  A **VarExpr** represents a variable reference expression. It extends **Expr** and contains an optional reference object (**Expr**, which might be **null**) and a variable name (**String**).

- **ArrayExpr**

  An **ArrayExpr** represents an array index expression. It extends **Expr** and contains an optional reference object (**Expr**, which might be **null**), a variable name (**String**), and an index expression (**Expr**).

- **ListNode**

  A **ListNode** is an abstract class that represents a generic list of nodes in the AST. All lists of nodes must extend **ListNode**. These include the following:

  - **ClassList**

    A **ClassList** node extends **ListNode** and contains a list of **Class_** nodes.

  - **MemberList**

    A **MemberList** node extends **ListNode** and contains a list of **Member** nodes.

- **StmtList**

  A **StmtList** node extends **ListNode** and contains a list of **Stmt** nodes.

- **ExprList**

  An **ExprList** node extends **ListNode** and contains a list of **Expr** nodes.

For more information on the API of each different class of AST node, view the API documentation off of the Bantam Java website at *http://www.bantamjava.com/*.

Visitor pattern. The last data structure is the Visitor design pattern [13]. The Visitor pattern allows us to traverse a tree (such as an AST) from a separate class. Before describing how it works, let's look at why we need it.

In order to implement the various phases of the compiler (in the upcoming sections), it will be necessary to traverse the AST several times. For example, to print the AST (if the compiler option "-sp" or "-ss" is set), we must walk over each AST node and print out a representation of it. In Java, there are essentially two ways we can do this traversal. We will first discuss a naïve approach and then discuss a better alternative. One approach is to add a method to each AST node that when called will print that particular node. The problem with this approach is that we have to add a method to every class in the ast package for printing that particular node. If printing was the only reason for traversing the AST then this solution might be acceptable. However, we will also need to traverse the AST to type check the program and to generate assembly code (and these might be further broken down into multiple traversals). Therefore, we would have to add a lot of code to each of the classes in the ast package (and there are more than 50 classes in the ast package!). Furthermore, in this approach, we are mixing functionality into the same classes. The AST classes should only contain code for representing a node in the abstract syntax tree. It should not contain code for printing the AST, or performing semantic analysis, or performing code generation. Imagine that we used this approach, and we wanted to make a small change to the semantic analyzer. This change could require making modifications to each class in the ast package!

A better alternative is to put the traversal code in a separate class. If we are traversing the AST in order to type check it, then that class would be a part of the semant package. But care needs to be taken in designing this traversal code. For example, many AST nodes contain sub-nodes whose exact dynamic type is unknown. An if statement has a predicate sub-node, which has static type **Expr**. This sub-node could actually be an **instanceof** expression, a **==** expression, or any other

type of expression. We could potentially have some large nested case statements that perform the appropriate actions based on the dynamic type of each sub-node. This would quickly become far too error-prone and nearly impossible to maintain.

Instead, we can use the Visitor pattern, which works as follows. We first build an abstract class called **Visitor**. Every class that needs to traverse the AST extends **Visitor**. The **Visitor** class defines a **visit** method for each type of AST node. For example, here is the **visit** method for **BinaryArithPlusExpr**:

```
public Object visit(BinaryArithPlusExpr node) {
    node.getLeftExpr().accept(this);
    node.getRightExpr().accept(this);
    return null;
}
```

This method gets the left and right expressions and calls an **accept** method on each. This **accept** method must be added to each class in the ast package, however, it only has to be done once, rather than each time we want to do a different traversal. The **accept** method works as follows:

```
public Object accept(Visitor v) {
    return v.visit(this);
}
```

By calling **visit** and then **accept**, we can vector back into the appropriate **visit** method for the sub-node. For example, if the left-hand expression of the **BinaryArithPlusExpr** were a **ConstIntExpr**, then after calling **visit** on the left-hand expression, and then **accept** in **ConstIntExpr**, we would end up in the method **visit(ConstIntExpr node)**.

Essentially, we are simulating multiple dispatch with single dispatch. We need to call a method based on both the runtime type of a node in the AST as well as a particular type of visitor. Since Java only supports single dispatch, we can achieve this by first dispatching into the particular AST node, and then dispatching back into the visitor object.

To build a new visitor class (for example, one that performs type checking), you will need to extend **Visitor**. You can then override each **visit** method with a method that traverses the AST and performs the additional functionality (*e.g.*, type checking). As stated earlier, you will need to do this when type checking the program during semantic analysis and when generating code

79

during code generation. In fact, each phase, *i.e.*, semantic analysis and code generation, may require multiple visitors. To view an example visitor, look at the **PrintVisitor** class in the visitor package, which traverses the AST and prints each node. Your visitors will look structurally similar to **PrintVisitor**.

4.3 Last Words

This section has given an overview of the Bantam Java compiler, which translates Bantam Java source code into assembly code. The student is responsible for implementing a working compiler, which must follow the language specifications layed out in the previous section and must use the data structures and code structure layed out in this section. In addition to these requirements, the assembly code generated by the student's compiler must interoperate with the runtime library code for handling memory allocation and deallocation, and built-in class methods, among other things. This is the topic of the next section.

5 Bantam Java Runtime System

This section describes the Bantam Java runtime system for the MIPS and x86 targets. If you are generating code for the Java Virtual Machine (JVM) then you can skip this section. For JVM the runtime library is very simple; it contains class files representing the **TextIO** and **Sys** built-in classes (all other functionality is already built in to the JVM itself). On the other hand, the runtime system for the MIPS and x86 targets are more complicated. They contain assembly code for performing various tasks such as garbage collection and error handling (among others). For MIPS and x86, in order to use the runtime system (*i.e.*, if you do not want to write your own runtime system), your compiler must use the same calling conventions and follow a specific format for encoding objects and primitives in memory.

The runtime system is provided with the Bantam Java compiler in the **lib/** directory within the main Bantam Java installation directory. If you are generating code for MIPS/SPIM then the runtime system is contained within the file **exceptions.s**. If you are generating code for x86/Linux then the runtime system is contained within the file **runtime.s**. Below we discuss each of the responsibilities of the runtime system, the calling conventions, and the format for encoding objects and primitives.

5.1 Runtime System Responsibilities

The runtime system contains assembly code for initiating the program, performing memory management, executing the built-in methods (*e.g.*, **Object.clone**, **String.concat**), and for error handling.

Program initiation. The compiled program begins executing in an initial subroutine within the runtime system (**__start** for Mips, **main** for x86). This subroutine does some initialization of runtime system data structures, and then calls the **main** method within the **Main** class, which is assembly code compiled by your compiler, *i.e.*, it's not part of the runtime system.

Memory management. The runtime system also contains several subroutines for allocating and deallocating memory. Deallocation is done via a simple, mark-and-sweep garbage collector, unless garbage collection is disabled in which case there is no deallocation. None of the methods have to be called by the compiled code, they are called automatically every time an object is created (below, we discuss how to create objects).

Built-in methods. The runtime system also contains the assembly code for every method of a

81

_null_pointer_error	Null pointer referenced	No additional arguments
_array_index_error	Array index out of bounds	Index passed in $t0 (Mips) or ebx (x86)
_array_size_error	Illegal array size	Size passed in $t0 (Mips) or ebx (x86)
_array_store_error	Array store error	Array type ID passed in $t0 (Mips) or ebx (x86),
		assigned object type ID passed in $t1 (Mips) or ecx (x86)
_class_cast_error	Class cast error	Object ID passed in $t0 (Mips) or ebx (x86),
		target ID passed in $t1 (Mips) or ecx (x86)
_divide_zero_error	Divide (or mod) by zero	No additional arguments

Table 5: Bantam Java runtime system error subroutines.

built-in class (methods within **Object**, **Sys**, **String**, and **TextIO**). Your compiler should not generate code for these methods.

Error subroutines. Finally, the runtime system contains subroutines for handling errors. Most error handling subroutines are called automatically by code within the runtime system. There are some that must be called by your compiled code. These are shown in Figure 5. For example, if the program attempts to use a **null** object (*e.g.*, dispatch on a **null** object), then the compiled code should call **_null_pointer_error**, which prints an error message and exits. If the program attempts to store a value into an array, whose type does not conform to the dynamic base type of the array, then the compiled code should call **_array_store_error**.

All of these subroutines take as input the current line number and a pointer to a string representing the filename, so that the line number and filename can be printed in an error message. For a Mips target, the line number is passed via the **$a1** register and the filename pointer is passed via the **$a2** register. For the x86 target, the line number is passed by storing it to the memory address **_current_line_number** and the filename pointer is passed by storing it to the memory address **_current_filename_ptr**. Both **_current_line_number** and **_current_filename_ptr** are defined within the runtime system (and are globals). The line number and filename must also be set when dispatching to a method or when constructing an object, since either of these operations could result in a call to an error subroutine.

In addition, some of the other subroutines take additional arguments. The error subroutines **_array_index_error** and **_array_size_error** take the illegal array index and size, respectively. For Mips, the index/size is passed via register **$t0**. For x86, the index/size is passed via register **ebx**. The error subroutine **_array_store_error** takes two additional arguments: a numeric identifier that indicates the type of the lefthand array in the assignment and a numeric identifier that indicates the type of the righthand object in the assignment. These identifiers are used to print out the types in

the error message. For Mips, the array identifier is passed via **$t0** and the righthand expression object identifier is passed via **$t1**. For x86, the array identifier is passed via **ebx** and the righthand expression object identifier is passed via **ecx**. The error subroutine **_class_cast_error** takes two additional arguments: a numeric identifier that indicates the type of object being casted and a numeric identifier that indicates the target type in the cast expression. These identifiers are also used to print out the types in the error message. For Mips, the object identifier is passed via register **$t0** and the target type identifier is passed via register **$t1**. For x86, the object identifier is passed via register **ebx** and the target type identifier is passed via register **ecx**.

5.2   Calling Conventions

The Bantam Java runtime system uses precise calling conventions that must be followed in order for the compiled program to work correctly. These conventions are discussed below. Note that these conventions only apply to Bantam Java methods. The calling conventions for error subroutines are discussed in the previous subsection.

Naming.   Methods that are compiled to assembly subroutines are named using the format: <**class**>.<**name**> where <**class**> is the name of the class containing the method and <**name**> is the name of the method. The assembly subroutine for each built-in method follows this convention. So, for example, if you want to call the **clone()** method in the **Object** class at the assembly level, you would call **Object.clone**. The runtime system also assumes that your compiler uses this convention, although the only method it will need to call is **main()** in the **Main** class. It will assume the corresponding compiled subroutine is named **Main.main**. This subroutine label must also be global (using the ".globl Main.main" directive) to allow the runtime system access from another assembly file.

Passing the reference object to the callee.   Because Bantam Java does not support 'static' all methods calls are dispatches using a reference object. The reference object is an implicit argument and must be passed to the called method. As in Java, Bantam Java supports only call by value. For objects, it is the reference to the object that is copied not the object itself. In Mips this address is passed in the register **$a0** and in x86 this address is passed in the accumulator register **eax**.

Passing parameters to the callee. All other arguments are passed via the stack. Starting from the leftmost argument and proceeding to the rightmost argument, each actual parameter is evaluated

83

Numeric identifier
Size in bytes
Dispatch table pointer
Value of field 1
Value of field 2
...
Value of field n

Figure 5: Bantam Java encoding of an object with **n** fields.

and the result is pushed onto the stack. To push it onto the stack, the value is stored to the address held in the stack pointer and the stack pointer is decremented by one word (x86 actually has an instruction for performing a push).

Passing the return value to the caller. The return value (if there is one) is passed via a register. In Mips it is passed via register **$v0** and in x86 it is passed via the accumulator register **eax**.

5.3   Encoding Objects and Primitives

To use the runtime system, objects and primitives must be encoded using the format discussed below.

Primitives. There are two types of primitives in Bantam Java: **int** and **boolean**. Nothing special needs to be done for **int** values. To generate a word (in assembly) with the value -10, the compiler would generate the following directive:

.word -10

Assembly language does not have support for **boolean** values, so **true** and **false** must be encoded using numeric values. **false** is encoded using 0 and **true** is encoded using a non-zero value. Probably the best non-zero value to use for true is -1 (we leave it to the reader to figure out why).

Objects. Figure 5 shows the encoding of an object. An object is encoded as a table with the following entries: (1) a numeric identifier that indicates the class or array type that the object was constructed from, (2) the size of the table in bytes, and (3) a pointer to the dispatch table (which are discussed below). The table also contains an entry per field defined (inherited or otherwise). If there are **n** fields then there are **n** words that follow the dispatch table pointer, each of which contains the value for the ith field. If a class has no fields then the dispatch table pointer is the last word in the object. In this case, the size of the object is 12 bytes. In general, for an object with **n**

```
.word 2          # String identifier
.word 24         # size of object in bytes
.word String_dispatch_table           .word 2          # String identifier
                 # pointer to String's     .word 24         # size of object in bytes
                 # dispatch table          .word String_dispatch_table
.word 4          # length of the string                     # pointer to String's
.byte 97         # ASCII encoding for 'a'                    # dispatch table
.byte 98         # ASCII encoding for 'b'   .word 4          # length of the string
.byte 99         # ASCII encoding for 'c'   .asciiz "abcd"   # ASCII encoded "abcd",
.byte 100        # ASCII encoding for 'd'                    # followed by null byte
.byte 0          # null terminating byte    .align 2         # word align by padding
.byte 0          # alignment byte                            # with null bytes
.byte 0          # alignment byte
.btye 0          # alignment byte
```

(a)   (b)

Figure 6: A string constant in Mips assembly code: (a) without the ".asciiz" and ".align" directives and (b) with the directives.

fields (inherited or otherwise), the size of the object is **12+n*4**.

Order within the list of fields is important. First, the order of fields defined within the same class should be the same as the order that they appeared in the class declaration. More importantly, fields inherited from classes higher in the class hierarchy tree (*i.e.*, closer to **Object**) should appear earlier in the list. If the program accesses a field in an object of type **X**, it should not matter whether the dynamic type of the object is **X** or some subtype of **X**; the field should reside at the same location in the object regardless of the dynamic type. By putting inherited fields earlier in the list, this property is guaranteed. Note: if an inherited field is redeclared, then it is treated as a separate variable, and so it has a separate entry in the list of fields. In this case, the static type of the object will determine which of the two fields is accessed at any point in the program.

Strings. Strings have a slightly different format than other objects. This is because unlike other objects, the size of string objects can vary. The size of the object depends on the length of the sequence of characters. Like other objects, a string object contains an identifier, size, and dispatch table pointer. The fourth entry is the length of the string (in number of characters). Following that is a byte holding each character encoded in ASCII, followed by a null terminating byte (*i.e.*, 0), and possibly followed by other null bytes so that the entire object is word-aligned.

Figure 6(a) shows the assembly code in Mips for encoding the string "abcd". If the ISA supports the directives ".asciiz" and ".align" (both Mips and x86 do), then we can use these to simplify

the assembly code as shown in Figure 6(b). The ".asciiz" directive automatically converts each character into the equivalent byte and terminates the string with a null byte. The ".align" directive automatically adds the necessary null bytes to word align the string object.

Arrays. Arrays are similar to **String** objects. Like **String** objects the size varies from array object to array object. Each array contains one field **length**, which holds the size of the array. The elements of the array follow the length field. The size of the object is then $16 + length * 4$ (note: unlike **String** objects each element takes a word in memory not a byte). Because the dispatch table of an array type is the same as **Object**, an array types does not need a separate dispatch table. Instead, each array type can use **Object**'s dispatch table.

Class name table. Your compiler will also need to generate a table that the runtime system uses for error reporting. Some error subroutines in the runtime system need to print class or array type names given an object (*e.g.*, on a class cast error). To find the name, it uses a special table called **class_name_table** (note: despite the name, this table is also for array types), which has an entry for each built-in and user-defined class or array type. The table is indexed by the object identifier. Each entry contains a pointer to a string representing the name of the class or array type. Your compiler will have to generate the **class_name_table** as well as the strings for each class name, including strings for built-in classes.

Object references and null. Objects are referred to using their memory address. An object reference (*e.g.*, a variable) is simply a pointer in memory (a word) that contains the memory address of the object that it refers to. If an object reference is **null**, then the pointer contains the value 0. Therefore, testing for **null** is achieved by comparing the reference to 0.

Dispatch tables. For each class, your compiler must generate a dispatch table, including for built-in classes. The dispatch table holds the addresses of every method that is visible from within that particular class. These should be named using the following format:

<center><class name>_dispatch_table</center>

In particular, the runtime system will need access to **String**'s dispatch table and will expect it to be named **String_dispatch_table**. In addition, the **String_dispatch_table** label will need to be made global (using ".globl") to allow access from the runtime system.

The dispatch table for each class contains memory addresses to all visible methods (inherited or otherwise). Each word in the dispatch table contains a memory address, specified as a label, to a particular method. As with the list of fields in an object, order within the dispatch table is

important. Methods from classes higher in the class hierarchy tree (*i.e.*, closer to **Object**) should appear earlier in the dispatch table. The order of methods defined within a particular class should be fixed for each dispatch table. In the absence of redefined methods, the order is the same as the order in which the method was defined within the class. If an extended class redefines an inherited method, then the address of the redefined method is placed in the entry that originally contained the inherited method. A new entry is not created. This last requirement is important for handling inheritance and object casting. If an object is casted to a super type, and a method is called that was overriden in the extended class, then this encoding ensures that the overriden method is the method that is called.

Object templates. Objects are not dynamically constructed from scratch. Instead, they are constructed from object templates. For each class, your compiler will generate an object template, including for built-in classes. These should be named using the following format:

<center><class name>_template</center>

In particular, the runtime system will need access to **String**'s template and **Main**'s template, and will expect these to be named **String_template** and **Main_template**, respectively. In addition, these labels will also need to be made global (using ".globl") to allow access from the runtime system.

Each template is encoded using the format shown in Figure 5. The templates contain the numeric identifier corresponding to that particular object type, the size of the object in bytes, the pointer to the dispatch table, and the default value for each field in the class (0 for **int**, **false** for **boolean**, **null** otherwise). When a new object is requested (using **new**), the template for the corresponding object type is copied using the built-in subroutine **Object.clone** provided in the runtime library. Therefore, your compiler does not have to explicitly manage memory allocation. Instead, your compiler can use templates and **Object.clone** to construct new objects.

Array construction is slightly more challenging than non-array object construction since the size of the constructed array can vary. However, by dynamically modifying the array template, we can use the same approach. Once the size of the array is known (which is not necessarily known until runtime), your compiler can compute the size of the array object in bytes and write this to the size entry in the array template. In addition, your compiler can write the length to the length field in the array template. The template can then be passed to **Object.clone**, which will copy the array.

In fact, with arrays, it is easier to use a single template for all array types. The pointer to the dispatch table is the same for all array types (**Object_dispatch_table**). The numeric identifier, which

<center>87</center>

Object_template:

0	
12	
Object_dispatch_table	→

Object_dispatch_table:

Object.clone

```
class Foo {
    int field1;
    boolean field2;
    void method1() { ... }
    void method2() { ... }
}

class Goo extends Foo {
    String field3;
    boolean field1;
    void method3() { ... }
    void method2() { ... }
}
```

(a)

Foo_template:

1
20
Foo_dispatch_table
0 (field1)
0 (field2 - false)

Foo_dispatch_table:

Object.clone
Foo.method1
Foo.method2

Goo_template:

2
28
Goo_dispatch_table
0 (field1)
0 (field2 - false)
0 (field3 - null)
0 (field1 - false)

Goo_dispatch_table:

| |
|---|
| Object.clone |
| Foo.method1 |
| Goo.method2 |
| Goo.method3 |

(b)

Figure 7: Some example templates and dispatch tables. (a) shows some Bantam Java code (which is incomplete) and (b) shows the templates and dispatch tables for three of the classes: Object, Foo, and Goo.

indicates the object type, does vary, however, it can be written during array construction. Since the size and length must be modified anyway, this adds very little additional overhead. This approach is used in the reference compiler.

Object initialization. Your compiler will also need to generate subroutines for initializing an object based on the initialization expressions of each field within the object. These should be named using the following format:

<class name>_init

In particular, the runtime system will need access to **Main**'s initialization subroutine in order to

88

create a **Main** object during program initialization. In addition, the **Main init** label will need to be made global (using ".globl").

After constructing an object (which is discussed above), the corresponding initialization subroutine is called. The initialization subroutine will execute each field's initialization expression and assign the resulting value to the corresponding field. Fields should be initialized in the same order in which they are defined within a class. Fields inherited from classes higher in the class hierarchy tree (*e.g.*, closer to **Object**) should be initialized earlier than those from classes lower in the tree. However, an initialization subroutine does not have to directly initialize inherited fields. Instead, it can call the initialization subroutine for the parent class (which can call the initialization subroutine of its parent class, and so forth).

Example. Figure 7(a) shows part of a (contrived) program. Figure 7(b) shows graphically three of the object templates and three of the dispatch tables that are generated by the compiler (note: other templates and dispatch tables would also be generated). Note that the method **method2** in class **Foo** is overriden in class **Goo**. It appears in the same place in the dispatch table in the **Goo** dispatch table as it does in the **Foo** dispatch table, however, in the **Goo** dispatch table it points to the method in **Goo** rather than in **Foo**. Note also that the field **field1** is redefined in class **Goo**. With fields, the redefined field does not replace the existing field. Instead a new entry is made for the redefined field. When executing code in class **Goo**, the second entry is used when **field1** is referenced. Alternatively, when executing code in class **Foo**, the first entry is used when **field1** is referenced.

5.4   Last Words

Over the past three sections, we looked at the Bantam Java language, compiler, and runtime system. If you are implementing an optimization component (an optional project), you will want to read the next section, which describes the Bantam Java optimization component. If not, then you are now ready to write your very own compiler. You can skip the next section and proceed to the following four sections where you will build the four components of a (non-optimizing) Bantam Java compiler: a lexer (Section 7), parser (Section 8), semantic analyzer (Section 9), and code generator (Section 10).

89

Figure 8: Bantam Java compiler phases including the optimizer.

6 Bantam Java Optimizer

This section describes the Bantam Java optimization component within the compiler. This section first gives an overview of the optimizer and then discusses the intermediate representation used by the optimizer. This section should be read after Sections 3-5. Moreover, the optimizer is an optional project so you should read this section only if you are doing the optimization project (described in Section 11).

6.1 Overview

As shown in Figure 8, the optimizer is a fifth phase in the Bantam Java compiler, which is positioned between the semantic analyzer and code generator. It takes as input the program represented as an abstract syntax tree (AST) and returns an optimized representation of the program. The program representation returned by the optimizer is a control flow graph [2, 22, 8, 19]. The Bantam Java optimizer focuses solely on reducing the overall execution time of the running program. Although commmercial compilers sometimes do other types of optimizations (*e.g.*, optimizing for code size or power), these are ignored in the Bantam Java compiler.

Optimization is performed on all user-defined Bantam Java methods as well as all initialization subroutines (see Section 5 for details on initialization subroutines). The built-in Bantam Java methods, error subroutines, and other code, which are a part of the runtime system are not optimized. In addition, the layout of the data section in the final generated assembly code is also not optimized.

Optimization is enabled and disabled on the command-line by the user (see Section 2 for details on the compiler command-line options). If enabled, the user sets the optimization level: 1-4 (0 for no optimization). Based on the specified optimization level, the optimizer performs the following tasks. If the optimization level is greater than or equal to 1, then the optimizer converts the AST into a control flow graph and sends this control flow graph (as opposed to the AST)on to the code generator. If the level is exactly 1, then no actual optimization is done, just the program representation conversion. If the level is greater than 1, then some high-level transformations are

done in the process of converting the AST into a control flow graph (*e.g.*, converting dynamic dispatches into direct calls when the called method is unambiguous). If the level is 2, then these are the only optimizations that are done. If the level is 3, then some low-level optimization is applied to the control flow graph (*i.e.*, common subexpression elimination [2, 22, 8, 19]). The low level optimization makes use of a data flow analysis. Finally, if the optimization level is 4, then any additional optimizations (high- or low-level) are applied to both the AST as it is converted to a control flow graph and to the control flow graph, itself.

The resulting control flow graph is passed to the code generator, which uses this representation to generate the final assembly code. The control flow graphs are passed to the code generator via the method symbol tables of each class. Each user-defined method is added to the method symbol table belonging to the class in which the method is defined. Likewise, the initialization subroutine is added to the method symbol table for the corresponding class. In both cases, the lookup key should be the final name used in the generated code (*e.g.*, **Main.main**, **Main_init**). See Section 5 on the runtime system for details on the naming of user-defined classes and initialization subroutines. The value added to the symbol table is simply the entrance node of the control flow graph.

6.2   Naming Scheme

Variables in the optimized representation must follow a specific naming scheme in order for the code generator to work correctly. In particular, the optimized representation does not contain scopes, therefore, it must be clear when a variable such as **foo** is used, the type of the variable: a local, a parameter, or a field. To distinguish between these different types of variables a special suffix is added to all variable names. For instance, the string "@l" (without the quotes) is added to all local variable names. If a variable **foo** is declared locally (within a method) then it is referred to as **foo@l** within the opimized representation. For formal parameters, the string "@p" (without the quotes) is added to the name. If **goo** is a formal parameter then it is referred to as **goo@p** in the optimized representation. For fields, the string "@f_<class name>" (without the quotes) where <class name> is replaced with the class name in which the field was defined. If **boo** is a field defined in the class **Main**, then it is referred to as **boo@f_Main** in the optimized representation.

In addition, the optimizer will need to generate some of its own temporary variables. These are named as "@t<num>" (without the quotes) where <num> is replaced with some non-negative integer. For temporary variables, there is no prefix in the name (since it did not come from a source

variable) just the suffix. For example, **@t0** and **@t15** are two different legal optimizer variables.

There is only one exception to this naming scheme: **this**. The special variable **this** does not require any suffix as it is unambiguous. On the other hand, the reference **super** cannot be used in the optimized representation. However, it is also not necessary. For example, assume the source program references a hidden field **x** using **super** from a parent class **Parent**. This code (**super.x**) gets replaced with **x@f_Parent**, which is still correct. If **super** is used to dispatch to an overriden method then this a direct call can be used to call the correct method in the parent.

Constants in the optimization representation are the same as in the source program. For example, **1** and **501** are legal **int** constants; **true** and **false** are legal **boolean** constants; and **"abc"** and **""** are legal **String** constants. In the optimized program representation, **null** is treated as a constant. It can be referred to simply as **null**.

6.3   Control Flow Graph

The optimizer converts the AST representation of the program (performing some high-level optimizations while doing this translation) into a control flow graph [2, 22, 8, 19]. This representation is contained in the package **cfg**. A control flow graph has nodes containing basic blocks and directed edges that indicate control transfers between blocks. A basic block is just a list of instructions with no interior control flow (only the last instruction can be a control instruction, control can only transfer to the first instruction). The instructions within the basic blocks are three-address code (3-AC) instructions. A 3-AC instruction is effectively a low-level but machine-independent instruction. For more details on control flow graphs, basic blocks, or 3-AC instructions see a compiler textbook [2, 22, 8, 19].

Basic blocks. Basic blocks in the Bantam Java compiler are represented by a class **BasicBlock** within the cfg package. To form a control flow graph, basic blocks are linked together. Each basic block has a set of out edges to successor blocks (*i.e.*, blocks that could immediately follow it in some execution of the program). Each block also has a set of in edges to predecessor blocks. The out and in edges can be retrieved, set, removed, *etc.* When an out edge in block **A** is set to block **B**, the corresponding in edge in **B** is also set to **A**. Likewise, when setting in edges. Therefore, you need only set an out edge from one block to another or the equivalent in edge but not both.

Each control flow graph has an *entrance* (or starting) basic block. This is the first basic block that is executed in the control flow graph. It should have no in edges. All other blocks should be

Class	Instruction	Example Use
If	Conditional branch	**if (@t0 < @t1) goto bb5**
Call	Indirect call	**@t1 = indircall @t0, 3**
	Direct call	**@t1 = dircall Main.foo, 3**
	Standard parameter	**stdparam @t0**
Parameter	Reference parameter	**refparam @t0**
	Error parameter	**errparam linenum 3**
Return	Return	**return @t0**
	Load variable	**@t0 = @t1**
Load	Load constant	**@t0 = 5**
	Load address	**@t0 = TextIO_init**
	Load entry	**@t0 = @t1[3]**
Store	Store entry	**@t0[3] = @t1**
Unary	Negation	**@t0 = - @t1**
	Not	**@t0 = ! @t1**
	Addition	**@t0 = @t1 + @t2**
	Subtraction	**@t0 = @t1 - @t2**
	Multiplication	**@t0 = @t1 * @t2**
Binary	Division	**@t0 = @t1 / @t2**
	Modulus	**@t0 = @t1 % @t2**
	AND	**@t0 = @t1 && @t2**
	OR	**@t0 = @t1 || @t2**

Table 6: Three-address code instructions.

reachable via the entrance block. As a result, the entrance block is used to represent the entire control flow graph. To add the control flow graph to the method symbol table, you should simply add the entrance block.

Each basic block contains a list of 3-AC instructions, which can be retrieved, set, removed, *etc.* Table 6 shows the various 3-AC instructions supported by the Bantam Java compiler. Each 3-AC instruction is defined in a separate Java class within the **cfg** package. Each of these classes is a subclass of the generic class **TACInst** (short for Three-Address Code Instruction). There are eight classes of instructions described below.

If instruction. First, there is an if instruction (**IfInst**) for performing conditional branches. The if instruction compares two operands and based on the comparison branches to one of two basic blocks. The comparison can be any one of the following: =, !=, <, <=, >, >=. More complicated predicates must be split into several instructions. In the example in Table 6, the instruction will branch to basic block **bb5** if the value in the temporary variable **@t0** is less than **@t1**, otherwise, it will fall through to the next basic block.

Call instructions. There are two types of call instructions: an indirect call (**InDirCallInst**) and a direct call (**DirCallInst**). Both classes representing each call extend a generic call class (**CallInst**), which contains the shared functionality. The indirect call can be used to call an address within a variable (like the call in a dynamic dispatch). The direct call can be used to jump to a specific target label. Both calls have an optional destination variable (**null** if unused) where the result of the call can be saved. Both calls also take the number of parameters required by the called method. Unlike in a source language dispatch, the reference object must be passed explicitly and it must be included in the parameter count. In fact, because Bantam Java has only instance methods, if at least one parameter is used, then this must include a reference parameter. In other words, all Bantam Java methods will have a parameter count of at least one. Only error subroutines (called when an error occurs such as a null pointer error) will have a 0 parameter count.

In the examples in Table 6, the indirect call instruction will call the method whose address is stored in variable **@t0**. The value returned from the called method will be written to variable **@t1**. The called method has 3 parameters (1 reference and 2 standard). The direct call instruction is similar to the indirect call instruction except that it directly calls method **Main.foo**.

Parameter instructions. There are also instructions for passing parameters to a called method. There are three types of parameter instructions: standard parameter (**StdParamInst**), reference parameter (**RefParamInst**), and special error subroutine parameters (**ErrParamInst**). All three classes representing each parameter instruction extend a generic class (**ParamInst**), which contains the shared functionality. Standard parameter instructions are used for passing conventional parameters (non-reference) to the called method. Reference parameter instructions are used for passing the reference parameter to a called method. In both instructions, the passed value must be either a variable (possibly a temporary variable generated by the optimizer) or a constant. Actual parameter expressions that are more complicated than this must be split into several instructions. In the examples in Table 6, the standard and reference parameters pass the values stored in variable **@t0** to the called method.

The legal ordering of standard and reference parameters is rigid. The reference parameter must come first followed by the standard parameters in the same order that they appear within the source program. Only one reference parameter can be used per called method. If any standard parameters are supplied, then a reference parameter must also be supplied (as all methods in Bantam Java are instance methods).

The error parameter instructions are for passing special error values to an error subroutine, which must be called when particular errors occur in the program, such as dispatching on a **null** object. These error subroutines are described in more detail in Section 5. Each subroutine needs to print out information about the error, such as the source line number where it occurred. This information is passed to the subroutine via the error parameters. There are six types of error parameters: filename, line number, object ID, target ID, array index, and array size. The first two are required in order to call any error subroutine. The others are needed for specific error subroutines. See Section 5 for more details. The passed value in the error subroutine must be a variable (possibly a temporary variable generated by the optimizer) or a constant. In the example in Table 6, the error parameter instruction passes line number **3** to the called error subroutine.

Return instruction. A return instruction (**ReturnInst**) is used for returning from a method. It contains an optional return value (**null** if unused). This value can be either a variable (possibly a temporary variable generated by the optimizer) or a constant. In the example in Table 6, the return instruction returns the value stored in variable **@t0** to the call site.

Load instructions. There are several load instructions for loading values into a variable (possibly a temporary variable generated by the optimizer). All of the classes representing these loads extend the class **LoadInst**, which contains the shared functionality. There is a load variable instruction (**LoadVarInst**) for copying the value from one variable to another. In the load variable example in Table 6 variable **@t0** is assigned to the value in variable **@t1**. There is also a load constant instruction (**LoadConstInst**) for loading a constant into a variable. In the load constant example in Table 6, variable **@t0** is assigned the value 5. Load constant instructions can also be used with **boolean** and **String** constants. There is also a load address instruction (**LoadAddrInst**) for loading the address of a label into a variable. In the load address example in Table 6, variable **@t0** is assigned the address of the label **TextIO_init**. Finally, there is a load entry instruction (**LoadEntryInst**) for loading an entry within a variable. For example, if the source variable is an object, this instruction can be used to access entries within the object, such as the dispatch table pointer (see Section 5 for details on the encoding of objects). This last load instruction has an additional argument: the index of the entry that you want to load. In the example in Table 6, variable **@t0** is assigned the value at index **3** within variable **@t1**. Note: **@t1** must refer to a structure that can be indexed such as an object. Note also: the index should not depend on the word size of the target machine.

Store instruction. There is one store instruction for setting an entry within a variable. This works

like the load entry instruction. In the example in Table 6, entry **3** within the structure referred to by variable **@t0** is assigned the value in variable **@t1**.

Unary instructions. There are two unary instructions for setting a variable (possibly a temporary variable generated by the optimizer) to the result of a unary operation. Both classes representing unary operations extend the class **UnaryInst**, which contains the shared functionality. There is a negation instruction (**UnaryNegInst**) for negating an **int** variable or constant, and putting the result in a variable. There is a not instruction (**UnaryNotInst**) for complementing a **boolean** variable or constant, and putting the result in a variable. In the examples in Table 6, the unary instructions perform the operation on the value in variable **@t1** and put the result in variable **@t0**.

Binary instructions. There are seven binary instructions for setting a variable (possibly a temporary variable generated by the optimizer) to the result of a binary operation. All of the classes representing binary operations extend the class **BinaryInst**, which contains the shared functionality. There are five instructions for performing arithmetic operations: addition (**BinaryAddInst**), subtraction (**BinarySubInst**), multiplication (**BinaryMulInst**), division (**BinaryDivInst**), and modulus (**BinaryModInst**). Each of these instructions performs the corresponding operation on two source operands, which can be either an **int** variable or **int** constant. There are also two instructions for performing boolean logic operations: AND (**BinaryAndInst**) and OR (**BinaryOrInst**). Both instructions performs the corresponding operation on two source operands, which can be either a **boolean** variable or **boolean** constant. In the examples in Table 6, the binary instructions perform the operation on the values in variable **@t1** and **@t2**, and put the result in variable **@t0**.

You may notice that there is no unconditional goto or jump instruction. These are handled implicitly by linking basic blocks. If a basic block has one out edge to another basic block and the code for the successor block does not immediately follow the former block, then an unconditional jump is automatically inserted. Therefore, you need only correctly link basic blocks. Furthermore, linking basic blocks must always be done explicitly. This is not always obvious. For instance, an if instruction (**IfInst**) contains the blocks that are branched to when the comparison condition is true and when it is false. Even still, the block containing the if instruction must still be linked to these succesor blocks. A block without the correct linking will result in an error if printed or passed to the code generator.

There are other correctness checks with regards to basic blocks and instructions. A basic block can have at most two out edges. If a block has two out edges, it must end with an if instruction. A

96

basic block with no out edges should end with either a return instruction or a call instruction to an error subroutine. In addition, an instruction cannot be added after an if instruction within a basic block since this instruction represents a control flow change. Every block must have at least one in edge except the starting block (otherwise, it would be unreachable).

For more information on the API of each different 3-AC class and basic block, view the API documentation off of the Bantam Java website at *http://www.bantamjava.com/*.

6.4 Last Words

This section has given an overview of the Bantam Java optimizer, which performs various transformations on the program representation in order to reduce the overall execution time of the source program. The student is responsible for implementing a working optimizer, which must follow the specifications layed out in this section. You are now ready to implement the Bantam Java optimizer, which is the project described in Section 11.

Figure 9: Bantam Java compiler phases. The lexer phase is shaded.

7 Project 1: Building a Lexer

In the first project, you will build a lexer using an automatic scanner generator. You will need to complete the code within the lexer package of your compiler source code. Make sure that you have read Section 3 on the Bantam Java programming language and Section 4 on the Bantam Java compiler before starting this project.

7.1 Overview

The lexer is the phase of the compiler that builds tokens from a stream of characters. Figure 9 shows the lexer (shaded) along with the other compiler phases. The lexer matches on all keywords, identifiers, constants, and special symbols (*e.g.*, ';') in the program and builds a token for each match that is passed on to the parser for syntactic analysis. For example, if the lexer scans the text " 99 " (without the quotes), then it would pass an integer token to the parser. The lexeme (*i.e.*, the matched text) of the token would be "99". The lexer also removes all spaces, tabs, newlines, and comments from the program, simplifying the work of later compiler phases. Finally, the lexer finds all lexical errors, errors in the lexemes. For example, string constants cannot span multiple lines; a string that spans multiple lines is a lexical error. See a compiler textbook for more information on lexers [2, 22, 8, 19].

7.2 Files and Directories

All of the source code you will need to edit is within the lexer package. The particular files that you will need to modify depend on whether you're using JLex or JavaCC to generate a lexer.

JLex. If you're using JLex then you will need to modify the file **lexer.jlex** (and only the file **lexer.jlex**). This file contains an incomplete specification of the lexer, which you will need to complete. See the JLex manual [7] for more information about building JLex specification files.

There is also a (working) file called **Token.java** that contains a class for representing tokens. The lexer will need to pass a **Token** object to the parser, for each recognized token. See the API

documentation for how to construct a **Token** object. It will require knowing the identifier of the token, which can be found in the class **TokenIds** in the parser class.

Tokens should still be sent to the parser even when errors are discovered. There is a special token identifier (**LEX_ERROR**) that indicates a lexical error occurred.

JavaCC. If you're using JavaCC then you will need to modify the file **lexer.jj** (and only the file **lexer.jj**). This file contains an incomplete specification of the lexer, which you will need to complete. See the JavaCC manual [6] for more information about building JavaCC specification files.

7.3 Lexer Details

Your lexer must be able to handle the following types of tokens:

- Special symbols (*e.g.*, '{', '}', ';', *etc.*)

- Keywords (*e.g.*, **class**, **while**, *etc.*)

- Identifiers (*e.g.*, **foo**, **Foo**, *etc.*)
  Identifiers are any non-keyword that starts with an uppercase or lowercase letter and is followed by a sequence of letters (upper or lowercase), digits, and '_'.

- Integer constants (*e.g.*, 9, 32, *etc.*)
  Integer constants are any sequence of digits (with possibly leading zeros) that is between 0 and $(2^{31} - 1)$. Note: other integer constant forms such as hexadecimal are not supported.

- Boolean constants (**true** or **false**)

- String constants (*e.g.*, "abc")
  String constants start and end with double quotes. They may contain the following special symbols: \n (newline), \t (tab), \" (double quote), \\ (backslash), and \f (form feed). A string constant cannot exceed 5000 characters and cannot span multiple lines.

Your lexer should remove all spaces, tabs, and newlines, as well as all comments. Bantam supports both multi-line comments (starting with /* and ending with */) and single line comments (starting with // and going until the end of the line).

Your lexer should identify the following lexical errors:

- Unterminated multi-line comments

- Unterminated string constants

- String constants spanning multiple lines

- String constants that are too long

- Unsupported escape characters within a string (*e.g.*, '\0')

- Integer constants that are too large

- Unsupported characters (*e.g.*, '?', '@') that are not within a comment or string constant

7.4   Testing the Lexer

You can test your lexer in one of two ways. First, you can supply the "-sl" flag to the Bantam compiler, which stops compilation after lexical analysis and prints the scanned tokens. You can compare the printed tokens with the printed tokens from the reference compiler. If they are different then this probably indicates an error in your lexer. A second way to test your lexer is to use the provided class files for performing syntactic and semantic analysis, and code generation, and fully compile some programs. If this produces incorrect assembly code, then this probably indicates an error in your lexer.

As part of the assignment you should submit a Bantam soure program that tests your lexer (**Test.btm**). It should contain as many kinds of lexical errors as you can code in a single program. It should also contain each type of token.

Important: take testing seriously in this project and all of the remaining projects. It is a critical part of the software design process. There is a strong correlation between students with working compilers and students who thoroughly test their code.

7.5   Last Words

Remember to start early on this project and all the remaining projects. It will take a significant amount of time (more than you probably think) and the sooner you finish the more time you can spend testing your lexer. Good luck and remember to have fun!

Figure 10: Bantam Java compiler phases. The parser phase is shaded.

8 Project 2: Building a Parser

In the second project, you will build a parser using an automatic parser generator. You will need to complete the code within the parser package of your compiler source code. Make sure that you have read Section 3 on the Bantam Java programming language and Section 4 on the Bantam Java compiler before starting this project.

8.1 Overview

The parser is the phase of the compiler that checks that the program is syntactically correct and, in addition, builds an internal representation of the program. Figure 10 shows the parser phase (shaded) along with the other compiler phases.

The parser uses the lexer to build tokens from streams of characters (see Section 7). It then checks that these tokens form a syntactically correct sentence, *i.e.*, a correct program. For example, if a Bantam program includes a class without an end brace ('}') the parser should catch this error, print an appropriate error message, and stop compilation after parsing is complete.

As the parser is checking correctness, it builds an intermediate representation of the program. In the Bantam Java compiler, the program is represented using an abstract syntax tree, or AST. An AST has a node for every type of construct in the language (*e.g.*, class, if statement, multiplication expression, int constant, *etc.*). These nodes may contain sub-nodes depending on the particular construct. For example, a class AST node contains a list of member definitions (methods and fields). See Section 4 for the details of the AST. If the program is syntactically correct then the AST representing the source program is passed on to later phases of the compiler (*e.g.*, semantic analyzer, code generator). For more information on parsers and ASTs see a compiler textbook [2, 22, 8, 19].

8.2 Files and Directories

All of the source code you will need to edit is within the parser package. The particular files that you will need to modify depend on whether you're using Java Cup or JavaCC to generate a parser.

Java Cup. If you're using Java Cup then you will need to modify the file **parser.cup** (and only the file **parser.cup**). This file contains an incomplete specification of the parser, which you will need to complete. See the Java Cup manual [14] for more information about building Java Cup specification files.

JavaCC. If you're using JavaCC then you will need to modify the file **parser.jj** (and only the file **parser.jj**). This file contains an incomplete specification of the parser, which you will need to complete. See the JavaCC manual [6] for more information about building JavaCC specification files.

8.3 Parser Details

Figure 11 shows a specification of the Bantam Java syntax. The specification is similar to a context-free grammar, which is described in detail in any compiler textbook [2, 22, 8, 19]. This specification shows how we can generate any syntactically-correct program starting from the non-terminal $Program$. Note that the terminals $identifier$, int_const, $boolean\_const$, and $string\_const$ are lexical tokens for identifiers (*e.g.*, variable names, class names), int constants (*e.g.*, 1, 29), boolean constants (*e.g.*, true, false), and string constants (*e.g.*, "abc", "def"), respectively.

Parser generators use context-free grammars to automatically construct a parser that can recognize syntactically-correct programs. You will write a specification for the context-free grammar in either the Java Cup or JavaCC format (depending on which parser generator you are using). You will then run the parser generator and build your Bantam parser by building the source (see Section 2 for details on building the source). Once you have a working specification, your parser should recognize any program that can be generated from this specification. Any program that cannot be generated from this specification is syntactically incorrect. For these programs, your parser should halt compilation and print out appropriate error messages.

There are a few differences between the specification in Figure 11 and a context-free grammar. First, $\lceil$ and $\rceil$ (not to be confused with '[' and ']') are used to denote an optional sequence of terminals and non-terminals within the righthand side of some production. Second, $*$ and $+$ as

$Program \rightarrow Class^+$

$Class \rightarrow class\ identifier\ [\ extends\ identifier\ ]\ \{\ Member^*\ \}$

$Member \rightarrow Method\ |\ Field$

$Method \rightarrow identifier\ [[\ ]]\ identifier\ (\ Formal^*_,\ )\ \{Stmt^*\ \}$

$Field \rightarrow identifier\ [\ [\ ]\ ]\ identifier\ [\ =\ Expr\ ]\ ;$

$Formal \rightarrow identifier\ [\ [\ ]\ ]\ identifier$

$Stmt \rightarrow ExprStmt\ |\ DeclStmt\ |\ IfStmt\ |\ WhileStmt\ |\ ForStmt\ |\ BreakStmt\ |\ ReturnStmt\ |\ BlockStmt$

$ExprStmt \rightarrow Expr\ ;$

$DeclStmt \rightarrow identifier\ [\ [\ ]\ ]\ identifier\ =\ Expr\ ;$

$IfStmt \rightarrow if\ (\ Expr\ )\ Stmt\ [\ else\ Stmt\ ]$

$WhileStmt \rightarrow while\ (\ Expr\ )\ Stmt$

$ForStmt \rightarrow for\ (\ [\ Expr\ ]\ ;\ [\ Expr\ ]\ ;\ [\ Expr\ ]\ )\ Stmt$

$BreakStmt \rightarrow break\ ;$

$ReturnStmt \rightarrow return\ [\ Expr\ ]\ ;$

$BlockStmt \rightarrow \{\ Stmt^*\ \}$

$Expr \rightarrow AssignExpr\ |\ DispatchExpr\ |\ NewExpr\ |\ InstanceofExpr\ |\ CastExpr\ |$
$\qquad BinaryExpr\ |\ UnaryExpr\ |\ ConstExpr\ |\ VarExpr\ |\ (\ Expr\ )$

$AssignExpr \rightarrow VarExpr\ =\ Expr$

$DispatchExpr \rightarrow [\ Expr\ .\ ]\ identifier\ (\ Expr^*_,\ )$

$NewExpr \rightarrow new\ identifier\ (\ )\ |\ new\ identifier\ [\ Expr\ ]$

$InstanceofExpr \rightarrow Expr\ instanceof\ identifier\ [\ [\ ]\ ]$

$CastExpr \rightarrow (\ identifier\ [\ [\ ]\ ]\ )\ (\ Expr\ )$

$BinaryExpr \rightarrow BinaryArithExpr\ |\ BinaryCompExpr\ |\ BinaryLogicExpr$

$UnaryExpr \rightarrow UnaryNegExpr\ |\ UnaryNotExpr\ |\ UnaryIncrExpr\ |\ UnaryDecrExpr$

$ConstExpr \rightarrow int\_const\ |\ boolean\_const\ |\ string\_const$

$BinaryArithExpr \rightarrow Expr\ +\ Expr\ |\ Expr\ -\ Expr\ |\ Expr\ *\ Expr\ |\ Expr\ /\ Expr\ |\ Expr\ \%\ Expr$

$BinaryCompExpr \rightarrow Expr\ ==\ Expr\ |\ Expr\ !=\ Expr\ |\ Expr\ <\ Expr\ |$
$\qquad\qquad Expr\ <=\ Expr\ |\ Expr\ >\ Expr\ |\ Expr\ >=\ Expr$

$BinaryLogicExpr \rightarrow Expr\ \&\&\ Expr\ |\ Expr\ ||\ Expr$

$UnaryNegExpr \rightarrow -\ Expr$

$UnaryNotExpr \rightarrow !\ Expr$

$UnaryIncrExpr \rightarrow ++\ Expr\ |\ Expr\ ++$

$UnaryDecrExpr \rightarrow --\ Expr\ |\ Expr\ --$

$VarExpr \rightarrow [\ Expr\ .\ ]\ identifier\ [\ [\ Expr\ ]\ ]$

Figure 11: Bantam Java syntax specification. Non-terminals start with capital letters and terminals start with lowercase letters or a non-alphabetic character. $Program$ is the starting non-terminal.

superscripts are used to denote zero or more occurrences or one or more occurrences, respectively. Finally, a comma as a subscript denotes comma-separated lists. A comma must appear between any two elements, but does not appear after the last element (if there are at least two elements – no comma is used for one or zero elements). You will need to figure out how to encode these into an unambiguous context-free grammar, which the parser generator will use to automatically construct the Bantam parser.

Your parser generator (Java Cup or JavaCC) also allows you to execute programmer-specified Java code, called *actions*, whenever a particular production is applied. You will use this feature to perform syntax-directed translation, in this case, translating the Bantam source program into an abstract syntax tree. See the Java Cup or JavaCC manual for the details on how to define actions.

8.4   Error Handling

Error handling is an important component of the compiler. The compiler must identify errors and print meaningful error messages so that the Bantam programmer can fix any errors. At a bare minimum, if your parser encounters a token that it was not expecting, it should print "Unexpected token" along with the filename and line number where the error occurred. For extra credit, identify specific errors such as a missing return statement, a nested class or method, or an unterminated class.

Your parser must halt compilation if at least one syntactic error is encountered, *i.e.*, the semantic analyzer and code generator should not run. For syntactic analysis, it is enough to stop parsing after one error is encountered (this will not be acceptable for semantic analysis). Extra credit will be awarded to parsers that can catch multiple syntactic errors during a single compilation. You will need to be able to recover from an error in order to implement this feature. See the Java Cup or JavaCC manual for details on implementing error recovery.

An **ErrorHandler** class is provided in the util package for printing and managing errors. See the API html documentation (click on class **ErrorHandler**) for details on using it.

8.5   Testing the Parser

You can test your parser in one of two ways. First, you can supply the "-sp" flag to the Bantam compiler, which stops compilation after syntactic analysis and prints the parsed program as a Bantam source program. You can compare the parsed program with the parsed program from

the reference compiler. If they are different then this probably indicates an error in your parser. A second way to test your parser is to use the provided class files for performing lexical and semantic analysis, and code generation, and fully compile some programs. If this produces incorrect assembly code, then this probably indicates an error in your parser.

As part of the assignment you should submit two Bantam soure programs that test your parser (**BadTest.btm** and **GoodTest.btm**). The first, **BadTest.btm**, should contain as many kinds of syntactic errors as you can code in a single program. The second, **GoodTest.btm**, should contain a working Bantam source program that tests all of the Bantam syntax rules.

Important: take testing seriously in this project. A parser is more complicated and thus more error-prone than a lexer, so be prepared to spend even more time debugging than in the previous project. Remember, testing is a critical part of the software design process. There is a strong correlation between students with working compilers and students who thoroughly test their code.

8.6 Last Words

Remember to start early on this project. It will take a significant amount of time (more time than the previous project) and the sooner you finish the more time you can spend testing your parser. Good luck and remember to have fun!

Figure 12: Bantam Java compiler phases. The semantic analysis phase is shaded.

9 Project 3: Building a Semantic Analyzer

In the third project, you will build a semantic analyzer. You will need to complete the code within the semant package of your compiler source code. Make sure that you have read Section 3 on the Bantam Java programming language and Section 4 on the Bantam Java compiler before starting this project.

9.1 Overview

The semantic analyzer is the phase of the compiler that checks that the program is semantically correct. Your semantic analyzer will also annotate the intermediate representation with type information needed by the code generator, and will build a data structure for representing the class hierarchy tree, which is important for both semantic analysis and code generation. Figure 12 shows the semantic analysis phase (shaded) along with the other compiler phases.

The semantic analyzer is passed the abstract syntax tree or AST (the intermediate representation) from the parser (see Section 8), assuming the program is syntactically correct. The semantic analyzer traverses the AST and checks that the program is semantically sound. For example, the semantic analyzer must verify that the operands in a multiplication expression both have type int (*i.e.*, it must *type check* the expression). If this is not the case, then the semantic analyzer should catch this error, print an appropriate error message, and stop compilation after semantic analysis is complete.

Therefore, in order to check that the program is semantically correct, your semantic analyzer will need to determine the type of each expression in the AST. The code generator also needs this information in order to generate correct assembly code. So it is the responsibility of the semantic analyzer to annotate each expression node with its discovered type.

Because Bantam Java supports inheritance, to check that the program is semantically correct, your semantic analyzer will need a data structure for representing the class hierarchy tree (see Section 3 for a description of the class hierarchy tree). For example, if an expression with type **A** (an object type) is used at a place in the Bantam source program where type **B** (also an object type)

is required, this code is legal only if **A** is a subclass of **B**. With a data structure for representing the class hierarchy tree, this property can be verified.

For more information on the theory of semantic analysis see a compiler textbook [2, 22, 8, 19].

9.2 Files and Directories

All of the source code you will need to edit is within the semant package. In particular, you will need to complete the file, SemanticAnalyzer.java, which contains some skeleton code for performing semantic analysis. You may also find it helpful to add other Java files to this package. For example, you might write a separate class (in a separate Java file) for traversing the AST and performing type checking (*i.e.*, checking that types are used correctly within the program).

In addition, you will find many of the files in the util and visitor packages useful, although you should not need to modify any of them. For example, in the util package, the file ClassTreeNode.java contains a class for representing nodes in the class hierarchy tree, the file SymbolTable.java contains a class for representing a symbol table, and the file ErrorHandler.java contains a class for managing and printing errors. In the visitor package, the file, Visitor.java contains an abstract class, which can be extended in order to traverse the AST using the visitor design pattern [13].

9.3 Semantic Analyzer Details

Your semantic analyzer will need to perform the following four steps:

1. Build a data structure representing the class hierarchy tree and check that it is well-formed (*e.g.*, there are no cycles, no repeated classes, *etc.*). You should use the class **ClassTreeNode** in the util package for representing each node in the tree. Each node contains some meta-information about the class (*e.g.*, a link to the parent and children nodes) and also the AST node for that particular class. The code generator takes the root of this tree (*i.e.*, the **Object** node) as input. Note: array types do not need their own class tree nodes (in fact, the reference compiler does not create nodes for array types). By using the base type of an array, you will be able to determine, for example, whether one type conforms to another.

2. Build class environments. In order to type check the program (in the final step) you will need to know what methods and fields are defined in a particular class. For example, if

107

your semantic analyzer encounters **x.foo()** during type checking, where **x** has type **X**, your semantic analyzer will need to look up **foo** in class **X** and verify that it exists and it is used correctly. So in this step (2), you will add all class members (inherited or otherwise) to the class's symbol table (using the symbol tables is discussed below). In addition, your semantic analyzer will need to check that methods and fields are defined correctly (*e.g.*, there are no methods or fields defined with the same name). Note: you will not need to explicitly build the environment for array types as every array type has the same environment: a method **clone** inherited from **Object** and a field **length**.

3. Check that the **Main** class is defined and contains a **main()** method, which has no parameters and no return value. Note: the **Main** class can extend another class and the **main()** method could be inherited.

4. Type check each field initialization expression and each method, annotating expressions as you type check. Anywhere a particular type is expected, you will need to verify that that type is used or a subtype is used (if the type is an object type). You will also annotate all expression nodes with their type, which is needed by the code generator. In addition, you will need to set a flag in cast and **instanceof** expressions, which indicates whether the cast/check is an upcast/upcheck (in which cast it is trivial to generate code for).

Read Section 3 closely for details on the semantics of the Bantam Java language. For instance, Section 3 specifies the correct operand types for each particular expression. It also specifies the resulting type for each particular expression. In addition, Section 3 specifies the rules for method calls and for redefining methods and variables in an extended class.

Data structures. Your semantic analyzer will need to build and use two important data structures: a class hierarchy tree and a symbol table. Both of these are provided in the util package.

The class **ClassTreeNode** is used for representing a node in the class hierarchy tree. After (or while) checking that the class hierarchy tree is well formed, your semantic analyzer should then build a representation of the tree using **ClassTreeNode**. Each node in this tree contains all the information pertinent to that corresponding class including the name, the links to the parent and children class tree nodes, and a link to the AST node representing that class. The root of this tree, which is the node representing the built-in class **Object**, is passed to the code generator. See Section 4.2 or the API for more information on how to use this class.

The class **SymbolTable** is used for representing a symbol table. Each class tree node has two symbol tables: one for variables and one for methods. Variables and methods can be added to these symbol tables in order to later type check the program. For example, during type checking, if your semantic analyzer encounters an assignment to a variable **x**, it can look up **x** in the variable symbol table to make sure that it exists and that the type matches the type on the righthand side of the assignment. A similar strategy can be used when your semantic analyzer encounters a method call.

These symbol tables support inheritance; each symbol table has a link to the corresponding symbol table in the parent class. When looking up a variable or method, first the current class's symbol table is searched, and then if the symbol is not found, the parent class's symbol table is searched. This process continues until the top of the class hierarchy tree is reached.

The symbol table also supports nested scopes. For example, a block statement defines a new nested name space (nested within the current name space). When searching for symbols in the symbol table, your semantic analyzer can search either in the current scope (called a *peek*) or in any scope (called a *lookup*). See Section 4.2 or the API for more information on how to use this class.

Visitor pattern. Your semantic analyzer will need to traverse the abstract syntax tree (AST) in order to perform semantic analysis. While you should not walk the AST needlessly, do not feel as though you need to optimize the total number of traversals (perhaps, by combining multiple traversals into one). Such optimizations will not reduce the compile time by all that much in most cases, and the performance of the compiler is generally not critical (as opposed to the performance of the compiled code). One approach might be to traverse the AST once and build the symbol tables for each class (bullet two in the four steps described above) and a second time to type check each class (bullet four in the four steps described above).

As described in Section 4.2, the best approach for traversing an AST from a software engineering perspective is to use the visitor design pattern. This approach allows us to keep the AST functionality separate from the semantic analysis functionality. It also makes it easier to modify and maintain the semantic analyzer. A change in the semantic analyzer has no impact on the AST, which is important since there are more than 50 classes representing each particular type of AST node. To write a visitor class, you will write a class that extends the abstract class **Visitor**. This class will override each **visit** method with a method that traverses the AST and performs the ad-

109

ditional functionality (*e.g.*, type checking). For an example, look at the **PrintVisitor** class in the visitor package, which traverses the AST and prints each node. For a more detailed description of the visitor design pattern, see Section 4.2.

9.4   Error Handling

Like in the parser, error handling is an important component of the semantic analyzer. The compiler must identify errors and print meaningful error messages so that the Bantam programmer can fix any errors. Unlike with parsing, error identification in the semantic analyzer should be more precise and more thorough. Your semantic analyzer should find all semantic errors within the program, assuming that one error does not hide another error. For example, imagine a programmer attempts to call a method **foo**, which takes an **int** as a parameter, and mistakenly mistypes the method name and forgets to pass any parameter in the call. Only the first error can be discovered in this case, the second error will only be discovered after the programmer fixes the first error. But excluding errors hidden by other errors, your semantic analyzer should find all semantic errors within the program.

For each discovered error, you should print a meaningful error message and also print the file-name and line number where the error occurred. For semantic analysis, unlike syntactic analysis, it is usually much easier to give a precise error message. Run the reference compiler (discussed underneath the heading "Testing the Semantic Analyzer") to see some examples of good error messages.

An **ErrorHandler** class is provided in the util package for printing and managing errors. See the API documentation for details on using **ErrorHandler**.

9.5   Testing the Semantic Analyzer

You can test your semantic analyzer in one of two ways. First, you can supply the "-ss" flag to the Bantam compiler, which stops compilation after semantic analysis and prints the intermediate representation as a Bantam source program. Annotations (*i.e.*, expression types) are provided within comments in the source program. You can compare this output with the corresponding output from the reference compiler. If the output is different then this probably indicates an error in your semantic analyzer. A second way to test your semantic analyzer is to use the provided class files for performing lexical and syntactic analysis, and code generation, and fully compile some

programs. If this produces incorrect assembly code, then this probably indicates an error in your semantic analyzer.

Important: take testing seriously in this project. A semantic analyzer is more complicated and thus more error-prone than a parser or a lexer, so be prepared to spend even more time debugging than in the previous two projects. Remember, testing is a critical part of the software design process. There is a strong correlation between students with working compilers and students who thoroughly test their code.

As part of the assignment you should submit two Bantam soure programs that test your semantic analyzer (**BadTest.btm** and **GoodTest.btm**). The first, **BadTest.btm**, should contain as many kinds of semantic errors as you can code in a single program. The second, **GoodTest.btm**, should contain a working Bantam source program that tests all (or most) of the Bantam semantic rules described in Section 3.

9.6 Last Words

Remember to start early on this project. It will take a significant amount of time (more time than the previous two projects) and the sooner you finish the more time you can spend testing your semantic analyzer. Good luck and remember to have fun!

Figure 13: Bantam Java compiler phases. The code generation phase is shaded.

10   Project 4: Building a Code Generator

In the fourth and final core project, you will build a code generator. You will need to complete the code in one of the three code generator packages: codegenmips (if targeting MIPS), codegenx86 (if targeting x86), or codegenjvm (if targeting the JVM). Before starting this project, make sure that you have read Section 3 on the Bantam Java programming language, Section 4 on the Bantam Java compiler, and most importantly, Section 5 on the Bantam Java runtime system.

10.1   Overview

The code generator is the final phase of the compiler that takes as input the internal representation of the program and produces the target code (either assembly code or text-based java bytecode, depending on the target). Figure 13 shows the code generation phase (shaded) along with the other compiler phases.

The code generator is passed the root node of the class hierarchy tree from the semantic analyzer. From the root node, we can of course reach any other class tree node. Each node contains meta-information about the class as well as a link to the AST node representing that particular class. Any expression AST nodes embedded within the class AST node is annotated with its type, which was done by the semantic analyzer. The code generator can assume that the program it is passed is a legal, Bantam program. The code generator traverses the class hierarchy tree and AST nodes (perhaps, multiple times) and generates target code.

For more information on the theory of code generation see a compiler textbook [2, 22, 8, 19].

10.2   Files and Directories

All of the source code you will need to edit is within the codegenmips package, codegenx86 package, or codegenjvm package (depending on the particular target you are using). In particular, you will need to complete the file, **MipsCodeGenerator.java**, **X86CodeGenerator.java**, **JVMCodeGenerator.java**, (again, depending on the target), which contains some skeleton code for performing code

112

generation. You may also find it helpful to add other Java files to this package. For example, you might write a separate class (in a separate Java file) for traversing each AST node and generating target code.

In addition, you will find many of the files in the util and visitor packages useful, although you should not need to modify any of them. For example, in the util package, the file **ClassTreeNode.java** contains a class for representing nodes in the class hierarchy tree and the file **SymbolTable.java** contains a class for representing a symbol table. In the visitor package, the file, **Visitor.java** contains an abstract class, which can be extended in order to traverse the AST using the visitor design pattern [13].

10.3   MIPS/x86 Code Generator Details

If you are targeting MIPS or x86 then read this subsection for details on implementing your code generator. If you are targeting the JVM then skip this subsection and proceed to the next subsection.

The code generator will need to work closely with the provided runtime system, which is discussed in Section 5. In particular, the code generator will need to use the same calling conventions as the runtime system and to encode data structures in memory as the runtime system expects. Many of the details are discussed in Section 5.

Generated data/code. Your code generator will need to generate several different data structures and subroutines. The data structures must appear within the data section, *i.e.*, they must follow the **.data** directive. The subroutines must appear within the text section, *i.e.*, they must follow the **.text** directive. A few of these are already implemented within the compiler, but most are not. In particular, your code generator will need to generate the following:

- **Globals.** Several of the data structures and subroutines discussed below must be made available to the runtime system. The data structures that must be made global include **gc_flag**, **class_name_table**, **Main_template**, **String_template**, and **String_dispatch_table**. The subroutines that must be made global include **Main_init** and **Main.main**.

- **Garbage collection.** Your code generator will need to set a flag called **gc_flag** to 1 or 0 depending on whether garbage collection is enabled or disabled, respectively. For example, if garbage collection is disabled, then the following code should appear in the data section:

113

```
gc_flag:      .word 0
```

- **String constants.** Your code generator will need to scan the program to find all string constants (*e.g.*, "abc"). For each string constant, you should generate a string object within the data section of the assembly program. Give each string constant a unique name. For example, you could use **String_const_<id>** to name each string constant where **<id>** is replaced with a unique numeric identifier. When that string constant is used in the program, you can use the address of the string object (which in assembly is simply the name). You will probably also need a data structure for mapping a string constant (*e.g.*, "abc") to its string object reference name (*e.g.*, **String_const_5**). As discussed in Section 5, you will also need to generate string objects for each built-in and user-defined class name as well as each program filename for use in error reporting.

- **Class name table.** Your code generator will need to build a class name table (called **class_name_table**) that contains pointers to the strings representing the class names. This table is used by the runtime system for error reporting when a runtime error occurs. It is ordered by class identifier. For example, if the class **Object** has identifier 0, then a pointer to a string representing "Object" should be contained in the first entry.

- **Object templates.** Your code generator will need to generate object templates (*e.g.*, **TextIO_template**) for each built-in and user-defined class. These are used to construct new objects, *i.e.*, they are cloned during object construction. See Section 5 for details on how to construct these templates.

- **Dispatch tables.** Your code generator will need to generate dispatch tables for each built-in and user-defined class (*e.g.*, **String_dispatch_table**). These are needed for implementing dynamic dispatch. See Section 5 for details on how to construct these tables.

- **Initialization subroutines.** Your code generator will need to generate initialization subroutines for each built-in and user-defined class (*e.g.*, **Main_init**). These are called to initialize each field in an object after an object is constructed. See Section 5 for details on how to construct these subroutines.

114

- **User-defined methods.** Finally, your code generator will need to generate code for each user-defined method (*e.g.*, **Main.main**). It will not need to generate code for built-in methods since these are defined within the runtime system.

  These methods should use the calling conventions discussed in Section 5.

Traversing the AST. Your code generator will probably need to perform several traversals over the internal representation of the program. For example, you could traverse the program once to find string constants, a second time to generate initialization subroutines, and a third time to generate user-defined methods. Create visitor classes to traverse the program as in project 3.

It is much easier to implement the code generator if you can generate the code for each node independently of where it is positioned in the AST. To do this you will have to set up some conventions between AST nodes. For example, if one expression is embedded within a larger expression (*e.g.*, a variable within an multiplication expression) where can the encompassing expression (*e.g.*, multiplication expression) expect the computed value from the inner expression (*e.g.*, variable)? You should think carefully about this issue and related issues.

Storing values. Another issue you will need to think about is where to store variables and other temporary values. Each variable or temporary value can either be stored in an available register or on the stack in memory. The easiest approach, and the one used by the reference compiler, is to put all variables and temporary values on the stack until they are needed in some computation (*e.g.*, a multiplication operation). When they are needed, they can be retrieved from memory, put in a temporary register, and used in the computation. This approach avoids the complexities of register allocation. For extra credit (see below), you can look at implementing a more efficient approach.

Regardless of where variables are stored, your code generator will need to keep track of their location. The symbol table, which you used in project 3 to keep track of types, can also be used to keep track of the location of each variable. To represent the location, you can use the **Location** class defined within the **util** package. **Location** can represent a location in memory or in a register. See the API for details.

Testing conformance. One important task for your code generator is to generate code to test whether an object conforms to a particular type. This problem arises when the program executes an **instanceof** expression and when the program performs a down cast (to identify a cast error). In such cases, your generated code will need to use the numeric identifier that indicates the class that

115

an object was constructed from (see Section 5). This identifier can be compared with the identifier of the target type in either the **instanceof** expression or the cast expression.

Unfortunately, this will not work in all cases. For example, the object's dynamic type might be a strict subtype of the target type, in which case the identifiers will be different even though the object conforms to the target type. One solution is to compare the object's identifier with the identifiers of all conforming types. But the cost of this check would depend on the number of conforming types and could be quite high in general. We can avoid this high-cost check if we use care in assigning numeric identifiers to object types. In fact, this check can be done in constant time. We leave this problem as an exercise to the student.

Your generated code will also need to test conformance when assigning to an array element (to identify an array store error). This test is more difficult than in an **instanceof** or cast expression as the target type is not known statically. Instead, you will need to find the numeric identifier of the base type of the array given the numeric identifier of the array type. Again, if you use care with assigning numeric identifiers, then this can be done in constant time. Your generated code will then need to determine if the righthand value in the assignment conforms to this particular base type. This check will require at least one additional data structure. We also leave this problem as an exercise to the student.

Error handling. Although the program was checked for lexical, syntactical, and semantic errors in previous phases, there are several errors that can still occur at runtime. Your code generator will also need to generate code to catch these runtime errors. These include:

- **Null pointer error.** This error occurs when the program attempts to access a member of a **null** object.

- **Array index error.** This error occurs when the program attempts to index an array using an index that is out of bounds (*i.e.*, < 0 or $>=$length).

- **Array size error.** This error occurs when the program attempts to construct an array with an invalid size. If garbage collection is disabled, then the size must be $>= 0$. Otherwise, the size must be $>= 0$ and $<= 1500$.

- **Array store error.** This error occurs when the program attempts to assign an array element a value that does not conform to the dynamic base type of the array.

116

- **Class cast error.** This error occurs when the program attempts to cast an object to a type and the object's type does not conform to the target type.

- **Divide by zero.** This error occurs when the program performs a divide ('/') or modulus ('%') with a divisor of 0.

If one of the errors above occurs within the program, then the code generated from your compiler should call the corresponding error subroutine in the runtime system. See Section 5 for more details on how to call each of these subroutines.

10.4 JVM Code Generator Details

If you are targeting the JVM then read this subsection for details on implementing your code generator. If you are targeting MIPS or x86 then skip this subsection and proceed to the next subsection.

As mentioned above, the JVM code generator will produce several Jasmin input files [18, 17], which are simply files containing Java assembly code (text-based Java bytecode). See the Jasmin documentation [18, 17] for more details on the format of this code. The Jasmin tool is then used to convert each input file into a Java class file, which can then be run via the Java Virtual Machine.

Generated classes. Your code generator will need to create a Jasmin assembly file for each user-defined class in the program (built-in classes are defined in the runtime library). Each Jasmin file must be named using the following format (otherwise, Jasmin will not properly assemble it):

<class name>.j

where <**class name**> is replaced with the actual name of the class (*e.g.*, **Main**). Each file will contain the Java assembly code pertaining to that particular class. In particular, it will contain code defining all the fields and methods that were specified in that particular class.

Your code generator must add one additional construct to the **Main** class: a static **main** method. This method must be added to **Main** since Java (unlike Bantam Java) requires a static **main** method as the entrance point to the program. The static **main** method will simply create a **Main** object and call the non-static **main** method.

Generated code within classes. Your code generator will need to translate each class in the source program into the corresponding Jasmin directives, instructions, and labels. Directives can

117

be used to declare the class as well as the methods, fields, and local variables defined within the class (you may also find other uses for directives). Source-level Bantam Java statements will need to be translated into Jasmin instructions. Labels will be required when translating Bantam Java control flow such if statements and while loops. See the Jasmin documentation [18, 17] for more information on how to formulate Jasmin directives, instructions, and labels.

The JVM [16] is a stack-based architecture, meaning that it uses a stack rather than registers to hold local data values retrieved from memory as well as intermediate values computed at runtime. So most of the generated instructions will be manipulating this stack. For example, there are instructions for pushing variables or other values onto the stack, instructions for popping values off the stack, and instructions for performing operations on the top element(s) of the stack (*e.g.*, adding the top two elements).

Although the JVM has some built-in support for manipulating integers and objects, you will find that in some cases, you will need to use several Jasmin instructions when translating certain Bantam Java language features (*e.g.*, dynamic dispatches). In addition, the JVM does not have built-in support for manipulating booleans so you will need to encode and operate on booleans as integers.

Traversing the AST. Your code generator will probably need to perform one or more traversals over the internal representation of the program. At the very least, you will need one visitor class to traverse the program and translate each element into Java assembly code. When visiting expressions, which can be embedded within larger expressions, the node representing the expression should push its result on the stack. This result can then be popped off of the stack by the node representing an encompassing expression.

Generating verifiable code. To run on the JVM, your generated code must not only satisfy the constraints specified by the Java language (in our case, the Bantam Java language), it must also successfully go through the verification process. Every time one invokes the JVM (*e.g.*, by typing the **java** command), the JVM loads the necessary class files, verifies them, links them together, and finally runs them. The verification step is crucial to ensuring safe execution. Without going into detail about the algorithm ([16], Section 4.9), the JVM makes sure that each class (except "Object") has a direct superclass, that methods are invoked with the correct number of arguments and that these are appropriately typed, that a local variable is not used unless it has verifiably been assigned a value of the correct type, *etc.* Of course, the code that your compiler generates

118

must satisfy these constraints. However, there is one additional constraint that is not so obvious, namely: "*The verifier ensures that at any given point in the program, no matter what code path is taken to reach that point, [. . .] the stack is always the same size*" ([16], Section 4.9.1). Recall that the stack contains temporary values needed during execution. For example, when executing the second statement in the compiled code for the following Bantam Java source program fragment:

```
TextIO io = new TextIO();
io.putString("hello");
```

the JVM will first push a reference to the object "io" onto the stack. This reference will then be used and popped off the stack by the subsequent evaluation of the "putString" method call. Once the "putString" method terminates, its return value (namely, a reference to the "io" object) is pushed onto the stack. This is needed since calls to "putString" can be chained, and the next call (of which there is none in this example) would need this reference for its method dispatch. In conclusion, after this code fragment is executed, there is (at least) one extra object reference remaining on top of the stack, namely, the reference to "io" that is returned by the last call to "putString" in the call chain. This is not always a problem, since the JVM automatically pops elements left unused on top of the stack at the end of the execution of each method body.

Unfortunately, these extra stack elements sometimes violate the rule stated above. We now need to explain one central phrase in that rule: "*no matter what code path is taken.*" The compiled code for some Bantam Java language constructs will contain "goto" instructions. For example, an "if-then-else" construct embodies a decision that translates into branching code. (Can you think of other Bantam Java language constructs that will also translate into branching code?) Now, the rule implies that, whichever one of the "then" or "else" branches is taken, the stack must contain the same number of elements before executing the statement following the "if." But if the code fragment above is part of the "then" branch (say), while the "else" branch is empty (say), the stack size will be different, and larger by (at least) one, when the "then" branch is taken. This program will not pass the verification process and will be rejected by the JVM even before it runs the program.[1]

To surmount this obstacle, you need to make sure and eliminate all extraneous stack elements. It is probably easiest to delete (*i.e.*, pop) them as soon as they are pushed. Therefore, you need to

---

[1]Can you figure out why having extra stack elements lying around on the stack is a security risk?

come up with a complete list of all statements that can potentially leave unused elements on the stack. Hint: Study the list of expression statements in Bantam Java.

Optimizing memory use. When a method is invoked, the JVM starts by allocating memory for an activation record (or stack frame) that contains the stack space and memory for the local variables and parameters of the method. Therefore, the JVM must know in advance how much memory to allocate for the stack and for the local variables. Your compiler must compute these two numbers and include them in each method of the Jasmin file(s) it outputs (see the ".limit stack" and ".limit locals" directives described in [18]).

To determine the stack size, your compiler must compute the maximum height reached by the stack during the execution of the method, which can be easily achieved given that the net effect on the stack height of each instruction that your compiler outputs is known ([16], section 6).

To determine the number of slots needed for the parameters and local variables of the method, you need to keep in mind that:

1. each method in Bantam Java is non-static, and thus takes as an implicit first parameter the reference to the calling object or "this." This reference must be allocated memory, just like for any other method parameter,

2. some local variables can share the same memory locations, namely when it is known that they are never active at the same time. For example, variables declared in a "then" branch of an "if" statement can occupy the same memory space as the variables in the "else" branch of that "if" statement. Can you think of other Bantam Java language constructs where similar reuse of memory space is possible?

In conclusion, the number of slots allocated by the JVM should be large enough to contain all of the local variables (including parameters) that can be active at any given point in the method body. In fact, your compiler should compute the *exact* number of memory locations needed, since allocating extra memory would be wasteful. For example, allocating a constant, arbitrarily large number of memory locations, say 100, would be extremely wasteful, since most (arguably, all) well-designed methods use fewer local variables. Furthermore, doing so would be incorrect, since your compiler may one day be fed a program that does use more than 100 local variables. After all, your compiler should not impose arbitrary constraints that are not part of the source language specification.

10.5   Extra credit

For extra credit, add some optimizations to improve the performance of the generated code. For example, you could perform some data flow analysis (*e.g.*, available expressions) and use this analysis to perform an appropriate optimization (*e.g.*, common-subexpression elimination). A second (easier) possibility is to convert dynamic dispatches into direct calls where ever possible. A third possibility is to perform register allocation and minimize the number of times values have to be put on the stack. The amount of extra credit will vary greatly based on the amount of work that you do.

Important: under no circumstances should you attempt the extra credit until your code generator is working! Building a working code generator will be very challenging without worrying about optimization.

10.6   Testing the Code Generator

Unlike the previous phases, there is only one way to test your code generator. You can use the provided classes for performing lexical, syntactic, and semantic analysis, and fully compile some programs. You can then try comparing the generated program to the generated program from the reference compiler. However, the output from your code generator could be different from the reference output, and still be correct. If the two generated programs are different, then you can run them and see if the programs have the same behavior. If they do not have the same behavior, then this probably indicates an error in your code generator.

It might be worth using the same format for your generated program as the reference compiler. If you are able to achieve this, then you will be able to simply compare the two generated files, which is much easier than comparing the behavior of the programs when they are run. However, if you find this too difficult to do, then you may want to simply compare the compiled program's behavior when executed.

As part of the assignment you should submit a Bantam source program that tests your code generator (**Test.btm**). The file should contain a working Bantam source program that uses as much of the Bantam language as you can code within a single program.

Important: take testing seriously in this project. A code generator is by far the most complicated phase of the compiler, so be prepared to spend even more time debugging than in the previous projects. Remember, testing is a critical part of the software design process. There is a

strong correlation between students with working compilers and students who thoroughly test their code.

10.7   Last Words

Remember to start early on this project. It will take a significant amount of time (more time than the previous projects) and the sooner you finish the more time you can spend testing your code generator. Good luck and remember to have fun!

Figure 14: Bantam Java compiler phases. The optimizer phase is shaded.

11 (Optional) Project 5: Building an Optimizer

In this optional project, you will build an optimizer. You will need to complete the code within the opt package of your compiler source code. Before starting this project, make sure that you have read Section 6, which discusses the conventions and program representation used by the optimizer.

Note: this assignment is currently not supported with the JVM target. We hope to add this support in the near future.

11.1 Overview

Although optimization is often done throughout many phases of a commercial compiler, optimization in the Bantam Java compiler is restricted to a single phase of compilation. Figure 14 shows the optimizer (shaded) along with the other compiler phases. The optimizer takes as input the abstract syntax tree representation of the program and produces a set of optimized control flow graphs (one for each method). The Bantam Java optimizer works exclusively to reduce the overall execution time of the running program (it ignores other kinds of optimizations such as reducing code size).

The optimizer is passed the root node of the class hierarchy tree from the semantic analyzer. From the root node, we can of course reach any other class tree node. Each node contains meta-information about the class as well as a link to the AST node representing that particular class. Any expression AST nodes embedded within the class AST node is annotated with its type, which was done by the semantic analyzer. The optimizer can assume that the program it is passed is a legal, Bantam program. The optimizer traverses the class hierarchy tree and AST nodes (perhaps, multiple times), converts it to a set of control flow graphs (one for each method), and then optimizes these control flow graphs.

The optimizer will need to use global (intraprocedural) and local analyses to perform its transformations. In particular, the optimizer will need to perform one data flow analysis to find redundant expressions within the program. Your optimizer will not need to do any whole program (interprocedural) analyses (unless you want to do this for extra credit).

For more information on the theory of optimization see a compiler textbook [2, 22, 8, 19].

123

11.2 Files and Directories

All of the source code you will need to edit is within the opt package. In particular, you will need to complete the file, **Optimizer.java**, which contains some skeleton code for performing optimization. You may also find it helpful to add other Java files to this package. For example, you might write a separate class (in a separate Java file) for traversing AST nodes.

You will also need to make use of the files in the cfg package, which contains the Java classes for representating a control flow graph. As in previous assignments, you will find many of the files in the ast, util, and visitor packages useful. Note: you should not need to modify any of the files within these packages.

11.3 Optimizer Details

The optimizer converts the AST representation of the program into a set of control flow graphs. As it does this conversion, it performs some high-level optimizations on the AST. Once it has the control flow graphs and then performs some low-level optimizations (at least one) on each control graph.

Control flow graph. One difference between this assignment and previous assignments is that it uses a different representation for the program: control flow graphs. A control flow graph has nodes containing basic blocks and directed edges that indicate control transfers between blocks. A basic block is just a list of instructions with no interior control flow (only the last instruction can be a control instruction, control can only transfer to the first instruction). See Section 6 for more details on control flow graphs, basic blocks, and three-address code instructions.

Optimization steps. Your optimizer will need to perform the following three steps:

1. Convert the AST representation of the program into a set of control flow graphs, one for each user-defined method and initialization subroutine.

2. During conversion in the task above, perform some high-level optimizations such as dynamic dispatch conversion (discussed below).

3. Perform at least one low-level optimization on each control flow graph.

For extra credit, you can add additional high-level and/or low-level optimizations.

The level of optimization is specified by the user on the command-line. You should use this level to determine how many of the steps above your optimizer performs. For example, a level of 0 specifies no optimization at all. A level of 1 specifies that only step 1 above is done. A level of 2 specifies that only steps 1 and 2 above are done. A level of 3 specifies that steps 1, 2, and 3, are done. Finally, a level of 4 specifies that steps 1-3 are done plus any additional (extra credit) optimizations.

High-level optimization. Your optimizer will need to perform two high-level optimizations on the AST as it translates the AST into a control flow graph. In particular, your optimizer will need to do the following optimizations:

- **Dynamic dispatch conversion.** Your optimizer should convert dynamic dispatches into direct calls when the called method is known statically. For example, a dispatch to the method **putString** on a reference object with type **TextIO** will always result in a call of the method **TextIO.putString**. In this case, you should convert the dispatch into a direct call.

- **Divide by zero check removal.** Your optimizer should remove a divide by zero check when the right-hand operand is a non-zero constant. In this case, the check is unnecessary. This should be done for both division and modulus operations.

Low-level optimization. You will also need to perform some low-level optimization as well. In particular, you will need to perform a data flow analysis on each control flow graph in order to find redundant expressions. You will use this analysis to perform common subexpression elimination in each control flow block both within a basic block and across basic blocks (although within a single subroutine). See a compiler textbook for more details on data flow analysis [2, 22, 8, 19].

Extra credit. For extra credit, you can add other high-level or low-level optimizations. For example, the reference compiler also does constant folding, constant propagation, dead code elimination, and dead assignment elimination. The amount of extra credit will vary greatly based on the amount of work that you do as well as the achieved speedups of your optimizer.

11.4   Testing the Optimizer

You can test your optimizer in one of two ways. First, you can supply the "-so" flag to the Bantam compiler, which stops compilation after optimization and prints out each control flow graph. You

can compare this output with the corresponding output from the reference compiler. If the output is different then this probably indicates an error in your optimizer. A second way to test your optimizer is to use the provided class files for performing lexical, syntactic, and semantic analyses, and code generation, and fully compile some programs. If this produces incorrect target code, then this probably indicates an error in your optimizer.

Important: take testing seriously in this project. A optimizer is quite complicated, so be prepared to spend a significant amount of time debugging. Remember, testing is a critical part of the software design process. There is a strong correlation between students with working compilers and students who thoroughly test their code.

As part of the assignment you should submit a Bantam Java source program (**Test.btm**) that tests your optimizer. This file should contain a working Bantam source program that tests all (or most) of the source language features. It should also contain optimizable code. You should fully compile this program and test all of the optimization levels.

11.5   Last Words

Remember to start early on this project. It will take a significant amount of time and the sooner you finish the more time you can spend testing your optimizer. Good luck and remember to have fun!

Figure 15: Bantam Java interpreter phases. The interpreter phase (which replaces code generation) is shaded.

12 (Optional) Project 6: Building a Bantam Java Interpreter

In this optional project, you will build an interpreter that can execute Bantam Java programs. Rather than translate the source program to target code, the interpreter will instead execute the program directly (*i.e.*, interpret it). You will need to complete the code within the interp package of your compiler source code. Make sure that you have read Section 3 on the Bantam Java programming language and Section 4 on the Bantam Java compiler before starting this project.

12.1 Overview

The Bantam Java compiler has support for interpreting programs in addition to compiling programs. As shown in Figure 15, when in interpreter mode the interpreter simply relaces the code generator. It is the final phase and it executes (*i.e.*, interprets) the Bantam Java source program.

The interpreter is passed the root node of the class hierarchy tree from the semantic analyzer. (Note: optimization is not supported when in interpreter mode so the semantic analyzer always preceeds the interpreter.) From the root node, we can of course reach any other class tree node. Each node contains meta-information about the class as well as a link to the AST node representing that particular class. Any expression AST nodes embedded within the class AST node is annotated with its type, which was done by the semantic analyzer. The code generator can assume that the program it is passed is a legal, Bantam Java program. The code generator traverses the class hierarchy tree and AST nodes (probably many times for most programs) and interprets them.

12.2 Files and Directories

All of the source code you will need to edit is within the interp package. In particular, you will need to complete the file, **Interpreter.java**, which contains some skeleton code for interpreting Bantam Java programs. You may also find it helpful to add other Java files to this package. For example, you might write a separate class (in a separate Java file) for traversing AST nodes. In addition,

127

you will find many of the files in the ast, util, and visitor packages useful, although you should not need to modify any of them.

12.3   Interpreter Details

To build the interpreter you will need to figure out how to handle the following:

- **Code interpretation.** You will of course need some way to interpret (*i.e.*, execute) Bantam Java source code. You will probably want to build a visitor for this purpose. Every time a method or field initialization expression is executed, you can visit the corresponding part of the AST and interpret this source code.

- **Variable values.** You will need to keep track of the values of variables. You might leverage the symbol table from the util package for keeping track of variable values. Be careful though, symbol tables in the interpreter need to be per object/method, and not per class.

- **Objects.** You will need some class that can represent the Bantam Java objects created in the source program. You can decide whether you want the same representation for built-in and user-defined Bantam Java objects. You could potentially combine this representation with the visitor for interpreting methods and field initializations.

You will also need to think carefully about how you are going to deal with things such as the runtime stack, object heap, and garbage collection. If you design your interpreter carefully, you may be able to leverage the Java Virtual Machine's implementation of these three components. Otherwise, you will need to build these into your interpreter.

12.4   Testing the Interpreter

To test your interpreter, you can supply the "-int" flag (short for interpret) to the Bantam Java compiler, which causes it to interpret rather than compile the program. Using the provided classes for performing lexical, syntactic, and semantic analyses, you will be able to execute some programs and test your implementation. If the behavior of any the interpreted programs is different than the behavior exhibited by the reference compiler then this probably indicates an error in your interpreter.

Important: take testing seriously in this project. An interpreter is quite complicated, so be prepared to spend a significant amount of time debugging. Remember, testing is a critical part of the software design process. There is a strong correlation between students with working compilers and students who thoroughly test their code.

12.5 Last Words

Remember to start early on this project. It will take a significant amount of time and the sooner you finish the more time you can spend testing your interpreter. Good luck and remember to have fun!

13 Closing Remarks

Hopefully if you have reached this point then you have successfully implemented your very own compiler. If so, congratulations! There are few accomplishments in an undergraduate computer science program as significant as implementing a compiler. Nice work.

We hope that you have enjoyed the Bantam Java course projects and found this manual helpful and informative. Please let us know if you found any bugs or errors along the way, in either the Bantam Java toolset or in this manual. In addition, please contact us if you have questions about any part of this work or if you have requests for extensions. Finally, we welcome contributions from others, so please let us know if you would like to contribute in some way. To contact us you can email the first author at *corliss@hws.edu* or you can visit *http://www.bantamjava.com*, which has a link to an online bulletin board for submitting questions, comments, or requests.

Thanks,

Marc Corliss, David Furcy, and E Christopher Lewis

References

[1] GCC, the GNU compiler collection. URL: *http://gcc.gnu.org*.

[2] Alfred V. Aho, Monica S. Lam, Ravi Sethi, and Jeffrey D. Ullman. *Compilers: Principles, Techniques, & Tools*. Addison-Wesley, 2nd edition, 2007.

[3] Alexander Aiken. Cool: A portable project for teaching compiler construction. *SIGPLAN Notices*, 31(7):19–24, 1996.

[4] Andrew Appel. *Modern Compiler Implementation in Java*. Cambridge, 1st edition, 1998.

[5] Ken Arnold, James Gosling, and David Holmes. *The Java Programming Language*. Addison-Wesley, 4th edition, 2005.

[6] Elliot Berk. JavaCC documentation. URL: *https://javacc.dev.java.net/doc/docindex.html*.

[7] Elliot Berk. *JLex: A lexical analyzer generator for Java*, 1997. URL: *http://www.cs.princeton.edu/~appel/modern/java/JLex/current/manual.html*.

[8] Keith D. Cooper and Linda Torczon. *Engineering a Compiler*. Morgan Kaufman, 2004.

[9] Marc L. Corliss. *Larc – A Little Architecture for the Classroom: Lab Manual*, 2009. URL: *http://math.hws.edu/larc*.

[10] Marc L. Corliss and Robert Hendry. Larc: A little architecture for the classroom. *Journal of Computing Sciences in Small Colleges*, 24(6):15–20, 2009.

[11] Marc L. Corliss and E Christopher Lewis. Bantam: A customizable, Java-based, classroom compiler. In *Proc. of the 39th SIGCSE Tech. Symp. on Computer Science Education*, Mar. 2008.

[12] David Eck. *Introduction to Programming using Java - Part 1*. Creative Commons, 5th edition, 2006.

[13] Erich Gamma, Richard Helm, Ralph Johnson, and John Vlissides. *Design Patterns: Elements of Reusable Object-Oriented Software*. Addison-Wesley, 1995.

[14] Scott Hudson. *Cup User's Manual*, 1999. URL: *http://www.cs.princeton.edu/~appel/modern/java/CUP/manual.html*.

[15] Jim Larus. Spim s20: A mips r2000 simulator. Technical Report TR-966, Computer Sciences Department, University of Wisconsin-Madison, Sep. 1990.

[16] Tim Lindholm and Frank Yellin. *The Java Virtual Machine Specification*. Sun Microsystems, Inc., 2nd edition, 1999. URL: *http://java.sun.com/docs/books/jvms/second_edition/html/VMSpecTOC.doc.html*.

[17] J. Meyer. *Java Virtual Machine*. O'Reilly Associates, 1997.

[18] J. Meyer and D. Reynaud. Jasmin home page. URL: *http://jasmin.sourceforge.net/*, 2005.

[19] Thomas W. Parsons. *An Introduction to Compiler Construction*. W H Freeman and Co, 1992.

[20] Sun Microsystems, Inc. *javadoc - The Java API Documentation Generator*, 2002. URL: *http://java.sun.com/j2se/1.5.0/docs/tooldocs/solaris/javadoc.html*.

[21] Niklaus Wirth. *Compiler Construction*. Addison-Wesley, 1996.

[22] Andrew Appel with Jens Palsberg. *Modern Compiler Implementation in Java*. Cambridge, 2nd edition, 2002.